CUDNN LIBRARY

User Guide

Chapter 1.
INTRODUCTION

NVIDIA® cuDNN is a GPU-accelerated library of primitives for deep neural networks. It provides highly tuned implementations of routines arising frequently in DNN applications:

▸ Convolution forward and backward, including cross-correlation
▸ Pooling forward and backward
▸ Softmax forward and backward
▸ Neuron activations forward and backward:

    ▸ Rectified linear (ReLU)
    ▸ Sigmoid
    ▸ Hyperbolic tangent (TANH)

▸ Tensor transformation functions
▸ LRN, LCN and batch normalization forward and backward

cuDNN's convolution routines aim for performance competitive with the fastest GEMM (matrix multiply) based implementations of such routines while using significantly less memory.

cuDNN features customizable data layouts, supporting flexible dimension ordering, striding, and subregions for the 4D tensors used as inputs and outputs to all of its routines. This flexibility allows easy integration into any neural network implementation and avoids the input/output transposition steps sometimes necessary with GEMM-based convolutions.

cuDNN offers a context-based API that allows for easy multithreading and (optional) interoperability with CUDA streams.

Chapter 2.
GENERAL DESCRIPTION

2.1. Programming Model

The cuDNN Library exposes a Host API but assumes that for operations using the GPU, the necessary data is directly accessible from the device.

An application using cuDNN must initialize a handle to the library context by calling **cudnnCreate()**. This handle is explicitly passed to every subsequent library function that operates on GPU data. Once the application finishes using cuDNN, it can release the resources associated with the library handle using **cudnnDestroy()**. This approach allows the user to explicitly control the library's functioning when using multiple host threads, GPUs and CUDA Streams. For example, an application can use **cudaSetDevice()** to associate different devices with different host threads and in each of those host threads, use a unique cuDNN handle which directs library calls to the device associated with it. cuDNN library calls made with different handles will thus automatically run on different devices. The device associated with a particular cuDNN context is assumed to remain unchanged between the corresponding **cudnnCreate()** and **cudnnDestroy()** calls. In order for the cuDNN library to use a different device within the same host thread, the application must set the new device to be used by calling **cudaSetDevice()** and then create another cuDNN context, which will be associated with the new device, by calling **cudnnCreate()**.

2.2. Notation

As of CUDNN v4 we have adopted a mathematicaly-inspired notation for layer inputs and outputs using **x,y,dx,dy,b,w** for common layer parameters. This was done to improve readability and ease of understanding of parameters meaning. All layers now follow a uniform convention that during inference

x = layerFunction(y, otherParams).

And during backpropagation

(dx, dOtherParams) = layerFunctionGradient(x,y,dy,otherParams)

For convolution the notation is

```
y = x*w+b
```

where **w** is the matrix of filter weights, **x** is the previous layer's data (during inference), **y** is the next layer's data, **b** is the bias and ***** is the convolution operator. In backpropagation routines the parameters keep their meanings. **dx,dy,dw,db** always refer to the gradient of the final network error function with respect to a given parameter. So **dy** in all backpropagation routines always refers to error gradient backpropagated through the network computation graph so far. Similarly other parameters in more specialized layers, such as, for instance, **dMeans** or **dBnBias** refer to gradients of the loss function wrt those parameters.

> **w** is used in the API for both the width of the **x** tensor and convolution filter matrix. To resolve this ambiguity we use **w** and **filter** notation interchangeably for convolution filter weight matrix. The meaning is clear from the context since the layer width is always referenced near it's height.

2.3. Tensor Descriptor

The cuDNN Library describes data holding images, videos and any other data with contents with a generic n-D tensor defined with the following parameters :

▸ a dimension **dim** from 3 to 8
▸ a data type (32-bit floating point, 64 bit-floating point, 16 bit floating point...)
▸ **dim** integers defining the size of each dimension
▸ **dim** integers defining the stride of each dimension (e.g the number of elements to add to reach the next element from the same dimension)

The first two dimensions define respectively the batch number **n** and the number of features maps **c**. This tensor definition allows for example to have some dimensions overlapping each others within the same tensor by having the stride of one dimension smaller than the product of the dimension and the stride of the next dimension. In cuDNN, unless specified otherwise, all routines will support tensors with overlapping dimensions for forward pass input tensors, however, dimensions of the output tensors cannot overlap. Even though this tensor format supports negative strides (which can be useful for data mirroring), cuDNN routines do not support tensors with negative strides unless specified otherwise.

2.3.1. WXYZ Tensor Descriptor

Tensor descriptor formats are identified using acronyms, with each letter referencing a corresponding dimension. In this document, the usage of this terminology implies :

▸ all the strides are strictly positive
▸ the dimensions referenced by the letters are sorted in decreasing order of their respective strides

2.3.2. 4-D Tensor Descriptor

A 4-D Tensor descriptor is used to define the format for batches of 2D images with 4 letters : N,C,H,W for respectively the batch number, the number of feature maps, the height and the width. The letters are sorted in decreasing order of the strides. The commonly used 4-D tensor formats are :

- ▸ NCHW
- ▸ NHWC
- ▸ CHWN

2.3.3. 5-D Tensor Description

A 5-D Tensor descriptor is used to define the format of batch of 3D images with 5 letters : N,C,D,H,W for respectively the batch number, the number of feature maps, the depth, the height and the width. The letters are sorted in descreasing order of the strides. The commonly used 5-D tensor formats are called :

- ▸ NCDHW
- ▸ NDHWC
- ▸ CDHWN

2.3.4. Fully-packed tensors

A tensor is defined as **XYZ-fully-packed** if and only if :

- ▸ the number of tensor dimensions is equal to the number of letters preceding the **fully-packed** suffix.
- ▸ the stride of the i-th dimension is equal to the product of the (i+1)-th dimension by the (i+1)-th stride.
- ▸ the stride of the last dimension is 1.

2.3.5. Partially-packed tensors

The partially 'XYZ-packed' terminology only applies in a context of a tensor format described with a superset of the letters used to define a partially-packed tensor. A WXYZ tensor is defined as **XYZ-packed** if and only if :

- ▸ the strides of all dimensions NOT referenced in the -packed suffix are greater or equal to the product of the next dimension by the next stride.
- ▸ the stride of each dimension referenced in the -packed suffix in position i is equal to the product of the (i+1)-st dimension by the (i+1)-st stride.
- ▸ if last tensor's dimension is present in the -packed suffix, it's stride is 1.

For example a NHWC tensor WC-packed means that the c_stride is equal to 1 and w_stride is equal to c_dim x c_stride. In practice, the -packed suffix is usually with slowest changing dimensions of a tensor but it is also possible to refer to a NCHW tensor that is only N-packed.

2.3.6. Spatially packed tensors

Spatially-packed tensors are defined as partially-packed in spatial dimensions.

For example a spatially-packed 4D tensor would mean that the tensor is either NCHW HW-packed or CNHW HW-packed.

2.3.7. Overlapping tensors

A tensor is defined to be overlapping if a iterating over a full range of dimensions produces the same address more than once.

In practice an overlapped tensor will have stride[i-1] < stride[i]*dim[i] for some of the i from [1,nbDims] interval.

2.4. Thread Safety

The library is thread safe and its functions can be called from multiple host threads, even with the same handle. When sharing a handle across host threads, extreme care needs to be taken to ensure that any changes to the handle configuration in one thread do not adversely affect cuDNN function calls in others. This is especially true for the destruction of the handle. It is not recommended that multiple threads share the same cuDNN handle.

2.5. Reproducibility (determinism)

By design, most of cuDNN's routines from a given version generate the same bit-wise results at every run when executed on GPUs with the same architecture and the same number of SMs. However, bit-wise reproducibility is not guaranteed across versions, as the implementation of a given routine may change. With the current release, the following routines do not guarantee reproducibility because they use atomic add operations:

- ▶ `cudnnConvolutionBackwardFilter` when `CUDNN_CONVOLUTION_BWD_FILTER_ALGO_0` or `CUDNN_CONVOLUTION_BWD_FILTER_ALGO_3` is used
- ▶ `cudnnConvolutionBackwardData` when `CUDNN_CONVOLUTION_BWD_DATA_ALGO_0` is used
- ▶ `cudnnPoolingBackward` when `CUDNN_POOLING_MAX` is used

2.6. Scaling parameters `alpha` and `beta`

Many cuDNN routines like `cudnnConvolutionForward` take pointers to scaling factors (in host memory) used to blend the source value with prior value in the destination tensor as follows: dstValue = alpha[0]*srcValue + beta[0]*priorDstValue. When beta[0] is zero, the output is not read and can contain any uninitialized data (including NaN). The

storage data type for alpha[0], beta[0] is float for HALF and SINGLE tensors, and double for DOUBLE tensors. These parameters are passed using a host memory pointer.

> 💬 For improved performance it is advised to use beta[0] = 0.0. Use a non-zero value for beta[0] only when blending with prior values stored in the output tensor is needed.

2.7. GPU and driver requirements

cuDNN v5.0 supports NVIDIA GPUs of compute capability 3.0 and higher and requires an NVIDIA Driver compatible with CUDA Toolkit 7.5 (for ARM platforms CUDA Toolkit 7.0).

2.8. Backward compatibility and deprecation policy

When we want to change the API of an already implemented functionnality, we will create a new routine with the same name appended with a suffix _v**<n>** where **n** represents the cuDNN version where that new API was introduced. This way, we can guarantee backward compatibility with the previous version of cuDNN Library. However, in order to facilitate the maintenance and keep the API as clean as possible, we want to be able to deprecate and then eliminate the previous API entries and eventually remove also the new suffixed ones to only keep the unsuffixed entries mapped to the new API. To achieve this, we introduce the following deprecation policy:

▸ the former API entries will be renamed with a suffix _v**<f>** where **f** represents a cuDNN version anterior to the new API.
▸ the unsuffixed API entry will be mapped to the deprecated entry until the former API entry with an anterior suffix _v**<f>** appeared in the API header file.
▸ the deprecated former API entry with an anterior suffix _v**<f>** and new API entry with suffix _v**<n>** will be present for at least one release
▸ In one of the subsequent cuDNN release, the suffixed entries of a given entry will be removed simultaneously.

To sum up, when a routine appears in two forms, one with a suffix and one with no suffix, it means that the non-suffixed entry is deprecated. In this case, it is stongly advised to users to use the new suffixed API entry to guarantee backwards compatibility with the next cuDNN release. When a routine appears with multiples suffixes, the unsuffixed API entry will be mapped to the higher number. In that case it is strongly advised to use the non-suffixed API entry to guarantee backward compatibiliy with the next cuDNN release.

Chapter 3.
CUDNN DATATYPES REFERENCE

This chapter describes all the types and enums of the cuDNN library API.

3.1. cudnnHandle_t

cudnnHandle_t is a pointer to an opaque structure holding the cuDNN library context. The cuDNN library context must be created using **cudnnCreate()** and the returned handle must be passed to all subsequent library function calls. The context should be destroyed at the end using **cudnnDestroy()**. The context is associated with only one GPU device, the current device at the time of the call to **cudnnCreate()**. However multiple contexts can be created on the same GPU device.

3.2. cudnnStatus_t

cudnnStatus_t is an enumerated type used for function status returns. All cuDNN library functions return their status, which can be one of the following values:

Value	Meaning
CUDNN_STATUS_SUCCESS	The operation completed successfully.
CUDNN_STATUS_NOT_INITIALIZED	The cuDNN library was not initialized properly. This error is usually returned when a call to **cudnnCreate()** fails or when **cudnnCreate()** has not been called prior to calling another cuDNN routine. In the former case, it is usually due to an error in the CUDA Runtime API called by **cudnnCreate()** or by an error in the hardware setup.
CUDNN_STATUS_ALLOC_FAILED	Resource allocation failed inside the cuDNN library. This is usually caused by an internal **cudaMalloc()** failure. To correct: prior to the function call, deallocate previously allocated memory as much as possible.

Value	Meaning
`CUDNN_STATUS_BAD_PARAM`	An incorrect value or parameter was passed to the function. To correct: ensure that all the parameters being passed have valid values.
`CUDNN_STATUS_ARCH_MISMATCH`	The function requires a feature absent from the current GPU device. Note that cuDNN only supports devices with compute capabilities greater than or equal to 3.0. To correct: compile and run the application on a device with appropriate compute capability.
`CUDNN_STATUS_MAPPING_ERROR`	An access to GPU memory space failed, which is usually caused by a failure to bind a texture. To correct: prior to the function call, unbind any previously bound textures. Otherwise, this may indicate an internal error/bug in the library.
`CUDNN_STATUS_EXECUTION_FAILED`	The GPU program failed to execute. This is usually caused by a failure to launch some cuDNN kernel on the GPU, which can occur for multiple reasons. To correct: check that the hardware, an appropriate version of the driver, and the cuDNN library are correctly installed. Otherwise, this may indicate a internal error/bug in the library.
`CUDNN_STATUS_INTERNAL_ERROR`	An internal cuDNN operation failed.
`CUDNN_STATUS_NOT_SUPPORTED`	The functionality requested is not presently supported by cuDNN.
`CUDNN_STATUS_LICENSE_ERROR`	The functionality requested requires some license and an error was detected when trying to check the current licensing. This error can happen if the license is not present or is expired or if the environment variable NVIDIA_LICENSE_FILE is not set properly.

3.3. cudnnTensorDescriptor_t

cudnnCreateTensorDescriptor_t is a pointer to an opaque structure holding the description of a generic n-D dataset. **cudnnCreateTensorDescriptor()** is used to create one instance, and one of the routines **cudnnSetTensorNdDescriptor()**, **cudnnSetTensor4dDescriptor()** or **cudnnSetTensor4dDescriptorEx()** must be used to initialize this instance.

3.4. cudnnFilterDescriptor_t

cudnnFilterDescriptor_t is a pointer to an opaque structure holding the description of a filter dataset. **cudnnCreateFilterDescriptor()** is used to create one instance, and **cudnnSetFilterDescriptor()** must be used to initialize this instance.

3.5. cudnnConvolutionDescriptor_t

cudnnConvolutionDescriptor_t is a pointer to an opaque structure holding the description of a convolution operation. **cudnnCreateConvolutionDescriptor()** is used to create one instance, and **cudnnSetConvolutionNdDescriptor()** or **cudnnSetConvolution2dDescriptor()** must be used to initialize this instance.

3.6. cudnnNanPropagation_t

cudnnNanPropagation_t is an enumerated type used to indicate if some routines should propagate **Nan** numbers. This enumerated type is used as a field for the **cudnnActivationDescriptor_t** descriptor and **cudnnPoolingDescriptor_t** descriptor.

Value	Meaning
CUDNN_NOT_PROPAGATE_NAN	**Nan** numbers are not propagated
CUDNN_PROPAGATE_NAN	**Nan** numbers are propagated

3.7. cudnnActivationDescriptor_t

cudnnActivationDescriptor_t is a pointer to an opaque structure holding the description of a activation operation. **cudnnCreateActivationDescriptor()** is used to create one instance, and **cudnnSetActivationDescriptor()** must be used to initialize this instance.

3.8. cudnnPoolingDescriptor_t

cudnnPoolingDescriptor_t is a pointer to an opaque structure holding the description of a pooling operation. **cudnnCreatePoolingDescriptor()** is used to create one instance, and **cudnnSetPoolingNdDescriptor()** or **cudnnSetPooling2dDescriptor()** must be used to initialize this instance.

3.9. cudnnOpTensorOp_t

cudnnOpTensorOp_t is an enumerated type used to indicate the tensor operation to be used by the **cudnnOpTensor()** routine. This enumerated type is used as a field for the **cudnnOpTensorDescriptor_t** descriptor.

Value	Meaning
CUDNN_OP_TENSOR_ADD	The operation to be performed is addition
CUDNN_OP_TENSOR_MUL	The operation to be performed is multiplication
CUDNN_OP_TENSOR_MIN	The operation to be performed is a minimum comparison
CUDNN_OP_TENSOR_MAX	The operation to be performed is a maximum comparison

3.10. cudnnOpTensorDescriptor_t

cudnnOpTensorDescriptor_t is a pointer to an opaque structure holding the description of a tensor operation, used as a parameter to **cudnnOpTensor()**. **cudnnCreateOpTensorDescriptor()** is used to create one instance, and **cudnnSetOpTensorDescriptor()** must be used to initialize this instance.

3.11. cudnnDataType_t

cudnnDataType_t is an enumerated type indicating the data type to which a tensor descriptor or filter descriptor refers.

Value	Meaning
CUDNN_DATA_FLOAT	The data is 32-bit single-precision floating point (**float**).
CUDNN_DATA_DOUBLE	The data is 64-bit double-precision floating point (**double**).
CUDNN_DATA_HALF	The data is 16-bit floating point.

3.12. cudnnTensorFormat_t

cudnnTensorFormat_t is an enumerated type used by **cudnnSetTensor4dDescriptor()** to create a tensor with a pre-defined layout.

Value	Meaning
CUDNN_TENSOR_NCHW	This tensor format specifies that the data is laid out in the following order: image, features map,

Value	Meaning
	rows, columns. The strides are implicitly defined in such a way that the data are contiguous in memory with no padding between images, feature maps, rows, and columns; the columns are the inner dimension and the images are the outermost dimension.
`CUDNN_TENSOR_NHWC`	This tensor format specifies that the data is laid out in the following order: image, rows, columns, features maps. The strides are implicitly defined in such a way that the data are contiguous in memory with no padding between images, rows, columns, and features maps; the feature maps are the inner dimension and the images are the outermost dimension.

3.13. cudnnConvolutionMode_t

cudnnConvolutionMode_t is an enumerated type used by **cudnnSetConvolutionDescriptor()** to configure a convolution descriptor. The filter used for the convolution can be applied in two different ways, corresponding mathematically to a convolution or to a cross-correlation. (A cross-correlation is equivalent to a convolution with its filter rotated by 180 degrees.)

Value	Meaning
`CUDNN_CONVOLUTION`	In this mode, a convolution operation will be done when applying the filter to the images.
`CUDNN_CROSS_CORRELATION`	In this mode, a cross-correlation operation will be done when applying the filter to the images.

3.14. cudnnConvolutionFwdPreference_t

cudnnConvolutionFwdPreference_t is an enumerated type used by **cudnnGetConvolutionForwardAlgorithm()** to help the choice of the algorithm used for the forward convolution.

Value	Meaning
`CUDNN_CONVOLUTION_FWD_NO_WORKSPACE`	In this configuration, the routine **cudnnGetConvolutionForwardAlgorithm()** is guaranteed to return an algorithm that does not require any extra workspace to be provided by the user.
`CUDNN_CONVOLUTION_FWD_PREFER_FASTEST`	In this configuration, the routine **cudnnGetConvolutionForwardAlgorithm()** will return the fastest algorithm regardless how much workspace is needed to execute it.

Value	Meaning
`CUDNN_CONVOLUTION_FWD_SPECIFY_WORKSPACE_LIMIT`	In this configuration, the routine `cudnnGetConvolutionForwardAlgorithm()` will return the fastest algorithm that fits within the memory limit that the user provided.

3.15. cudnnConvolutionFwdAlgo_t

`cudnnConvolutionFwdAlgo_t` is an enumerated type that exposes the different algorithms available to execute the forward convolution operation.

Value	Meaning
`CUDNN_CONVOLUTION_FWD_ALGO_IMPLICIT_GEMM`	This algorithm expresses the convolution as a matrix product without actually explicitly form the matrix that holds the input tensor data.
`CUDNN_CONVOLUTION_FWD_ALGO_IMPLICIT_PRECOMP_GEMM`	This algorithm expresses the convolution as a matrix product without actually explicitly form the matrix that holds the input tensor data, but still needs some memory workspace to precompute some indices in order to facilitate the implicit construction of the matrix that holds the input tensor data
`CUDNN_CONVOLUTION_FWD_ALGO_GEMM`	This algorithm expresses the convolution as an explicit matrix product. A significant memory workspace is needed to store the matrix that holds the input tensor data.
`CUDNN_CONVOLUTION_FWD_ALGO_DIRECT`	This algorithm expresses the convolution as a direct convolution (e.g without implicitly or explicitly doing a matrix multiplication).
`CUDNN_CONVOLUTION_FWD_ALGO_FFT`	This algorithm uses a Fast-Fourier Transform approach to compute the convolution. A significant memory workspace is needed to store intermediate results.
`CUDNN_CONVOLUTION_FWD_ALGO_FFT_TILING`	This algorithm uses a Fast-Fourier Transform approach but splits the inputs into 32x32 tiles. A significant memory workspace is needed to store intermediate results but significantly less than `CUDNN_CONVOLUTION_FWD_ALGO_FFT` for big size images.
`CUDNN_CONVOLUTION_FWD_ALGO_WINOGRAD`	This algorithm uses a Winograd Transform approach to compute the convolution. A reasonably sized workspace is needed to store intermediate results.

3.16. cudnnConvolutionFwdAlgoPerf_t

`cudnnConvolutionFwdAlgoPerf_t` is a structure containing performance results returned by `cudnnFindConvolutionForwardAlgorithm()`.

Member Name	Explanation
`cudnnConvolutionFwdAlgo_t algo`	The algorithm run to obtain the associated performance metrics.
`cudnnStatus_t status`	If any error occurs during the workspace allocation or timing of `cudnnConvolutionForward()`, this status will represent that error. Otherwise, this status will be the return status of `cudnnConvolutionForward()`. ▸ `CUDNN_STATUS_ALLOC_FAILED` if any error occured during workspace allocation or if provided workspace is insufficient. ▸ `CUDNN_STATUS_INTERNAL_ERROR` if any error occured during timing calculations or workspace deallocation. ▸ Otherwise, this will be the return status of `cudnnConvolutionForward()`.
`float time`	The execution time of `cudnnConvolutionForward()` (in milliseconds).
`size_t memory`	The workspace size (in bytes).

3.17. cudnnConvolutionBwdFilterPreference_t

`cudnnConvolutionBwdFilterPreference_t` is an enumerated type used by `cudnnGetConvolutionBackwardFilterAlgorithm()` to help the choice of the algorithm used for the backward filter convolution.

Value	Meaning
`CUDNN_CONVOLUTION_BWD_FILTER_NO_WORKSPACE`	In this configuration, the routine `cudnnGetConvolutionBackwardFilterAlgorithm()` is guaranteed to return an algorithm that does not require any extra workspace to be provided by the user.
`CUDNN_CONVOLUTION_BWD_FILTER_PREFER_FASTEST`	In this configuration, the routine `cudnnGetConvolutionBackwardFilterAlgorithm()` will return the fastest algorithm regardless how much workspace is needed to execute it.
`CUDNN_CONVOLUTION_BWD_FILTER_SPECIFY_WORKSPACE_LIMIT`	In this configuration, the routine `cudnnGetConvolutionBackwardFilterAlgorithm()` will return the fastest algorithm that fits within the memory limit that the user provided.

3.18. cudnnConvolutionBwdFilterAlgo_t

`cudnnConvolutionBwdFilterAlgo_t` is an enumerated type that exposes the different algorithms available to execute the backward filter convolution operation.

Value	Meaning
`CUDNN_CONVOLUTION_BWD_FILTER_ALGO_0`	This algorithm expresses the convolution as a sum of matrix product without actually explicitly form the matrix that holds the input tensor data. The sum is done using atomic adds operation, thus the results are non-deterministic.
`CUDNN_CONVOLUTION_BWD_FILTER_ALGO_1`	This algorithm expresses the convolution as a matrix product without actually explicitly form the matrix that holds the input tensor data. The results are deterministic.
`CUDNN_CONVOLUTION_BWD_FILTER_ALGO_FFT`	This algorithm uses a Fast-Fourier Transform approach to compute the convolution. A significant memory workspace is needed to store intermediate results. The results are deterministic.
`CUDNN_CONVOLUTION_BWD_FILTER_ALGO_3`	This algorithm is similar to `CUDNN_CONVOLUTION_BWD_FILTER_ALGO_0` but uses some small workspace to precomputes some indices. The results are also non-deterministic.

3.19. cudnnConvolutionBwdFilterAlgoPerf_t

`cudnnConvolutionBwdFilterAlgoPerf_t` is a structure containing performance results returned by `cudnnFindConvolutionBackwardFilterAlgorithm()`.

Member Name	Explanation
`cudnnConvolutionBwdFilterAlgo_t algo`	The algorithm run to obtain the associated performance metrics.
`cudnnStatus_t status`	If any error occurs during the workspace allocation or timing of `cudnnConvolutionBackwardFilter_v3()`, this status will represent that error. Otherwise, this status will be the return status of `cudnnConvolutionBackwardFilter_v3()`. ▸ `CUDNN_STATUS_ALLOC_FAILED` if any error occured during workspace allocation or if provided workspace is insufficient. ▸ `CUDNN_STATUS_INTERNAL_ERROR` if any error occured during timing calculations or workspace deallocation. ▸ Otherwise, this will be the return status of `cudnnConvolutionBackwardFilter_v3()`.
`float time`	The execution time of `cudnnConvolutionBackwardFilter_v3()` (in milliseconds).
`size_t memory`	The workspace size (in bytes).

3.20. cudnnConvolutionBwdDataPreference_t

cudnnConvolutionBwdDataPreference_t is an enumerated type used by **cudnnGetConvolutionBackwardDataAlgorithm()** to help the choice of the algorithm used for the backward data convolution.

Value	Meaning
CUDNN_CONVOLUTION_BWD_DATA_NO_WORKSPACE	In this configuration, the routine **cudnnGetConvolutionBackwardDataAlgorithm()** is guaranteed to return an algorithm that does not require any extra workspace to be provided by the user.
CUDNN_CONVOLUTION_BWD_DATA_PREFER_FASTEST	In this configuration, the routine **cudnnGetConvolutionBackwardDataAlgorithm()** will return the fastest algorithm regardless how much workspace is needed to execute it.
CUDNN_CONVOLUTION_BWD_DATA_SPECIFY_WORKSPACE_LIMIT	In this configuration, the routine **cudnnGetConvolutionBackwardDataAlgorithm()** will return the fastest algorithm that fits within the memory limit that the user provided.

3.21. cudnnConvolutionBwdDataAlgo_t

cudnnConvolutionBwdDataAlgo_t is an enumerated type that exposes the different algorithms available to execute the backward data convolution operation.

Value	Meaning
CUDNN_CONVOLUTION_BWD_DATA_ALGO_0	This algorithm expresses the convolution as a sum of matrix product without actually explicitly form the matrix that holds the input tensor data. The sum is done using atomic adds operation, thus the results are non-deterministic.
CUDNN_CONVOLUTION_BWD_DATA_ALGO_1	This algorithm expresses the convolution as a matrix product without actually explicitly form the matrix that holds the input tensor data. The results are deterministic.
CUDNN_CONVOLUTION_BWD_DATA_ALGO_FFT	This algorithm uses a Fast-Fourier Transform approach to compute the convolution. A significant memory workspace is needed to store intermediate results. The results are deterministic.
CUDNN_CONVOLUTION_BWD_DATA_ALGO_FFT_TILING	This algorithm uses a Fast-Fourier Transform approach but splits the inputs into 32x32 tiles. A significant memory workspace is needed to store intermediate results but significantly less than CUDNN_CONVOLUTION_BWD_DATA_ALGO_FFT for big size images. A significant memory workspace is

Value	Meaning
	needed to store intermediate results. The results are deterministic.
`CUDNN_CONVOLUTION_BWD_DATA_ALGO_WINOGRAD`	This algorithm uses a Winograd Transform approach to compute the convolution. A reasonably sized workspace is needed to store intermediate results. The results are deterministic.

3.22. cudnnConvolutionBwdDataAlgoPerf_t

cudnnConvolutionBwdDataAlgoPerf_t is a structure containing performance results returned by **cudnnFindConvolutionBackwardDataAlgorithm()**.

Member Name	Explanation
`cudnnConvolutionBwdDataAlgo_t algo`	The algorithm run to obtain the associated performance metrics.
`cudnnStatus_t status`	If any error occurs during the workspace allocation or timing of `cudnnConvolutionBackwardData_v3()`, this status will represent that error. Otherwise, this status will be the return status of `cudnnConvolutionBackwardData_v3()`. ▸ `CUDNN_STATUS_ALLOC_FAILED` if any error occured during workspace allocation or if provided workspace is insufficient. ▸ `CUDNN_STATUS_INTERNAL_ERROR` if any error occured during timing calculations or workspace deallocation. ▸ Otherwise, this will be the return status of `cudnnConvolutionBackwardData_v3()`.
`float time`	The execution time of `cudnnConvolutionBackwardData_v3()` (in milliseconds).
`size_t memory`	The workspace size (in bytes).

3.23. cudnnSoftmaxAlgorithm_t

cudnnSoftmaxAlgorithm_t is used to select an implementation of the softmax function used in **cudnnSoftmaxForward()** and **cudnnSoftmaxBackward()**.

Value	Meaning
`CUDNN_SOFTMAX_FAST`	This implementation applies the straightforward softmax operation.
`CUDNN_SOFTMAX_ACCURATE`	This implementation scales each point of the softmax input domain by its maximum value to

Value	Meaning
	avoid potential floating point overflows in the softmax evaluation.
`CUDNN_SOFTMAX_LOG`	This entry performs the Log softmax operation, avoiding overflows by scaling each point in the input domain as in `CUDNN_SOFTMAX_ACCURATE`

3.24. cudnnSoftmaxMode_t

`cudnnSoftmaxMode_t` is used to select over which data the `cudnnSoftmaxForward()` and `cudnnSoftmaxBackward()` are computing their results.

Value	Meaning
`CUDNN_SOFTMAX_MODE_INSTANCE`	The softmax operation is computed per image (N) across the dimensions C,H,W.
`CUDNN_SOFTMAX_MODE_CHANNEL`	The softmax operation is computed per spatial location (H,W) per image (N) across the dimension C.

3.25. cudnnPoolingMode_t

`cudnnPoolingMode_t` is an enumerated type passed to `cudnnSetPoolingDescriptor()` to select the pooling method to be used by `cudnnPoolingForward()` and `cudnnPoolingBackward()`.

Value	Meaning
`CUDNN_POOLING_MAX`	The maximum value inside the pooling window will be used.
`CUDNN_POOLING_AVERAGE_COUNT_INCLUDE_PADDING`	The values inside the pooling window will be averaged. The number of padded values will be taken into account when computing the average value
`CUDNN_POOLING_AVERAGE_COUNT_EXCLUDE_PADDING`	The values inside the pooling window will be averaged. The number of padded values will not be taken into account when computing the average value

3.26. cudnnActivationMode_t

`cudnnActivationMode_t` is an enumerated type used to select the neuron activation function used in `cudnnActivationForward()` and `cudnnActivationBackward()`.

Value	Meaning
`CUDNN_ACTIVATION_SIGMOID`	Selects the sigmoid function.

Value	Meaning
`CUDNN_ACTIVATION_RELU`	Selects the rectified linear function.
`CUDNN_ACTIVATION_TANH`	Selects the hyperbolic tangent function.
`CUDNN_ACTIVATION_CLIPPED_RELU`	Selects the clipped rectified linear function

3.27. cudnnLRNMode_t

`cudnnLRNMode_t` is an enumerated type used to specify the mode of operation in `cudnnLRNCrossChannelForward()` and `cudnnLRNCrossChannelBackward()`.

Value	Meaning
`CUDNN_LRN_CROSS_CHANNEL_DIM1`	LRN computation is performed across tensor's dimension dimA[1].

3.28. cudnnDivNormMode_t

`cudnnDivNormMode_t` is an enumerated type used to specify the mode of operation in `cudnnDivisiveNormalizationForward()` and `cudnnDivisiveNormalizationBackward()`.

Value	Meaning
`CUDNN_DIVNORM_PRECOMPUTED_MEANS`	The means tensor data pointer is expected to contain means or other kernel convolution values precomputed by the user. The means pointer can also be NULL, in that case it's considered to be filled with zeroes. This is equivalent to spatial LRN. Note that in the backward pass the means are treated as independent inputs and the gradient over means is computed independently. In this mode to yield a net gradient over the entire LCN computational graph the destDiffMeans result should be backpropagated through the user's means layer (which can be impelemented using average pooling) and added to the destDiffData tensor produced by cudnnDivisiveNormalizationBackward.

3.29. cudnnBatchNormMode_t

`cudnnBatchNormMode_t` is an enumerated type used to specify the mode of operation in `cudnnBatchNormalizationForwardInference()`, `cudnnBatchNormalizationForwardTraining()`, `cudnnBatchNormalizationBackward()` and `cudnnDeriveBNTensorDescriptor()` routines.

Value	Meaning
`CUDNN_BATCHNORM_PER_ACTIVATION`	Normalization is performed per-activation. This mode is intended to be used after non-convolutional network layers. In this mode bnBias and bnScale tensor dimensions are 1xCxHxW.
`CUDNN_BATCHNORM_SPATIAL`	Normalization is performed over N+spatial dimensions. This mode is intended for use after convolutional layers (where spatial invariance is desired). In this mode bnBias, bnScale tensor dimensions are 1xCx1x1.

3.30. cudnnRNNDescriptor_t

`cudnnRNNDescriptor_t` is a pointer to an opaque structure holding the description of an RNN operation. `cudnnCreateRNNDescriptor()` is used to create one instance, and `cudnnSetRNNDescriptor()` must be used to initialize this instance.

3.31. cudnnRNNMode_t

`cudnnRNNMode_t` is an enumerated type used to specify the type of network used in the `cudnnRNNForwardInference()`, `cudnnRNNForwardTraining()`, `cudnnRNNBackwardData()` and `cudnnRNNBackwardWeights()` routines.

Value	Meaning
`CUDNN_RNN_RELU`	A single-gate recurrent neural network with a ReLU activation function. In the forward pass the output h_t for a given iteration can be computed from the recurrent input h_{t-1} and the previous layer input x_t given matrices W, R and biases b_W, b_R from the following equation: $h_t = \text{ReLU}(W_i x_t + R_i h_{t-1} + b_{Wi} + b_{Ri})$ Where `ReLU(x) = max(x, 0)`.
`CUDNN_RNN_TANH`	A single-gate recurrent neural network with a tanh activation function. In the forward pass the output h_t for a given iteration can be computed from the recurrent input h_{t-1} and the previous layer input x_t given matrices W, R and biases b_W, b_R from the following equation: $h_t = \tanh(W_i x_t + R_i h_{t-1} + b_{Wi} + b_{Ri})$ Where `tanh` is the hyperbolic tangent function.
`CUDNN_LSTM`	A four-gate Long Short-Term Memory network with no peephole connections.

Value	Meaning

	In the forward pass the output h_t and cell output c_t for a given iteration can be computed from the recurrent input h_{t-1}, the cell input c_{t-1} and the previous layer input x_t given matrices W, R and biases b_W, b_R from the following equations:
	
	`i`$_t$ `= σ(W`$_i$`x`$_t$ `+ R`$_i$`h`$_{t-1}$ `+ b`$_{Wi}$ `+ b`$_{Ri}$`)`
	`f`$_t$ `= σ(W`$_f$`x`$_t$ `+ R`$_f$`h`$_{t-1}$ `+ b`$_{Wf}$ `+ b`$_{Rf}$`)`
	`o`$_t$ `= σ(W`$_o$`x`$_t$ `+ R`$_o$`h`$_{t-1}$ `+ b`$_{Wo}$ `+ b`$_{Ro}$`)`
	`c'`$_t$ `= tanh(W`$_c$`x`$_t$ `+ R`$_c$`h`$_{t-1}$ `+ b`$_{Wc}$ `+ b`$_{Rc}$`)`
	`c`$_t$ `= f`$_t$`°c`$_{t-1}$ `+ i`$_t$`°c'`$_t$
	`h`$_t$ `= o`$_t$`°tanh(c`$_t$`)`
	
	Where σ is the sigmoid operator: `σ(x) = 1 / (1 + e`$^{-x}$`)`, ° represents a point-wise multiplication and `tanh` is the hyperbolic tangent function. i_t, f_t, o_t, c'_t represent the input, forget, output and new gates respectively.

`CUDNN_GRU`	A three-gate network consisting of Gated Recurrent Units.
	
	In the forward pass the output h_t for a given iteration can be computed from the recurrent input h_{t-1} and the previous layer input x_t given matrices W, R and biases b_W, b_R from the following equations:
	
	`i`$_t$ `= σ(W`$_i$`x`$_t$ `+ R`$_i$`h`$_{t-1}$ `+ b`$_{Wi}$ `+ b`$_{Ru}$`)`
	`r`$_t$ `= σ(W`$_r$`x`$_t$ `+ R`$_r$`h`$_{t-1}$ `+ b`$_{Wr}$ `+ b`$_{Rr}$`)`
	`h'`$_t$ `= tanh(W`$_h$`x`$_t$ `+ r`$_t$`°(R`$_h$`h`$_{t-1}$ `+ b`$_{Rh}$`) +`
	  `b`$_{Wh}$`)`
	`h`$_t$ `= (1 - i`$_t$`)°h'`$_t$ `+ i`$_t$`°h`$_{t-1}$
	
	Where σ is the sigmoid operator: `σ(x) = 1 / (1 + e`$^{-x}$`)`, ° represents a point-wise multiplication and `tanh` is the hyperbolic tangent function. i_t, r_t, h'_t represent the input, reset, new gates respectively.

3.32. cudnnDirectionMode_t

cudnnDirectionMode_t is an enumerated type used to specify the recurrence pattern in the **cudnnRNNForwardInference()**, **cudnnRNNForwardTraining()**, **cudnnRNNBackwardData()** and **cudnnRNNBackwardWeights()** routines.

Value	Meaning
`CUDNN_UNIDIRECTIONAL`	The network iterates recurrently from the first input to the last.
`CUDNN_BIDIRECTIONAL`	Each layer of the the network iterates recurrently from the first input to the last and separately from the last input to the first. The outputs of the two are concatenated at each iteration giving the output of the layer.

3.33. cudnnRNNInputMode_t

cudnnRNNInputMode_t is an enumerated type used to specify the behavior of the first layer in the **cudnnRNNForwardInference()**, **cudnnRNNForwardTraining()**, **cudnnRNNBackwardData()** and **cudnnRNNBackwardWeights()** routines.

Value	Meaning
CUDNN_LINEAR_INPUT	A biased matrix multiplication is performed at the input of the first recurrent layer.
CUDNN_SKIP_INPUT	No operation is performed at the input of the first recurrent layer. If **CUDNN_SKIP_INPUT** is used the leading dimension of the input tensor must be equal to the hidden state size of the network.

3.34. cudnnDropoutDescriptor_t

cudnnDropoutDescriptor_t is a pointer to an opaque structure holding the description of a dropout operation. **cudnnCreateDropoutDescriptor()** is used to create one instance, **cudnnSetDropoutDescriptor()** is be used to initialize this instance, **cudnnDestroyDropoutDescriptor()** is be used to destroy this instance.

3.35. cudnnSpatialTransformerDescriptor_t

cudnnSpatialTransformerDescriptor_t is a pointer to an opaque structure holding the description of a spatial transformation operation. **cudnnCreateSpatialTransformerDescriptor()** is used to create one instance, **cudnnSetSpatialTransformerNdDescriptor()** is used to initialize this instance, **cudnnDestroySpatialTransformerDescriptor()** is used to destroy this instance.

3.36. cudnnSamplerType_t

cudnnSamplerType_t is an enumerated type passed to **cudnnSetSpatialTransformerNdDescriptor()** to select the sampler type to be used by **cudnnSpatialTfSamplerForward()** and **cudnnSpatialTfSamplerBackward()**.

Value	Meaning
CUDNN_SAMPLER_BILINEAR	selects the bilinear sampler

Chapter 4.
CUDNN API REFERENCE

This chapter describes the API of all the routines of the cuDNN library.

4.1. cudnnGetVersion

```
size_t cudnnGetVersion()
```

This function returns the version number of the cuDNN Library. It returns the **CUDNN_VERSION** define present in the cudnn.h header file. Starting with release R2, the routine can be used to identify dynamically the current cuDNN Library used by the application. The define **CUDNN_VERSION** can be used to have the same application linked against different cuDNN versions using conditional compilation statements.

4.2. cudnnGetErrorString

```
const char * cudnnGetErrorString(cudnnStatus_t status)
```

This function returns a human-readable character string describing the **cudnnStatus_t** enumerate passed as input parameter.

4.3. cudnnCreate

```
cudnnStatus_t cudnnCreate(cudnnHandle_t *handle)
```

This function initializes the cuDNN library and creates a handle to an opaque structure holding the cuDNN library context. It allocates hardware resources on the host and device and must be called prior to making any other cuDNN library calls. The cuDNN library context is tied to the current CUDA device. To use the library on multiple devices, one cuDNN handle needs to be created for each device. For a given device, multiple cuDNN handles with different configurations (e.g., different current CUDA streams) may be created. Because **cudnnCreate** allocates some internal resources, the release of those resources by calling **cudnnDestroy** will implicitly call **cudaDeviceSynchronize**; therefore, the recommended best practice is to call **cudnnCreate/cudnnDestroy** outside of performance-critical code paths. For multithreaded applications that use the same device from different threads, the

recommended programming model is to create one (or a few, as is convenient) cuDNN handle(s) per thread and use that cuDNN handle for the entire life of the thread.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The initialization succeeded.
`CUDNN_STATUS_NOT_INITIALIZED`	CUDA Runtime API initialization failed.
`CUDNN_STATUS_ALLOC_FAILED`	The resources could not be allocated.

4.4. cudnnDestroy

```
cudnnStatus_t cudnnDestroy(cudnnHandle_t handle)
```

This function releases hardware resources used by the cuDNN library. This function is usually the last call with a particular handle to the cuDNN library. Because **cudnnCreate** allocates some internal resources, the release of those resources by calling **cudnnDestroy** will implicitly call **cudaDeviceSynchronize**; therefore, the recommended best practice is to call **cudnnCreate/cudnnDestroy** outside of performance-critical code paths.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The cuDNN context destruction was successful.
`CUDNN_STATUS_NOT_INITIALIZED`	The library was not initialized.

4.5. cudnnSetStream

```
cudnnStatus_t cudnnSetStream(cudnnHandle_t handle, cudaStream_t streamId)
```

This function sets the cuDNN library stream, which will be used to execute all subsequent calls to the cuDNN library functions with that particular handle. If the cuDNN library stream is not set, all kernels use the default (**NULL**) stream. In particular, this routine can be used to change the stream between kernel launches and then to reset the cuDNN library stream back to **NULL**.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The stream was set successfully.

4.6. cudnnGetStream

```
cudnnStatus_t cudnnGetStream(cudnnHandle_t handle, cudaStream_t *streamId)
```

This function gets the cuDNN library stream, which is being used to execute all calls to the cuDNN library functions. If the cuDNN library stream is not set, all kernels use the *default* **NULL** stream.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The stream was returned successfully.

4.7. cudnnCreateTensorDescriptor

```
cudnnStatus_t cudnnCreateTensorDescriptor(cudnnTensorDescriptor_t *tensorDesc)
```

This function creates a generic Tensor descriptor object by allocating the memory needed to hold its opaque structure. The data is initialized to be all zero.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was created successfully.
`CUDNN_STATUS_ALLOC_FAILED`	The resources could not be allocated.

4.8. cudnnSetTensor4dDescriptor

```
cudnnStatus_t
cudnnSetTensor4dDescriptor( cudnnTensorDescriptor_t   tensorDesc,
                            cudnnTensorFormat_t format,
                            cudnnDataType_t dataType,
                            int n,
                            int c,
                            int h,
                            int w )
```

This function initializes a previously created generic Tensor descriptor object into a 4D tensor. The strides of the four dimensions are inferred from the format parameter and set in such a way that the data is contiguous in memory with no padding between dimensions.

> The total size of a tensor including the potential padding between dimensions is limited to 2 Giga-elements of type `datatype`.

Param	In/out	Meaning
tensorDesc	input/output	Handle to a previously created tensor descriptor.
format	input	Type of format.
datatype	input	Data type.
n	input	Number of images.
c	input	Number of feature maps per image.
h	input	Height of each feature map.
w	input	Width of each feature map.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the parameters `n,c,h,w` was negative or `format` has an invalid enumerant value or `dataType` has an invalid enumerant value.
`CUDNN_STATUS_NOT_SUPPORTED`	The total size of the tensor descriptor exceeds the maximim limit of 2 Giga-elements.

4.9. cudnnSetTensor4dDescriptorEx

```
cudnnStatus_t
cudnnSetTensor4dDescriptorEx( cudnnTensorDescriptor_t tensorDesc,
                              cudnnDataType_t dataType,
                              int n,
                              int c,
                              int h,
                              int w,
                              int nStride,
                              int cStride,
                              int hStride,
                              int wStride );
```

This function initializes a previously created generic Tensor descriptor object into a 4D tensor, similarly to **cudnnSetTensor4dDescriptor** but with the strides explicitly passed as parameters. This can be used to lay out the 4D tensor in any order or simply to define gaps between dimensions.

> 💬 At present, some cuDNN routines have limited support for strides; Those routines will return CUDNN_STATUS_NOT_SUPPORTED if a Tensor4D object with an unsupported stride is used. **cudnnTransformTensor** can be used to convert the data to a supported layout.

> 💬 The total size of a tensor including the potential padding between dimensions is limited to 2 Giga-elements of type **datatype**.

Param	In/out	Meaning
tensorDesc	input/ output	Handle to a previously created tensor descriptor.
datatype	input	Data type.
n	input	Number of images.
c	input	Number of feature maps per image.
h	input	Height of each feature map.
w	input	Width of each feature map.
nStride	input	Stride between two consecutive images.

Param	In/out	Meaning
cStride	input	Stride between two consecutive feature maps.
hStride	input	Stride between two consecutive rows.
wStride	input	Stride between two consecutive columns.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the parameters `n,c,h,w` or `nStride,cStride,hStride,wStride` is negative or `dataType` has an invalid enumerant value.
`CUDNN_STATUS_NOT_SUPPORTED`	The total size of the tensor descriptor exceeds the maximim limit of 2 Giga-elements.

4.10. cudnnGetTensor4dDescriptor

```
cudnnStatus_t
cudnnGetTensor4dDescriptor( cudnnTensorDescriptor_t tensorDesc,
                            cudnnDataType_t *dataType,
                            int *n,
                            int *c,
                            int *h,
                            int *w,
                            int *nStride,
                            int *cStride,
                            int *hStride,
                            int *wStride )
```

This function queries the parameters of the previouly initialized Tensor4D descriptor object.

Param	In/out	Meaning
tensorDesc	input	Handle to a previously insitialized tensor descriptor.
datatype	output	Data type.
n	output	Number of images.
c	output	Number of feature maps per image.
h	output	Height of each feature map.
w	output	Width of each feature map.
nStride	output	Stride between two consecutive images.
cStride	output	Stride between two consecutive feature maps.
hStride	output	Stride between two consecutive rows.
wStride	output	Stride between two consecutive columns.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The operation succeeded.

4.11. cudnnSetTensorNdDescriptor

```
cudnnStatus_t
cudnnSetTensorNdDescriptor( cudnnTensorDescriptor_t   tensorDesc,
                            cudnnDataType_t dataType,
                            int nbDims,
                            int dimA[],
                            int strideA[])
```

This function initializes a previously created generic Tensor descriptor object.

> The total size of a tensor including the potential padding between dimensions is limited to 2 Giga-elements of type `datatype`.

Param	In/out	Meaning
tensorDesc	input/output	Handle to a previously created tensor descriptor.
datatype	input	Data type.
nbDims	input	Dimension of the tensor.
dimA	input	Array of dimension `nbDims` that contain the size of the tensor for every dimension.
strideA	input	Array of dimension `nbDims` that contain the stride of the tensor for every dimension.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the elements of the array `dimA` was negative or zero, or `dataType` has an invalid enumerant value.
`CUDNN_STATUS_NOT_SUPPORTED`	the parameter `nbDims` exceeds CUDNN_DIM_MAX or the total size of the tensor descriptor exceeds the maximim limit of 2 Giga-elements.

4.12. cudnnGetTensorNdDescriptor

```
cudnnStatus_t
cudnnGetTensorNdDescriptor( const cudnnTensorDescriptor_t    tensorDesc,
                            int nbDimsRequested,
                            cudnnDataType_t *dataType,
                            int *nbDims,
                            int dimA[],
                            int strideA[])
```

This function retrieves values stored in a previously initialized Tensor descriptor object.

Param	In/out	Meaning
tensorDesc	input	Handle to a previously initialized tensor descriptor.
nbDimsRequested	input	Number of dimensions to extract from a given tensor descriptor. It is also the minimum size of the arrays **dimA** and **strideA**. If this number is greater than the resulting nbDims[0], only nbDims[0] dimensions will be returned.
datatype	output	Data type.
nbDims	output	Actual number of dimensions of the tensor will be returned in nbDims[0].
dimA	output	Array of dimension of at least **nbDimsRequested** that will be filled with the dimensions from the provided tensor descriptor.
strideA	input	Array of dimension of at least **nbDimsRequested** that will be filled with the strides from the provided tensor descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The results were returned successfully.
CUDNN_STATUS_BAD_PARAM	Either **tensorDesc** or **nbDims** pointer is NULL.

4.13. cudnnDestroyTensorDescriptor

```
cudnnStatus_t cudnnDestroyTensorDescriptor(cudnnTensorDescriptor_t tensorDesc)
```

This function destroys a previously created Tensor descriptor object.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was destroyed successfully.

4.14. cudnnTransformTensor

```
cudnnStatus_t
cudnnTransformTensor( cudnnHandle_t                 handle,
                      const void                   *alpha,
                      const cudnnTensorDescriptor_t  xDesc,
                      const void                   *x,
                      const void                   *beta,
                      const cudnnTensorDescriptor_t  yDesc,
                      void                         *y )
```

This function copies the scaled data from one tensor to another tensor with a different layout. Those descriptors need to have the same dimensions but not necessarily the same strides. The input and output tensors must not overlap in any way (i.e., tensors cannot be transformed in place). This function can be used to convert a tensor with an unsupported format to a supported one.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the source value with prior value in the destination tensor as follows: dstValue = alpha[0]*srcValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to a previously initialized tensor descriptor.
x	input	Pointer to data of the tensor described by the **xDesc** descriptor.
yDesc	input	Handle to a previously initialized tensor descriptor.
y	output	Pointer to data of the tensor described by the **yDesc** descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The function launched successfully.
CUDNN_STATUS_BAD_PARAM	The dimensions **n,c,h,w** or the **dataType** of the two tensor descriptors are different.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.15. cudnnAddTensor

```
cudnnStatus_t
cudnnAddTensor_(    cudnnHandle_t                 handle,
                    const void                   *alpha,
                    const cudnnTensorDescriptor_t  aDesc,
                    const void                   *A,
                    const void                   *beta,
                    const cudnnTensorDescriptor_t  cDesc,
                    void                         *C )
```

This function adds the scaled values of a bias tensor to another tensor. Each dimension of the bias tensor `A` must match the corresponding dimension of the destination tensor `C` or must be equal to 1. In the latter case, the same value from the bias tensor for those dimensions will be used to blend into the `C` tensor.

> Up to dimension 5, all tensor formats are supported. Beyond those dimensions, this routine is not supported

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the source value with prior value in the destination tensor as follows: dstValue = alpha[0]*srcValue + beta[0]*priorDstValue. Please refer to this section for additional details.
aDesc	input	Handle to a previously initialized tensor descriptor.
A	input	Pointer to data of the tensor described by the `aDesc` descriptor.
cDesc	input	Handle to a previously initialized tensor descriptor.
C	input/ output	Pointer to data of the tensor described by the `cDesc` descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function executed successfully.
`CUDNN_STATUS_NOT_SUPPORTED`	The dimensions of the bias tensor and the output tensor dimensions are above 5.
`CUDNN_STATUS_BAD_PARAM`	The dimensions of the bias tensor refer to an amount of data that is incompatible the output tensor dimensions or the `dataType` of the two tensor descriptors are different.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.16. cudnnOpTensor

```
cudnnStatus_t
cudnnOpTensor(      cudnnHandle_t                    handle,
               const cudnnOpTensorDescriptor_t   opTensorDesc,
               const void                       *alpha1,
               const cudnnTensorDescriptor_t     aDesc,
               const void                       *A,
               const void                       *alpha2,
               const cudnnTensorDescriptor_t     bDesc,
               const void                       *B,
               const void                       *beta,
               const cudnnTensorDescriptor_t     cDesc,
               void                             *C )
```

This function implements the equation C = op (alpha1[0] * A, alpha2[0] * B) + beta[0] * C, given tensors A, B, and C and scaling factors alpha1, alpha2, and beta. The op to use is indicated by the descriptor **opTensorDesc**. Currently-supported ops are listed by the **cudnnOpTensorOp_t** enum.

Each dimension of the input tensor **A** must match the corresponding dimension of the destination tensor **C**, and each dimension of the input tensor **B** must match the corresponding dimension of the destination tensor **C** or must be equal to 1. In the latter case, the same value from the input tensor **B** for those dimensions will be used to blend into the **C** tensor.

The data types of the input tensors **A** and **B** must match. If the data type of the destination tensor **C** is double, then the data type of the input tensors also must be double.

If the data type of the destination tensor **C** is double, then **opTensorCompType** in **opTensorDesc** must be double. Else **opTensorCompType** must be float.

If the input tensor **B** is the same tensor as the destination tensor **C**, then the input tensor **A** also must be the same tensor as the destination tensor **C**.

> Up to dimension 5, all tensor formats are supported. Beyond those dimensions, this routine is not supported

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
opTensorDesc	input	Handle to a previously initialized op tensor descriptor.
alpha1, alpha2, beta	input	Pointers to scaling factors (in host memory) used to blend the source value with prior value in the destination tensor as indicated by the above op equation. Please refer to this section for additional details.
aDesc, bDesc, cDesc	input	Handle to a previously initialized tensor descriptor.
A, B	input	Pointer to data of the tensors described by the **aDesc** and **bDesc** descriptors, respectively.
C	input/ output	Pointer to data of the tensor described by the **cDesc** descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The function executed successfully.
CUDNN_STATUS_NOT_SUPPORTED	The dimensions of the bias tensor and the output tensor dimensions are above 5, or **opTensorCompType** is not set as stated above.
CUDNN_STATUS_BAD_PARAM	The data type of the destination tensor **C** is unrecognized or the conditions in the above paragraphs are unmet.

Return Value	Meaning
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.17. cudnnSetTensor

```
cudnnStatus_t cudnnSetTensor(
                cudnnHandle_t                    handle,
                const cudnnTensorDescriptor_t yDesc,
                void                             *y,
                const void                       *valuePtr );
```

This function sets all the elements of a tensor to a given value.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
yDesc	input	Handle to a previously initialized tensor descriptor.
y	input/output	Pointer to data of the tensor described by the `yDesc` descriptor.
valuePtr	input	Pointer in Host memory to a single value. All elements of the y tensor will be set to value[0]. The data type of the element in value[0] has to match the data type of tensor `y`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	one of the provided pointers is nil
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.18. cudnnScaleTensor

```
cudnnStatus_t cudnnScaleTensor( cudnnHandle_t                    handle,
                const cudnnTensorDescriptor_t yDesc,
                void                             *y,
                const void                       *alpha);
```

This function scale all the elements of a tensor by a given factor.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
yDesc	input	Handle to a previously initialized tensor descriptor.
y	input/ output	Pointer to data of the tensor described by the `yDesc` descriptor.
alpha	input	Pointer in Host memory to a single value that all elements of the tensor will be scaled with. Please refer to this section for additional details.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	one of the provided pointers is nil
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.19. cudnnCreateFilterDescriptor

`cudnnStatus_t cudnnCreateFilterDescriptor(cudnnFilterDescriptor_t *filterDesc)`

This function creates a filter descriptor object by allocating the memory needed to hold its opaque structure,

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was created successfully.
`CUDNN_STATUS_ALLOC_FAILED`	The resources could not be allocated.

4.20. cudnnSetFilter4dDescriptor

```
cudnnStatus_t
cudnnSetFilter4dDescriptor( cudnnFilterDescriptor_t filterDesc,
                            cudnnDataType_t dataType,
                            cudnnTensorFormat_t  format,
                            int k,
                            int c,
                            int h,
                            int w )
```

This function initializes a previously created filter descriptor object into a 4D filter. Filters layout must be contiguous in memory.

Tensor format CUDNN_TENSOR_NHWC has limited support in **cudnnConvolutionForward**, **cudnnConvolutionBackwardData** and **cudnnConvolutionBackwardFilter**; please refer to each function's documentation for more information.

Param	In/out	Meaning
filterDesc	input/ output	Handle to a previously created filter descriptor.
datatype	input	Data type.
format	input	Type of format.
k	input	Number of output feature maps.
c	input	Number of input feature maps.
h	input	Height of each filter.

Param	In/out	Meaning
w	input	Width of each filter.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the parameters `k,c,h,w` is negative or `dataType` or `format` has an invalid enumerant value.

4.21. cudnnGetFilter4dDescriptor

```
cudnnStatus_t
cudnnGetFilter4dDescriptor( cudnnFilterDescriptor_t filterDesc,
                            cudnnDataType_t *dataType,
                            cudnnTensorFormat_t  *format,
                            int *k,
                            int *c,
                            int *h,
                            int *w )
```

This function queries the parameters of the previouly initialized filter descriptor object.

Param	In/out	Meaning
filterDesc	input	Handle to a previously created filter descriptor.
datatype	output	Data type.
format	output	Type of format.
k	output	Number of output feature maps.
c	output	Number of input feature maps.
h	output	Height of each filter.
w	output	Width of each filter.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.

4.22. cudnnSetFilter4dDescriptor_v3

```
cudnnStatus_t
cudnnSetFilter4dDescriptor_v3( cudnnFilterDescriptor_t filterDesc,
                               cudnnDataType_t dataType,
                               int k,
                               int c,
                               int h,
                               int w )
```

This function initializes a previously created filter descriptor object into a 4D filter. Filters layout must be contiguous in memory. When using this routine to set up a filter descriptor, the filter format is set to CUDNN_TENSOR_NCHW.

> This routine is deprecated, **cudnnSetFilter4dDescriptor** should be used instead.

Param	In/out	Meaning
filterDesc	input/ output	Handle to a previously created filter descriptor.
datatype	input	Data type.
k	input	Number of output feature maps.
c	input	Number of input feature maps.
h	input	Height of each filter.
w	input	Width of each filter.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was set successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the parameters **k,c,h,w** is negative or **dataType** has an invalid enumerant value.

4.23. cudnnGetFilter4dDescriptor_v3

```
cudnnStatus_t
cudnnGetFilter4dDescriptor_v3( cudnnFilterDescriptor_t filterDesc,
                               cudnnDataType_t *dataType,
                               int *k,
                               int *c,
                               int *h,
                               int *w )
```

This function queries the parameters of the previouly initialized filter descriptor object.

> 💬 This routine is deprecated, `cudnnGetFilter4dDescriptor` should be used instead.

Param	In/out	Meaning
filterDesc	input	Handle to a previously created filter descriptor.
datatype	output	Data type.
k	output	Number of output feature maps.
c	output	Number of input feature maps.
h	output	Height of each filter.
w	output	Width of each filter.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.

4.24. cudnnSetFilter4dDescriptor_v4

```
cudnnStatus_t
cudnnSetFilter4dDescriptor_v4( cudnnFilterDescriptor_t filterDesc,
                               cudnnDataType_t dataType,
                               cudnnTensorFormat_t  format,
                               int k,
                               int c,
                               int h,
                               int w )
```

This function is equivalent to **cudnnSetFilter4dDescriptor**.

4.25. cudnnGetFilter4dDescriptor_v4

```
cudnnStatus_t
cudnnGetFilter4dDescriptor_v4( cudnnFilterDescriptor_t filterDesc,
                               cudnnDataType_t *dataType,
                               cudnnTensorFormat_t  *format,
                               int *k,
                               int *c,
                               int *h,
                               int *w )
```

This function is equivalent to **cudnnGetFilter4dDescriptor**.

4.26. cudnnSetFilterNdDescriptor

```
cudnnStatus_t
cudnnSetFilterNdDescriptor( cudnnFilterDescriptor_t filterDesc,
                            cudnnDataType_t  dataType,
                            cudnnTensorFormat_t  format,
                            int nbDims,
                            int filterDimA[])
```

This function initializes a previously created filter descriptor object. Filters layout must be contiguous in memory.

Tensor format CUDNN_TENSOR_NHWC has limited support in **cudnnConvolutionForward**, **cudnnConvolutionBackwardData** and **cudnnConvolutionBackwardFilter**; please refer to each function's documentation for more information.

Param	In/out	Meaning
filterDesc	input/ output	Handle to a previously created filter descriptor.
datatype	input	Data type.
format	input	Type of format.
nbDims	input	Dimension of the filter.
filterDimA	input	Array of dimension **nbDims** containing the size of the filter for each dimension.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was set successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the elements of the array **filterDimA** is negative or **dataType** or **format** has an invalid enumerant value.
CUDNN_STATUS_NOT_SUPPORTED	the parameter **nbDims** exceeds CUDNN_DIM_MAX.

4.27. cudnnGetFilterNdDescriptor

```
cudnnStatus_t
cudnnGetFilterNdDescriptor( const cudnnFilterDescriptor_t wDesc,
                            int nbDimsRequested,
                            cudnnDataType_t *dataType,
                            cudnnTensorFormat_t  *format,
                            int *nbDims,
                            int filterDimA[])
```

This function queries a previously initialized filter descriptor object.

Param	In/out	Meaning
wDesc	input	Handle to a previously initialized filter descriptor.
nbDimsRequested	input	Dimension of the expected filter descriptor. It is also the minimum size of the arrays `filterDimA` in order to be able to hold the results
datatype	input	Data type.
format	output	Type of format.
nbDims	input	Actual dimension of the filter.
filterDimA	input	Array of dimension of at least `nbDimsRequested` that will be filled with the filter parameters from the provided filter descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	The parameter `nbDimsRequested` is negative.

4.28. cudnnSetFilterNdDescriptor_v3

```
cudnnStatus_t
cudnnSetFilterNdDescriptor_v3( cudnnFilterDescriptor_t filterDesc,
                               cudnnDataType_t  dataType,
                               int nbDims,
                               int filterDimA[])
```

This function initializes a previously created filter descriptor object. Filters layout must be contiguous in memory. When using this routine to set up a filter descriptor, the filter format is set to CUDNN_TENSOR_NCHW.

> This routine is deprecated, `cudnnSetFilterNdDescriptor` should be used instead.

Param	In/out	Meaning
filterDesc	input/ output	Handle to a previously created filter descriptor.
datatype	input	Data type.
nbDims	input	Dimension of the filter.
filterDimA	input	Array of dimension `nbDims` containing the size of the filter for each dimension.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.

Return Value	Meaning
`CUDNN_STATUS_BAD_PARAM`	At least one of the elements of the array `filterDimA` is negative or `dataType` has an invalid enumerant value.
`CUDNN_STATUS_NOT_SUPPORTED`	the parameter `nbDims` exceeds CUDNN_DIM_MAX.

4.29. cudnnGetFilterNdDescriptor_v3

```
cudnnStatus_t
cudnnGetFilterNdDescriptor_v3( const cudnnFilterDescriptor_t wDesc,
                               int nbDimsRequested,
                               cudnnDataType_t *dataType,
                               int *nbDims,
                               int filterDimA[])
```

This function queries a previously initialized filter descriptor object.

> This routine is deprecated, **cudnnGetFilterNdDescriptor** should be used instead.

Param	In/out	Meaning
wDesc	input	Handle to a previously initialized filter descriptor.
nbDimsRequested	input	Dimension of the expected filter descriptor. It is also the minimum size of the arrays `filterDimA` in order to be able to hold the results
datatype	input	Data type.
nbDims	input	Actual dimension of the filter.
filterDimA	input	Array of dimension of at least `nbDimsRequested` that will be filled with the filter parameters from the provided filter descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	The parameter `nbDimsRequested` is negative.

4.30. cudnnSetFilterNdDescriptor_v4

```
cudnnStatus_t
cudnnSetFilterNdDescriptor_v4( cudnnFilterDescriptor_t filterDesc,
                               cudnnDataType_t  dataType,
                               cudnnTensorFormat_t  format,
                               int nbDims,
                               int filterDimA[])
```

This function is equivalent to **cudnnSetFilterNdDescriptor**.

4.31. cudnnGetFilterNdDescriptor_v4

```
cudnnStatus_t
cudnnGetFilterNdDescriptor_v4( const cudnnFilterDescriptor_t wDesc,
                               int nbDimsRequested,
                               cudnnDataType_t *dataType,
                               cudnnTensorFormat_t  *format,
                               int *nbDims,
                               int filterDimA[])
```

This function is equivalent to **cudnnGetFilterNdDescriptor**.

4.32. cudnnDestroyFilterDescriptor

```
cudnnStatus_t cudnnDestroyFilterDescriptor(cudnnFilterdDescriptor_t filterDesc)
```

This function destroys a previously created Tensor4D descriptor object.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was destroyed successfully.

4.33. cudnnCreateConvolutionDescriptor

```
cudnnStatus_t cudnnCreateConvolutionDescriptor(cudnnConvolutionDescriptor_t
 *convDesc)
```

This function creates a convolution descriptor object by allocating the memory needed to hold its opaque structure,

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was created successfully.
CUDNN_STATUS_ALLOC_FAILED	The resources could not be allocated.

4.34. cudnnSetConvolution2dDescriptor

```
cudnnStatus_t
cudnnSetConvolution2dDescriptor( cudnnConvolutionDescriptor_t convDesc,
                                 int pad_h,
                                 int pad_w,
                                 int u,
                                 int v,
                                 int upscalex,
                                 int upscaley,
                                 cudnnConvolutionMode_t mode )
```

This function initializes a previously created convolution descriptor object into a 2D correlation. This function assumes that the tensor and filter descriptors corresponds to the formard convolution path and checks if their settings are valid. That same convolution descriptor can be reused in the backward path provided it corresponds to the same layer.

Param	In/out	Meaning
convDesc	input/ output	Handle to a previously created convolution descriptor.
pad_h	input	zero-padding height: number of rows of zeros implicitly concatenated onto the top and onto the bottom of input images.
pad_w	input	zero-padding width: number of columns of zeros implicitly concatenated onto the left and onto the right of input images.
u	input	Vertical filter stride.
v	input	Horizontal filter stride.
upscalex	input	Upscale the input in x-direction.
upscaley	input	Upscale the input in y-direction.
mode	input	Selects between `CUDNN_CONVOLUTION` and `CUDNN_CROSS_CORRELATION`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The descriptor `convDesc` is nil. ▸ One of the parameters `pad_h`,`pad_v` is strictly negative. ▸ One of the parameters `u`,`v` is negative. ▸ The parameter `mode` has an invalid enumerant value.
`CUDNN_STATUS_NOT_SUPPORTED`	The parameter `upscalex` or `upscaley` is not 1.

4.35. cudnnGetConvolution2dDescriptor

```
cudnnStatus_t
cudnnGetConvolution2dDescriptor( const cudnnConvolutionDescriptor_t convDesc,
                                 int* pad_h,
                                 int* pad_w,
                                 int* u,
                                 int* v,
                                 int* upscalex,
                                 int* upscaley,
                                 cudnnConvolutionMode_t *mode )
```

This function queries a previously initialized 2D convolution descriptor object.

Param	In/out	Meaning
convDesc	input/ output	Handle to a previously created convolution descriptor.

Param	In/out	Meaning
pad_h	output	zero-padding height: number of rows of zeros implicitly concatenated onto the top and onto the bottom of input images.
pad_w	output	zero-padding width: number of columns of zeros implicitly concatenated onto the left and onto the right of input images.
u	output	Vertical filter stride.
v	output	Horizontal filter stride.
upscalex	output	Upscale the input in x-direction.
upscaley	output	Upscale the input in y-direction.
mode	output	convolution mode.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The operation was successful.
`CUDNN_STATUS_BAD_PARAM`	The parameter `convDesc` is nil.

4.36. cudnnGetConvolution2dForwardOutputDim

```
cudnnStatus_t
cudnnGetConvolution2dForwardOutputDim( const cudnnConvolutionDescriptor_t
 convDesc,
                                       const cudnnTensorDescriptor_t
 inputTensorDesc,

                                       const cudnnFilterDescriptor_t filterDesc,
                                       int *n,
                                       int *c,
                                       int *h,
                                       int *w )
```

This function returns the dimensions of the resulting 4D tensor of a 2D convolution, given the convolution descriptor, the input tensor descriptor and the filter descriptor This function can help to setup the output tensor and allocate the proper amount of memory prior to launch the actual convolution.

Each dimension `h and w` of the output images is computed as followed:

```
outputDim = 1 + (inputDim + 2*pad - filterDim)/convolutionStride;
```

Param	In/out	Meaning
convDesc	input	Handle to a previously created convolution descriptor.
inputTensorDesc	input	Handle to a previously initialized tensor descriptor.
filterDesc	input	Handle to a previously initialized filter descriptor.
n	output	Number of output images.

Param	In/out	Meaning
c	output	Number of output feature maps per image.
h	output	Height of each output feature map.
w	output	Width of each output feature map.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_BAD_PARAM`	One or more of the descriptors has not been created correctly or there is a mismatch between the feature maps of `inputTensorDesc` and `filterDesc`.
`CUDNN_STATUS_SUCCESS`	The object was set successfully.

4.37. cudnnSetConvolutionNdDescriptor

```
cudnnStatus_t
cudnnSetConvolutionNdDescriptor( cudnnConvolutionDescriptor_t convDesc,
                                 int arrayLength,
                                 int padA[],
                                 int filterStrideA[],
                                 int upscaleA[],
                                 cudnnConvolutionMode_t mode,
                                 cudnnDataType_t dataType )
```

This function initializes a previously created generic convolution descriptor object into a n-D correlation. That same convolution descriptor can be reused in the backward path provided it corresponds to the same layer. The convolution computation will done in the specified `dataType`, which can be potentially different from the input/output tensors.

Param	In/out	Meaning
convDesc	input/ output	Handle to a previously created convolution descriptor.
arrayLength	input	Dimension of the convolution.
padA	input	Array of dimension `arrayLength` containing the zero-padding size for each dimension. For every dimension, the padding represents the number of extra zeros implicitly concatenated at the start and at the end of every element of that dimension .
filterStrideA	input	Array of dimension `arrayLength` containing the filter stride for each dimension. For every dimension, the fitler stride represents the number of elements to slide to reach the next start of the filtering window of the next point.
upscaleA	input	Array of dimension `arrayLength` containing the upscale factor for each dimension.
mode	input	Selects between `CUDNN_CONVOLUTION` and `CUDNN_CROSS_CORRELATION`.

Param	In/out	Meaning
datatype	input	Selects the datatype in which the computation will be done.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was set successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The descriptor **convDesc** is nil. ▸ The **arrayLenghtRequest** is negative. ▸ The enumerant **mode** has an invalid value. ▸ The enumerant **datatype** has an invalid value. ▸ One of the elements of **padA** is strictly negative. ▸ One of the elements of **strideA** is negative or zero.
CUDNN_STATUS_NOT_SUPPORTED	At least one of the following conditions are met: ▸ The **arrayLenghtRequest** is greater than CUDNN_DIM_MAX. ▸ The array **upscaleA** contains an element different from 1.

4.38. cudnnGetConvolutionNdDescriptor

```
cudnnStatus_t
cudnnGetConvolutionNdDescriptor( const cudnnConvolutionDescriptor_t convDesc,
                                 int arrayLengthRequested,
                                 int *arrayLength,
                                 int padA[],
                                 int filterStrideA[],
                                 int upscaleA[],
                                 cudnnConvolutionMode_t *mode,
                                 cudnnDataType_t *dataType )
```

This function queries a previously initialized convolution descriptor object.

Param	In/out	Meaning
convDesc	input/output	Handle to a previously created convolution descriptor.
arrayLengthRequested	input	Dimension of the expected convolution descriptor. It is also the minimum size of the arrays **padA**, **filterStrideA** and **upscaleA** in order to be able to hold the results
arrayLength	output	actual dimension of the convolution descriptor.
padA	output	Array of dimension of at least **arrayLengthRequested** that will be filled with the padding parameters from the provided convolution descriptor.

Param	In/out	Meaning
filterStrideA	output	Array of dimension of at least `arrayLengthRequested` that will be filled with the filter stride from the provided convolution descriptor.
upscaleA	output	Array of dimension at least `arrayLengthRequested` that will be filled with the upscaling parameters from the provided convolution descriptor.
mode	output	convolution mode of the provided descriptor.
datatype	output	datatype of the provided descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The descriptor `convDesc` is nil. ▸ The `arrayLenghtRequest` is negative.
`CUDNN_STATUS_NOT_SUPPORTED`	The `arrayLenghtRequest` is greater than CUDNN_DIM_MAX

4.39. cudnnGetConvolutionNdForwardOutputDim

```
cudnnStatus_t
cudnnGetConvolutionNdForwardOutputDim( const cudnnConvolutionDescriptor_t
 convDesc,
                                        const cudnnTensorDescriptor_t
 inputTensorDesc,

                                        const cudnnFilterDescriptor_t filterDesc,
                                        int nbDims,
                                        int tensorOuputDimA[] )
```

This function returns the dimensions of the resulting n-D tensor of a `nbDims-2`-D convolution, given the convolution descriptor, the input tensor descriptor and the filter descriptor This function can help to setup the output tensor and allocate the proper amount of memory prior to launch the actual convolution.

Each dimension of the `(nbDims-2)-D` images of the output tensor is computed as followed:

```
outputDim = 1 + (inputDim + 2*pad - filterDim)/convolutionStride;
```

Param	In/out	Meaning
convDesc	input	Handle to a previously created convolution descriptor.
inputTensorDesc	input	Handle to a previously initialized tensor descriptor.
filterDesc	input	Handle to a previously initialized filter descriptor.
nbDims	input	Dimension of the output tensor

Param	In/out	Meaning
tensorOuputDimA	Output	Array of dimensions `nbDims` that contains on exit of this routine the sizes of the output tensor

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the parameters `convDesc`, `inputTensorDesc`, and `filterDesc`, is nil ▸ The dimension of the filter descriptor `filterDesc` is different from the dimension of input tensor descriptor `inputTensorDesc`. ▸ The dimension of the convolution descriptor is different from the dimension of input tensor descriptor `inputTensorDesc` -2 . ▸ The features map of the filter descriptor `filterDesc` is different from the one of input tensor descriptor `inputTensorDesc`. ▸ The size of the filter `filterDesc` is larger than the padded sizes of the input tensor. ▸ The dimension `nbDims` of the output array is negative or greater than the dimension of input tensor descriptor `inputTensorDesc`.
`CUDNN_STATUS_SUCCESS`	The routine exits successfully.

4.40. cudnnDestroyConvolutionDescriptor

```
cudnnStatus_t cudnnDestroyConvolutionDescriptor(cudnnConvolutionDescriptor_t
 convDesc)
```

This function destroys a previously created convolution descriptor object.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was destroyed successfully.

4.41. cudnnFindConvolutionForwardAlgorithm

```
cudnnStatus_t
cudnnFindConvolutionForwardAlgorithm( cudnnHandle_t                      handle,
                                      const cudnnTensorDescriptor_t      xDesc,
                                      const cudnnFilterDescriptor_t      wDesc,
                                      const cudnnConvolutionDescriptor_t
convDesc,
                                      const cudnnTensorDescriptor_t      yDesc,
                                      const int
requestedAlgoCount,
                                      int
*returnedAlgoCount,
                                      cudnnConvolutionFwdAlgoPerf_t
*perfResults
                    )
```

This function attempts all cuDNN algorithms for **cudnnConvolutionForward()**, using memory allocated via **cudaMalloc()**, and outputs performance metrics to a user-allocated array of **cudnnConvolutionFwdAlgoPerf_t**. These metrics are written in sorted fashion where the first element has the lowest compute time.

> 💬 This function is host blocking.

> 💬 It is recommend to run this function prior to allocating layer data; doing otherwise may needlessly inhibit some algorithm options due to resource usage.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized input tensor descriptor.
wDesc	input	Handle to a previously initialized filter descriptor.
convDesc	input	Previously initialized convolution descriptor.
yDesc	input	Handle to the previously initialized output tensor descriptor.
requestedAlgoCount	input	The maximum number of elements to be stored in perfResults.
returnedAlgoCount	output	The number of output elements stored in perfResults.
perfResults	output	A user-allocated array to store performance metrics sorted ascending by compute time.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.

Return Value	Meaning
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `handle` is not allocated properly. ▸ `xDesc`, `wDesc` or `yDesc` is not allocated properly. ▸ `xDesc`, `wDesc` or `yDesc` has fewer than 1 dimension. ▸ Either `returnedCount` or `perfResults` is nil. ▸ `requestedCount` is less than 1.
`CUDNN_STATUS_ALLOC_FAILED`	This function was unable to allocate memory to store sample input, filters and output.
`CUDNN_STATUS_INTERNAL_ERROR`	At least one of the following conditions are met: ▸ The function was unable to allocate neccesary timing objects. ▸ The function was unable to deallocate neccesary timing objects. ▸ The function was unable to deallocate sample input, filters and output.

4.42. cudnnFindConvolutionForwardAlgorithmEx

```
cudnnStatus_t
cudnnFindConvolutionForwardAlgorithmEx( cudnnHandle_t
 handle,
                                        const cudnnTensorDescriptor_t
 xDesc,
                                        const void                    *x,
                                        const cudnnFilterDescriptor_t
 wDesc,
                                        const void                    *w,
                                        const cudnnConvolutionDescriptor_t
 convDesc,
                                        const cudnnTensorDescriptor_t
 yDesc,
                                        void                          *y,
                                        const int
 requestedAlgoCount,
                                        int
 *returnedAlgoCount,
                                        cudnnConvolutionFwdAlgoPerf_t
 *perfResults,
                                        void
 *workSpace,
                                        size_t
 workSpaceSizeInBytes
                        )
```

This function attempts all available cuDNN algorithms for **cudnnConvolutionForward**, using user-allocated GPU memory, and outputs performance metrics to a user-allocated array of **cudnnConvolutionFwdAlgoPerf_t**.

These metrics are written in sorted fashion where the first element has the lowest compute time.

> 💬 This function is host blocking.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor `xDesc`.
wDesc	input	Handle to a previously initialized filter descriptor.
w	input	Data pointer to GPU memory associated with the filter descriptor `wDesc`.
convDesc	input	Previously initialized convolution descriptor.
yDesc	input	Handle to the previously initialized output tensor descriptor.
y	input/ output	Data pointer to GPU memory associated with the tensor descriptor `yDesc`. The content of this tensor will be overwritten with arbitary values.
requestedAlgoCount	input	The maximum number of elements to be stored in perfResults.
returnedAlgoCount	output	The number of output elements stored in perfResults.
perfResults	output	A user-allocated array to store performance metrics sorted ascending by compute time.
workSpace	input	Data pointer to GPU memory that is a necessary workspace for some algorithms. The size of this workspace will determine the availability of algorithms. A nil pointer is considered a workSpace of 0 bytes.
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workSpace`

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `handle` is not allocated properly. ▸ `xDesc, wDesc` or `yDesc` is not allocated properly. ▸ `xDesc, wDesc` or `yDesc` has fewer than 1 dimension.

Return Value	Meaning
	▸ **x**, **w** or **y** is nil.
	▸ Either **returnedCount** or **perfResults** is nil.
	▸ **requestedCount** is less than 1.
CUDNN_STATUS_INTERNAL_ERROR	At least one of the following conditions are met: ▸ The function was unable to allocate neccesary timing objects. ▸ The function was unable to deallocate neccesay timing objects. ▸ The function was unable to deallocate sample input, filters and output.

4.43. cudnnGetConvolutionForwardAlgorithm

```
cudnnStatus_t
cudnnGetConvolutionForwardAlgorithm( cudnnHandle_t                     handle,
                                     const cudnnTensorDescriptor_t     xDesc,
                                     const cudnnFilterDescriptor_t     wDesc,
                                     const cudnnConvolutionDescriptor_t
convDesc,
                                     const cudnnTensorDescriptor_t     yDesc,
                                     cudnnConvolutionFwdPreference_t
preference,
                                     size_t
memoryLimitInbytes,
                                     cudnnConvolutionFwdAlgo_t         *algo
                                   )
```

This function serves as a heuristic for obtaining the best suited algorithm for **cudnnConvolutionForward** for the given layer specifications. Based on the input preference, this function will either return the fastest algorithm or the fastest algorithm within a given memory limit. For an exhaustive search for the fastest algorithm, please use **cudnnFindConvolutionForwardAlgorithm**.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized input tensor descriptor.
wDesc	input	Handle to a previously initialized convolution filter descriptor.
convDesc	input	Previously initialized convolution descriptor.
yDesc	input	Handle to the previously initialized output tensor descriptor.
preference	input	Enumerant to express the preference criteria in terms of memory requirement and speed.
memoryLimitInBytes	input	It is used when enumerant **preference** is set to **CUDNN_CONVOLUTION_FWD_SPECIFY_WORKSPACE_LIMIT** to specify the maximum amount of GPU memory the user is willing to use as a workspace

Param	In/out	Meaning
algo	output	Enumerant that specifies which convolution algorithm should be used to compute the results according to the specified preference

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the parameters handle, xDesc, wDesc, convDesc, yDesc is NULL. ▸ Either yDesc or wDesc have different dimensions from xDesc. ▸ The data types of tensors xDesc, yDesc or wDesc are not all the same. ▸ The number of feature maps in xDesc and wDesc differs. ▸ The tensor xDesc has a dimension smaller than 3.

4.44. cudnnGetConvolutionForwardWorkspaceSize

```
cudnnStatus_t
cudnnGetConvolutionForwardWorkspaceSize( cudnnHandle_t    handle,
                                         const    cudnnTensorDescriptor_t
 xDesc,
                                         const    cudnnFilterDescriptor_t
 wDesc,
                                         const    cudnnConvolutionDescriptor_t
 convDesc,
                                         const    cudnnTensor4dDescriptor_t
 yDesc,
                                         cudnnConvolutionFwdAlgo_t
 algo,
                                         size_t
 *sizeInBytes
                             )
```

This function returns the amount of GPU memory workspace the user needs to allocate to be able to call **cudnnConvolutionForward** with the specified algorithm. The workspace allocated will then be passed to the routine **cudnnConvolutionForward**. The specified algorithm can be the result of the call to **cudnnGetConvolutionForwardAlgorithm** or can be chosen arbitrarily by the user. Note that not every algorithm is available for every configuration of the input tensor and/or every configuration of the convolution descriptor.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized x tensor descriptor.

Param	In/out	Meaning
wDesc	input	Handle to a previously initialized filter descriptor.
convDesc	input	Previously initialized convolution descriptor.
yDesc	input	Handle to the previously initialized y tensor descriptor.
algo	input	Enumerant that specifies the chosen convolution algorithm
sizeInBytes	output	Amount of GPU memory needed as workspace to be able to execute a forward convolution with the specified **algo**

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the parameters handle, xDesc, wDesc, convDesc, yDesc is NULL. ▸ The tensor yDesc or wDesc are not of the same dimension as xDesc. ▸ The tensor xDesc, yDesc or wDesc are not of the same data type. ▸ The numbers of feature maps of the tensor xDesc and wDesc differ. ▸ The tensor xDesc has a dimension smaller than 3.
`CUDNN_STATUS_NOT_SUPPORTED`	The combination of the tensor descriptors, filter descriptor and convolution descriptor is not supported for the specified algorithm.

4.45. cudnnConvolutionForward

```
cudnnStatus_t
cudnnConvolutionForward( cudnnHandle_t                      handle,
                         const void                        *alpha,
                         const cudnnTensorDescriptor_t      xDesc,
                         const void                        *x,
                         const cudnnFilterDescriptor_t      wDesc,
                         const void                        *w,
                         const cudnnConvolutionDescriptor_t  convDesc,
                         cudnnConvolutionFwdAlgo_t          algo,
                         void                              *workSpace,
                         size_t
workSpaceSizeInBytes,
                         const void                        *beta,
                         const cudnnTensorDescriptor_t      yDesc,
                         void                              *y )
```

This function executes convolutions or cross-correlations over **x** using filters specified with **w**, returning results in **y**. Scaling factors **alpha** and **beta** can be used to scale the input tensor and the output tensor respectively.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to a previously initialized tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor `xDesc`.
wDesc	input	Handle to a previously initialized filter descriptor.
w	input	Data pointer to GPU memory associated with the filter descriptor `wDesc`.
convDesc	input	Previously initialized convolution descriptor.
algo	input	Enumerant that specifies which convolution algorithm shoud be used to compute the results
workSpace	input	Data pointer to GPU memory to a workspace needed to able to execute the specified algorithm. If no workspace is needed for a particular algorithm, that pointer can be nil
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workSpace`
yDesc	input	Handle to a previously initialized tensor descriptor.
y	input/ output	Data pointer to GPU memory associated with the tensor descriptor `yDesc` that carries the result of the convolution.

This function supports only four specific combinations of data types for `xDesc`, `wDesc`, `convDesc` and `yDesc`. See the following for an exhaustive list of these configurations.

Data Type Configurations	`xDesc`'s, `wDesc`'s and `yDesc`'s Data Type	`convDesc`'s Data Type
TRUE_HALF_CONFIG	CUDNN_DATA_HALF	CUDNN_DATA_HALF
PSEUDO_HALF_CONFIG	CUDNN_DATA_HALF	CUDNN_DATA_FLOAT
FLOAT_CONFIG	CUDNN_DATA_FLOAT	CUDNN_DATA_FLOAT
DOUBLE_CONFIG	CUDNN_DATA_DOUBLE	CUDNN_DATA_DOUBLE

> TRUE_HALF_CONFIG is only supported on architectures with true fp16 support (compute capability 5.3 and 6.0).

For this function, all algorithms perform deterministic computations. Specifying a separate algorithm can cause changes in performance and support. See the following for an exhaustive list of algorithm options and their respective supported parameters.

`wDesc` may only have format CUDNN_TENSOR_NHWC when all of the following are true:

▸ `algo` is CUDNN_CONVOLUTION_FWD_ALGO_IMPLICIT_GEMM
▸ `xDesc` and `yDesc` is NHWC HWC-packed

▶ Data type configuration is PSEUDO_HALF_CONFIG or FLOAT_CONFIG
▶ The convolution is 2-dimensional

The following is an exhaustive list of algo support for 2-d convolutions.

▶ **CUDNN_CONVOLUTION_FWD_ALGO_IMPLICIT_GEMM**

  ▶ **xDesc** Format Support: All
  ▶ **yDesc** Format Support: All
  ▶ Data Type Config Support: All except TRUE_HALF_CONFIG

▶ **CUDNN_CONVOLUTION_FWD_ALGO_IMPLICIT_PRECOMPUTED_GEMM**

  ▶ **xDesc** Format Support: All
  ▶ **yDesc** Format Support: All
  ▶ Data Type Config Support: All

▶ **CUDNN_CONVOLUTION_FWD_ALGO_GEMM**

  ▶ **xDesc** Format Support: All
  ▶ **yDesc** Format Support: All
  ▶ Data Type Config Support: All except TRUE_HALF_CONFIG

▶ **CUDNN_CONVOLUTION_FWD_ALGO_DIRECT**

  ▶ This algorithm has no current implementation in cuDNN.

▶ **CUDNN_CONVOLUTION_FWD_ALGO_FFT**

  ▶ **xDesc** Format Support: NCHW HW-packed
  ▶ **yDesc** Format Support: NCHW HW-packed
  ▶ Data Type Config Support: PSEUDO_HALF_CONFIG, FLOAT_CONFIG
  ▶ Notes:

    ▶ **xDesc**'s feature map height + 2 * **convDesc**'s zero-padding height must equal 256 or less
    ▶ **xDesc**'s feature map width + 2 * **convDesc**'s zero-padding width must equal 256 or less
    ▶ **convDesc**'s vertical and horizontal filter stride must equal 1
    ▶ **wDesc**'s filter height must be greater than **convDesc**'s zero-padding height
    ▶ **wDesc**'s filter width must be greater than **convDesc**'s zero-padding width

▶ **CUDNN_CONVOLUTION_FWD_ALGO_FFT_TILING**

  ▶ **xDesc** Format Support: NCHW HW-packed
  ▶ **yDesc** Format Support: NCHW HW-packed
  ▶ Data Type Config Support: PSEUDO_HALF_CONFIG, FLOAT_CONFIG
  ▶ Notes:

    ▶ **wDesc**'s filter height must equal 32 or less
    ▶ **wDesc**'s filter width must equal 32 or less
    ▶ **convDesc**'s vertical and horizontal filter stride must equal 1
    ▶ **wDesc**'s filter height must be greater than **convDesc**'s zero-padding height
    ▶ **wDesc**'s filter width must be greater than **convDesc**'s zero-padding width

▶ **CUDNN_CONVOLUTION_FWD_ALGO_WINOGRAD**

- ▶ **xDesc** Format Support: All
- ▶ **yDesc** Format Support: All
- ▶ Data Type Config Support: PSEUDO_HALF_CONFIG, FLOAT_CONFIG
- ▶ Notes:

  - ▶ **convDesc**'s vertical and horizontal filter stride must equal 1
  - ▶ **wDesc**'s filter height must be 3
  - ▶ **wDesc**'s filter width must be 3

The following is an exhaustive list of algo support for 3-d convolutions.

- ▶ **CUDNN_CONVOLUTION_FWD_ALGO_IMPLICIT_GEMM**

  - ▶ **xDesc** Format Support: All
  - ▶ **yDesc** Format Support: All
  - ▶ Data Type Config Support: All except TRUE_HALF_CONFIG
- ▶ **CUDNN_CONVOLUTION_FWD_ALGO_IMPLICIT_PRECOMPUTED_GEMM**

  - ▶ **xDesc** Format Support: NCDHW-fully-packed
  - ▶ **yDesc** Format Support: NCDHW-fully-packed
  - ▶ Data Type Config Support: All except TRUE_HALF_CONFIG
- ▶ **CUDNN_CONVOLUTION_FWD_ALGO_FFT_TILING**

  - ▶ **xDesc** Format Support: NCDHW DHW-packed
  - ▶ **yDesc** Format Support: NCDHW DHW-packed
  - ▶ Data Type Config Support: All except TRUE_HALF_CONFIG
  - ▶ Notes:

    - ▶ **wDesc**'s filter height must equal 16 or less
    - ▶ **wDesc**'s filter width must equal 16 or less
    - ▶ **wDesc**'s filter depth must equal 16 or less
    - ▶ **convDesc**'s must have all filter strides equal to 1
    - ▶ **wDesc**'s filter height must be greater than **convDesc**'s zero-padding height
    - ▶ **wDesc**'s filter width must be greater than **convDesc**'s zero-padding width
    - ▶ **wDesc**'s filter depth must be greater than **convDesc**'s zero-padding width

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The operation was launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▶ At least one of the following is NULL: handle, `xDesc`, `wDesc`, `convDesc`, `yDesc`, `xData`, `w`, `yData`, `alpha`, `beta` ▶ `xDesc` and `yDesc` have a non-matching number of dimensions ▶ `xDesc` and `wDesc` have a non-matching number of dimensions ▶ `xDesc` has fewer than three number of dimensions

Return Value	Meaning
	▸ `xDesc`'s number of dimensions is not equal to `convDesc`'s array length + 2 ▸ `xDesc` and `wDesc` have a non-matching number of input feature maps per image ▸ `xDesc`, `wDesc` and `yDesc` have a non-matching data type ▸ For some spatial dimension, `wDesc` has a spatial size that is larger than the input spatial size (including zero-padding size)
`CUDNN_STATUS_NOT_SUPPORTED`	At least one of the following conditions are met: ▸ `xDesc` or `yDesc` have negative tensor striding ▸ `xDesc`, `wDesc` or `yDesc` has a number of dimensions that is not 4 or 5 ▸ The chosen algo does not support the parameters provided; see above for exhaustive list of parameter support for each algo
`CUDNN_STATUS_MAPPING_ERROR`	An error occured during the texture binding of the filter data.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.46. cudnnConvolutionBackwardBias

```
cudnnStatus_t
cudnnConvolutionBackwardBias( cudnnHandle_t                  handle,
                              const void                     *alpha,
                              const cudnnTensorDescriptor_t   dyDesc,
                              const void                     *dy,
                              const void                     *beta,
                              const cudnnTensorDescriptor_t   dbDesc,
                              void                           *db
                             )
```

This function computes the convolution function gradient with respect to the bias, which is the sum of every element belonging to the same feature map across all of the images of the input tensor. Therefore, the number of elements produced is equal to the number of features maps of the input tensor.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
dyDesc	input	Handle to the previously initialized input tensor descriptor.
dy	input	Data pointer to GPU memory associated with the tensor descriptor `dyDesc`.
dbDesc	input	Handle to the previously initialized output tensor descriptor.

Param	In/out	Meaning
db	output	Data pointer to GPU memory associated with the output tensor descriptor `dbDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The operation was launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the parameters `n`,`height`,`width` of the output tensor is not 1. ▸ The numbers of feature maps of the input tensor and output tensor differ. ▸ The `dataType` of the two tensor descriptors are different.

4.47. cudnnFindConvolutionBackwardFilterAlgorithm

```
cudnnStatus_t
cudnnFindConvolutionBackwardFilterAlgorithm( cudnnHandle_t
    handle,
                                            const cudnnTensorDescriptor_t
    xDesc,
                                            const cudnnTensorDescriptor_t
    dyDesc,
                                            const cudnnConvolutionDescriptor_t
    convDesc,
                                            const cudnnFilterDescriptor_t
    dwDesc,
                                            const int
     requestedAlgoCount,
                                            int
    *returnedAlgoCount,
                                            cudnnConvolutionBwdFilterAlgoPerf_t
    *perfResults
                                            )
```

This function attempts all cuDNN algorithms for
`cudnnConvolutionBackwardFilter()`, using GPU memory allocated via
`cudaMalloc()`, and outputs performance metrics to a user-allocated array of
`cudnnConvolutionBwdFilterAlgoPerf_t`. These metrics are written in sorted
fashion where the first element has the lowest compute time.

💬 This function is host blocking.

💬 It is recommend to run this function prior to allocating layer data; doing otherwise
may needlessly inhibit some algorithm options due to resource usage.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized input tensor descriptor.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
convDesc	input	Previously initialized convolution descriptor.
dwDesc	input	Handle to a previously initialized filter descriptor.
requestedAlgoCount	input	The maximum number of elements to be stored in perfResults.
returnedAlgoCount	output	The number of output elements stored in perfResults.
perfResults	output	A user-allocated array to store performance metrics sorted ascending by compute time.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `handle` is not allocated properly. ▸ `xDesc`, `dyDesc` or `dwDesc` is not allocated properly. ▸ `xDesc`, `dyDesc` or `dwDesc` has fewer than 1 dimension. ▸ Either `returnedCount` or `perfResults` is nil. ▸ `requestedCount` is less than 1.
`CUDNN_STATUS_ALLOC_FAILED`	This function was unable to allocate memory to store sample input, filters and output.
`CUDNN_STATUS_INTERNAL_ERROR`	At least one of the following conditions are met: ▸ The function was unable to allocate neccesary timing objects. ▸ The function was unable to deallocate neccesary timing objects. ▸ The function was unable to deallocate sample input, filters and output.

4.48. cudnnFindConvolutionBackwardFilterAlgorithmEx

```
cudnnStatus_t
cudnnFindConvolutionBackwardFilterAlgorithmEx( cudnnHandle_t
      handle,
                                               const cudnnTensorDescriptor_t
      xDesc,
                                               const void
      *x,
                                               const cudnnTensorDescriptor_t
      dyDesc,
                                               const void
      *dy,
                                               const
 cudnnConvolutionDescriptor_t     convDesc,
                                               const cudnnFilterDescriptor_t
      dwDesc,
                                               void
      *dw,
                                               const int
      requestedAlgoCount,
                                               int
      *returnedAlgoCount,
 cudnnConvolutionBwdFilterAlgoPerf_t    *perfResults,
                                               void
      *workSpace,
                                               size_t
      workSpaceSizeInBytes
                                             )
```

This function attempts all cuDNN algorithms for **cudnnConvolutionBackwardFilter**,
using user-allocated GPU memory, and outputs performance metrics to a user-allocated
array of **cudnnConvolutionBwdFilterAlgoPerf_t**. These metrics are written in
sorted fashion where the first element has the lowest compute time.

> This function is host blocking.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the filter descriptor **xDesc**.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the tensor descriptor **dyDesc**.
convDesc	input	Previously initialized convolution descriptor.
dwDesc	input	Handle to a previously initialized filter descriptor.
dw	input/ output	Data pointer to GPU memory associated with the filter descriptor **dwDesc**. The content of this tensor will be overwritten with arbitary values.

Param	In/out	Meaning
requestedAlgoCount	input	The maximum number of elements to be stored in perfResults.
returnedAlgoCount	output	The number of output elements stored in perfResults.
perfResults	output	A user-allocated array to store performance metrics sorted ascending by compute time.
workSpace	input	Data pointer to GPU memory that is a necessary workspace for some algorithms. The size of this workspace will determine the availabilty of algorithms. A nil pointer is considered a workSpace of 0 bytes.
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workSpace`

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `handle` is not allocated properly. ▸ `xDesc`, `dyDesc` or `dwDesc` is not allocated properly. ▸ `xDesc`, `dyDesc` or `dwDesc` has fewer than 1 dimension. ▸ `x`, `dy` or `dw` is nil. ▸ Either `returnedCount` or `perfResults` is nil. ▸ `requestedCount` is less than 1.
`CUDNN_STATUS_INTERNAL_ERROR`	At least one of the following conditions are met: ▸ The function was unable to allocate neccesary timing objects. ▸ The function was unable to deallocate neccesary timing objects. ▸ The function was unable to deallocate sample input, filters and output.

4.49. cudnnGetConvolutionBackwardFilterAlgorithm

```
cudnnStatus_t
cudnnGetConvolutionBackwardFilterAlgorithm( cudnnHandle_t
    handle,
                                            const cudnnTensorDescriptor_t
    xDesc,
                                            const cudnnTensorDescriptor_t
    dyDesc,
                                            const cudnnConvolutionDescriptor_t
    convDesc,
                                            const cudnnFilterDescriptor_t
    dwDesc,
 cudnnConvolutionBwdFilterPreference_t  preference,
                                            size_t
    memoryLimitInbytes,
                                            cudnnConvolutionBwdFilterAlgo_t
    *algo
                                )
```

This function serves as a heuristic for obtaining the best suited algorithm for **cudnnConvolutionBackwardFilter** for the given layer specifications. Based on the input preference, this function will either return the fastest algorithm or the fastest algorithm within a given memory limit. For an exhaustive search for the fastest algorithm, please use **cudnnFindConvolutionBackwardFilterAlgorithm**.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized input tensor descriptor.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
convDesc	input	Previously initialized convolution descriptor.
dwDesc	input	Handle to a previously initialized filter descriptor.
preference	input	Enumerant to express the preference criteria in terms of memory requirement and speed.
memoryLimitInbytes	input	It is to specify the maximum amount of GPU memory the user is willing to use as a workspace. This is currently a placeholder and is not used.
algo	output	Enumerant that specifies which convolution algorithm should be used to compute the results according to the specified preference

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.
CUDNN_STATUS_BADPARAM	At least one of the following conditions are met: ▸ The numbers of feature maps of the input tensor and output tensor differ. ▸ The **dataType** of the two tensor descriptors or the filter are different.

4.50. cudnnGetConvolutionBackwardFilterWorkspaceSize

```
cudnnStatus_t
cudnnGetConvolutionBackwardFilterWorkspaceSize( cudnnHandle_t    handle,
                                                const cudnnTensorDescriptor_t
    xDesc,
                                                const cudnnTensorDescriptor_t
    dyDesc,
                                                const
 cudnnConvolutionDescriptor_t  convDesc,
                                                const cudnnFilterDescriptor_t
    dwDesc,
                                                cudnnConvolutionFwdAlgo_t
    algo,
                                                size_t
  *sizeInBytes
                             )
```

This function returns the amount of GPU memory workspace the user needs to allocate to be able to call **cudnnConvolutionBackwardFilter** with the specified algorithm. The workspace allocated will then be passed to the routine **cudnnConvolutionBackwardFilter**. The specified algorithm can be the result of the call to **cudnnGetConvolutionBackwardFilterAlgorithm** or can be chosen arbitrarily by the user. Note that not every algorithm is available for every configuration of the input tensor and/or every configuration of the convolution descriptor.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
xDesc	input	Handle to the previously initialized input tensor descriptor.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
convDesc	input	Previously initialized convolution descriptor.
dwDesc	input	Handle to a previously initialized filter descriptor.
algo	input	Enumerant that specifies the chosen convolution algorithm
sizeInBytes	output	Amount of GPU memory needed as workspace to be able to execute a forward convolution with the specified **algo**

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The numbers of feature maps of the input tensor and output tensor differ. ▸ The **dataType** of the two tensor descriptors or the filter are different.

Return Value	Meaning
`CUDNN_STATUS_NOT_SUPPORTED`	The combination of the tensor descriptors, filter descriptor and convolution descriptor is not supported for the specified algorithm.

4.51. cudnnConvolutionBackwardFilter

```
cudnnStatus_t
cudnnConvolutionBackwardFilter  ( cudnnHandle_t                          handle,
                                   const void                           *alpha,
                                   const cudnnTensorDescriptor_t         xDesc,
                                   const void                           *x,
                                   const cudnnTensorDescriptor_t         dyDesc,
                                   const void                           *dy,
                                   const cudnnConvolutionDescriptor_t
convDesc,
                                   cudnnConvolutionBwdFilterAlgo_t       algo,
                                   void
*workSpace,
                                   size_t
workSpaceSizeInBytes,
                                   const void                           *beta,
                                   const cudnnFilterDescriptor_t         dwDesc,
                                   void                                 *dw )
```

This function computes the convolution gradient with respect to filter coefficients using the specified `algo`, returning results in `gradDesc`.Scaling factors `alpha` and `beta` can be used to scale the input tensor and the output tensor respectively.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to a previously initialized tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor `xDesc`.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the backpropagation gradient tensor descriptor `dyDesc`.
convDesc	input	Previously initialized convolution descriptor.
algo	input	Enumerant that specifies which convolution algorithm shoud be used to compute the results
workSpace	input	Data pointer to GPU memory to a workspace needed to able to execute the specified algorithm. If no workspace is needed for a particular algorithm, that pointer can be nil
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workSpace`
dwDesc	input	Handle to a previously initialized filter gradient descriptor.

Param	In/out	Meaning
dw	input/ output	Data pointer to GPU memory associated with the filter gradient descriptor `dwDesc` that carries the result.

This function supports only three specific combinations of data types for `xDesc`, `dyDesc`, `convDesc` and `dwDesc`. See the following for an exhaustive list of these configurations.

Data Type Configurations	xDesc's, dyDesc's and dwDesc's Data Type	convDesc's Data Type
PSEUDO_HALF_CONFIG	CUDNN_DATA_HALF	CUDNN_DATA_FLOAT
FLOAT_CONFIG	CUDNN_DATA_FLOAT	CUDNN_DATA_FLOAT
DOUBLE_CONFIG	CUDNN_DATA_DOUBLE	CUDNN_DATA_DOUBLE

Specifying a separate algorithm can cause changes in performance, support and computation determinism. See the following for an exhaustive list of algorithm options and their respective supported parameters and deterministic behavior.

`dwDesc` may only have format CUDNN_TENSOR_NHWC when all of the following are true:

▸ `algo` is CUDNN_CONVOLUTION_BWD_FILTER_ALGO_0 or CUDNN_CONVOLUTION_BWD_FILTER_ALGO_1
▸ `xDesc` and `dyDesc` is NHWC HWC-packed
▸ Data type configuration is PSEUDO_HALF_CONFIG or FLOAT_CONFIG
▸ The convolution is 2-dimensional

The following is an exhaustive list of algo support for 2-d convolutions.

▸ **CUDNN_CONVOLUTION_BWD_FILTER_ALGO_0**

  ▸ Deterministic: No
  ▸ `xDesc` Format Support: All
  ▸ `dyDesc` Format Support: NCHW CHW-packed
  ▸ Data Type Config Support: All

▸ **CUDNN_CONVOLUTION_BWD_FILTER_ALGO_1**

  ▸ Deterministic: Yes
  ▸ `xDesc` Format Support: All
  ▸ `dyDesc` Format Support: NCHW CHW-packed
  ▸ Data Type Config Support: All

▸ **CUDNN_CONVOLUTION_BWD_FILTER_ALGO_FFT**

  ▸ Deterministic: Yes
  ▸ `xDesc` Format Support: NCHW CHW-packed
  ▸ `dyDesc` Format Support: NCHW CHW-packed
  ▸ Data Type Config Support: PSEUDO_HALF_CONFIG, FLOAT_CONFIG
  ▸ Notes:

- ▶ **xDesc**'s feature map height + 2 * **convDesc**'s zero-padding height must equal 256 or less
- ▶ **xDesc**'s feature map width + 2 * **convDesc**'s zero-padding width must equal 256 or less
- ▶ **convDesc**'s vertical and horizontal filter stride must equal 1
- ▶ **dwDesc**'s filter height must be greater than **convDesc**'s zero-padding height
- ▶ **dwDesc**'s filter width must be greater than **convDesc**'s zero-padding width

- ▶ **CUDNN_CONVOLUTION_BWD_FILTER_ALGO_3**
  - ▶ Deterministic: No
  - ▶ **xDesc** Format Support: All
  - ▶ **dyDesc** Format Support: NCHW CHW-packed
  - ▶ Data Type Config Support: All

The following is an exhaustive list of algo support for 3-d convolutions.

- ▶ **CUDNN_CONVOLUTION_BWD_FILTER_ALGO_0**
  - ▶ Deterministic: No
  - ▶ **xDesc** Format Support: All
  - ▶ **dyDesc** Format Support: NCDHW CDHW-packed
  - ▶ Data Type Config Support: All

- ▶ **CUDNN_CONVOLUTION_BWD_FILTER_ALGO_3**
  - ▶ Deterministic: No
  - ▶ **xDesc** Format Support: NCDHW-fully-packed
  - ▶ **dyDesc** Format Support: NCDHW-fully-packed
  - ▶ Data Type Config Support: All

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The operation was launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▶ At least one of the following is NULL: `handle`, `xDesc`, `dyDesc`, `convDesc`, `dwDesc`, `xData`, `dyData`, `dwData`, `alpha`, `beta` ▶ `xDesc` and `dyDesc` have a non-matching number of dimensions ▶ `xDesc` and `dwDesc` have a non-matching number of dimensions ▶ `xDesc` has fewer than three number of dimensions ▶ `xDesc`, `dyDesc` and `dwDesc` have a non-matching data type. ▶ `xDesc` and `dwDesc` have a non-matching number of input feature maps per image.
`CUDNN_STATUS_NOT_SUPPORTED`	At least one of the following conditions are met:

Return Value	Meaning
	▸ **xDesc** or **dyDesc** have negative tensor striding ▸ **xDesc**, **dyDesc** or **dwDesc** has a number of dimensions that is not 4 or 5 ▸ The chosen algo does not support the parameters provided; see above for exhaustive list of parameter support for each algo
CUDNN_STATUS_MAPPING_ERROR	An error occurs during the texture binding of the filter data.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.52. cudnnFindConvolutionBackwardDataAlgorithm

```
cudnnStatus_t
cudnnFindConvolutionBackwardDataAlgorithm(cudnnHandle_t
 handle,
                                           const cudnnFilterDescriptor_t
 wDesc,
                                           const cudnnTensorDescriptor_t
 dyDesc,
                                           const cudnnConvolutionDescriptor_t
 convDesc,
                                           const cudnnTensorDescriptor_t
 dxDesc,
                                           const int
 requestedAlgoCount,
                                           int
 *returnedAlgoCount,
                                           cudnnConvolutionBwdFilterAlgoPerf_t
 *perfResults );
```

This function attempts all cuDNN algorithms for **cudnnConvolutionBackwardData()**, using memory allocated via **cudaMalloc()** and outputs performance metrics to a user-allocated array of **cudnnConvolutionBwdDataAlgoPerf_t**. These metrics are written in sorted fashion where the first element has the lowest compute time.

> 💬 This function is host blocking.

> 💬 It is recommend to run this function prior to allocating layer data; doing otherwise may needlessly inhibit some algorithm options due to resource usage.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
wDesc	input	Handle to a previously initialized filter descriptor.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
convDesc	input	Previously initialized convolution descriptor.

Param	In/out	Meaning
dxDesc	input	Handle to the previously initialized output tensor descriptor.
requestedAlgoCount	input	The maximum number of elements to be stored in perfResults.
returnedAlgoCount	output	The number of output elements stored in perfResults.
perfResults	output	A user-allocated array to store performance metrics sorted ascending by compute time.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `handle` is not allocated properly. ▸ `wDesc, dyDesc` or `dxDesc` is not allocated properly. ▸ `wDesc, dyDesc` or `dxDesc` has fewer than 1 dimension. ▸ Either `returnedCount` or `perfResults` is nil. ▸ `requestedCount` is less than 1.
`CUDNN_STATUS_ALLOC_FAILED`	This function was unable to allocate memory to store sample input, filters and output.
`CUDNN_STATUS_INTERNAL_ERROR`	At least one of the following conditions are met: ▸ The function was unable to allocate neccesary timing objects. ▸ The function was unable to deallocate neccesary timing objects. ▸ The function was unable to deallocate sample input, filters and output.

4.53. cudnnFindConvolutionBackwardDataAlgorithmEx

```
cudnnStatus_t
cudnnFindConvolutionBackwardDataAlgorithmEx(cudnnHandle_t
    handle,
                                            const cudnnFilterDescriptor_t
    wDesc,
                                            const void
    *w,
                                            const cudnnTensorDescriptor_t
    dyDesc,
                                            const void
    *dy,
                                            const cudnnConvolutionDescriptor_t
    convDesc,
                                            const cudnnTensorDescriptor_t
    dxDesc,
                                            void
    *dx,
                                            const int
    requestedAlgoCount,
                                            int
 *returnedAlgoCount,
                                            cudnnConvolutionBwdFilterAlgoPerf_t
 *perfResults,
                                            void
    *workSpace,
                                            size_t
    workSpaceSizeInBytes );
```

This function attempts all cuDNN algorithms for **cudnnConvolutionBackwardData**, using user-allocated GPU memory, and outputs performance metrics to a user-allocated array of **cudnnConvolutionBwdDataAlgoPerf_t**. These metrics are written in sorted fashion where the first element has the lowest compute time.

> 💬 This function is host blocking.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
wDesc	input	Handle to a previously initialized filter descriptor.
w	input	Data pointer to GPU memory associated with the filter descriptor **wDesc**.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the filter descriptor **dyDesc**.
convDesc	input	Previously initialized convolution descriptor.
dxDesc	input	Handle to the previously initialized output tensor descriptor.
dxDesc	input/ output	Data pointer to GPU memory associated with the tensor descriptor **dxDesc**. The content of this tensor will be overwritten with arbitary values.
requestedAlgoCount	input	The maximum number of elements to be stored in perfResults.

Param	In/out	Meaning
returnedAlgoCount	output	The number of output elements stored in perfResults.
perfResults	output	A user-allocated array to store performance metrics sorted ascending by compute time.
workSpace	input	Data pointer to GPU memory that is a necessary workspace for some algorithms. The size of this workspace will determine the availabilty of algorithms. A nil pointer is considered a workSpace of 0 bytes.
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workSpace`

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `handle` is not allocated properly. ▸ `wDesc, dyDesc` or `dxDesc` is not allocated properly. ▸ `wDesc, dyDesc` or `dxDesc` has fewer than 1 dimension. ▸ `w, dy` or `dx` is nil. ▸ Either `returnedCount` or `perfResults` is nil. ▸ `requestedCount` is less than 1.
`CUDNN_STATUS_INTERNAL_ERROR`	At least one of the following conditions are met: ▸ The function was unable to allocate neccesary timing objects. ▸ The function was unable to deallocate neccesary timing objects. ▸ The function was unable to deallocate sample input, filters and output.

4.54. cudnnGetConvolutionBackwardDataAlgorithm

```
cudnnStatus_t
cudnnGetConvolutionBackwardDataAlgorithm(    cudnnHandle_t
    handle,
                                             const cudnnFilterDescriptor_t
    wDesc,
                                             const cudnnTensorDescriptor_t
    dyDesc,
                                             const cudnnConvolutionDescriptor_t
    convDesc,
                                             const cudnnTensorDescriptor_t
    dxDesc,
                                             cudnnConvolutionBwdDataPreference_t
    preference,
                                             size_t
    memoryLimitInbytes,
                                             cudnnConvolutionBwdDataAlgo_t
    *algo
                                         )
```

This function serves as a heuristic for obtaining the best suited algorithm for **cudnnConvolutionBackwardData** for the given layer specifications. Based on the input preference, this function will either return the fastest algorithm or the fastest algorithm within a given memory limit. For an exhaustive search for the fastest algorithm, please use **cudnnFindConvolutionBackwardDataAlgorithm**.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
wDesc	input	Handle to a previously initialized filter descriptor.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
convDesc	input	Previously initialized convolution descriptor.
dxDesc	input	Handle to the previously initialized output tensor descriptor.
preference	input	Enumerant to express the preference criteria in terms of memory requirement and speed.
memoryLimitInbytes	input	It is to specify the maximum amount of GPU memory the user is willing to use as a workspace. This is currently a placeholder and is not used.
algo	output	Enumerant that specifies which convolution algorithm should be used to compute the results according to the specified preference

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The numbers of feature maps of the input tensor and output tensor differ.

Return Value	Meaning
	▸ The `dataType` of the two tensor descriptors or the filter are different.

4.55. cudnnGetConvolutionBackwardDataWorkspaceSize

```
cudnnStatus_t
cudnnGetConvolutionBackwardDataWorkspaceSize(    cudnnHandle_t
    handle,
                                                 const cudnnFilterDescriptor_t
    wDesc,
                                                 const cudnnTensorDescriptor_t
    dyDesc,
                                                 const
 cudnnConvolutionDescriptor_t  convDesc,
                                                 const cudnnTensorDescriptor_t
    dxDesc,
                                                 cudnnConvolutionFwdAlgo_t
    algo,
                                                 size_t
   *sizeInBytes
                                )
```

This function returns the amount of GPU memory workspace the user needs to allocate to be able to call **cudnnConvolutionBackwardData** with the specified algorithm. The workspace allocated will then be passed to the routine **cudnnConvolutionBackwardData**. The specified algorithm can be the result of the call to **cudnnGetConvolutionBackwardDataAlgorithm** or can be chosen arbitrarily by the user. Note that not every algorithm is available for every configuration of the input tensor and/or every configuration of the convolution descriptor.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
wDesc	input	Handle to a previously initialized filter descriptor.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
convDesc	input	Previously initialized convolution descriptor.
dxDesc	input	Handle to the previously initialized output tensor descriptor.
algo	input	Enumerant that specifies the chosen convolution algorithm
sizeInBytes	output	Amount of GPU memory needed as workspace to be able to execute a forward convolution with the specified `algo`

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The numbers of feature maps of the input tensor and output tensor differ.

Return Value	Meaning
	▸ The `dataType` of the two tensor descriptors or the filter are different.
`CUDNN_STATUS_NOT_SUPPORTED`	The combination of the tensor descriptors, filter descriptor and convolution descriptor is not supported for the specified algorithm.

4.56. cudnnConvolutionBackwardData

```
cudnnStatus_t
cudnnConvolutionBackwardData( cudnnHandle_t                   handle,
                              const void                          *alpha,
                              const cudnnFilterDescriptor_t        wDesc,
                              const void                          *w,
                              const cudnnTensorDescriptor_t        dyDesc,
                              const void                          *dy,
                              const cudnnConvolutionDescriptor_t   convDesc,
                              cudnnConvolutionBwdDataAlgo_t        algo,
                              void                                *workSpace,
                              size_t
workSpaceSizeInBytes,
                              const void                          *beta,
                              const cudnnTensorDescriptor_t        dxDesc,
                              void                                *dx );
```

This function computes the convolution gradient with respect to the output tensor using the specified `algo`, returning results in `gradDesc`. Scaling factors `alpha` and `beta` can be used to scale the input tensor and the output tensor respectively.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
wDesc	input	Handle to a previously initialized filter descriptor.
w	input	Data pointer to GPU memory associated with the filter descriptor `wDesc`.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the input differential tensor descriptor `dyDesc`.
convDesc	input	Previously initialized convolution descriptor.
algo	input	Enumerant that specifies which backward data convolution algorithm shoud be used to compute the results
workSpace	input	Data pointer to GPU memory to a workspace needed to able to execute the specified algorithm. If no workspace is needed for a particular algorithm, that pointer can be nil
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workSpace`

Param	In/out	Meaning
dxDesc	input	Handle to the previously initialized output tensor descriptor.
dx	input/ output	Data pointer to GPU memory associated with the output tensor descriptor `dxDesc` that carries the result.

This function supports only three specific combinations of data types for `wDesc`, `dyDesc`, `convDesc` and `dxDesc`. See the following for an exhaustive list of these configurations.

Data Type Configurations	wDesc's, dyDesc's and dxDesc's Data Type	convDesc's Data Type
PSEUDO_HALF_CONFIG	CUDNN_DATA_HALF	CUDNN_DATA_FLOAT
FLOAT_CONFIG	CUDNN_DATA_FLOAT	CUDNN_DATA_FLOAT
DOUBLE_CONFIG	CUDNN_DATA_DOUBLE	CUDNN_DATA_DOUBLE

Specifying a separate algorithm can cause changes in performance, support and computation determinism. See the following for an exhaustive list of algorithm options and their respective supported parameters and deterministic behavior.

`wDesc` may only have format CUDNN_TENSOR_NHWC when all of the following are true:

- `algo` is CUDNN_CONVOLUTION_BWD_DATA_ALGO_1
- `dyDesc` and `dxDesc` is NHWC HWC-packed
- Data type configuration is PSEUDO_HALF_CONFIG or FLOAT_CONFIG
- The convolution is 2-dimensional

The following is an exhaustive list of algo support for 2-d convolutions.

- **CUDNN_CONVOLUTION_BWD_DATA_ALGO_0**

    - Deterministic: No
    - `dyDesc` Format Support: NCHW CHW-packed
    - `dxDesc` Format Support: All
    - Data Type Config Support: All
- **CUDNN_CONVOLUTION_BWD_DATA_ALGO_1**

    - Deterministic: Yes
    - `dyDesc` Format Support: NCHW CHW-packed
    - `dxDesc` Format Support: All
    - Data Type Config Support: All
- **CUDNN_CONVOLUTION_BWD_DATA_ALGO_FFT**

    - Deterministic: Yes
    - `dyDesc` Format Support: NCHW CHW-packed
    - `dxDesc` Format Support: NCHW HW-packed
    - Data Type Config Support: PSEUDO_HALF_CONFIG, FLOAT_CONFIG
    - Notes:

- ▶ **dxDesc**'s feature map height + 2 * **convDesc**'s zero-padding height must equal 256 or less
- ▶ **dxDesc**'s feature map width + 2 * **convDesc**'s zero-padding width must equal 256 or less
- ▶ **convDesc**'s vertical and horizontal filter stride must equal 1
- ▶ **wDesc**'s filter height must be greater than **convDesc**'s zero-padding height
- ▶ **wDesc**'s filter width must be greater than **convDesc**'s zero-padding width

- ▶ **CUDNN_CONVOLUTION_BWD_DATA_ALGO_FFT_TILING**
  - ▶ Deterministic: Yes
  - ▶ **dyDesc** Format Support: NCHW CHW-packed
  - ▶ **dxDesc** Format Support: NCHW HW-packed
  - ▶ Data Type Config Support: PSEUDO_HALF_CONFIG, FLOAT_CONFIG
  - ▶ Notes:
    - ▶ **wDesc**'s filter height must equal 32 or less
    - ▶ **wDesc**'s filter width must equal 32 or less
    - ▶ **convDesc**'s vertical and horizontal filter stride must equal 1
    - ▶ **wDesc**'s filter height must be greater than **convDesc**'s zero-padding height
    - ▶ **wDesc**'s filter width must be greater than **convDesc**'s zero-padding width

- ▶ **CUDNN_CONVOLUTION_BWD_DATA_ALGO_WINOGRAD**
  - ▶ Deterministic: Yes
  - ▶ **xDesc** Format Support: NCHW CHW-packed
  - ▶ **yDesc** Format Support: All
  - ▶ Data Type Config Support: PSEUDO_HALF_CONFIG, FLOAT_CONFIG
  - ▶ Notes:
    - ▶ **convDesc**'s vertical and horizontal filter stride must equal 1
    - ▶ **wDesc**'s filter height must be 3
    - ▶ **wDesc**'s filter width must be 3

The following is an exhaustive list of algo support for 3-d convolutions.

- ▶ **CUDNN_CONVOLUTION_BWD_DATA_ALGO_0**
  - ▶ Deterministic: No
  - ▶ **dyDesc** Format Support: NCDHW CDHW-packed
  - ▶ **dxDesc** Format Support: All
  - ▶ Data Type Config Support: All

- ▶ **CUDNN_CONVOLUTION_BWD_DATA_ALGO_1**
  - ▶ Deterministic: Yes
  - ▶ **dyDesc** Format Support: NCDHW-fully-packed
  - ▶ **dxDesc** Format Support: NCDHW-fully-packed
  - ▶ Data Type Config Support: All

- ▶ **CUDNN_CONVOLUTION_BWD_DATA_ALGO_FFT_TILING**
  - ▶ Deterministic: Yes

- ▸ **dyDesc** Format Support: NCDHW CDHW-packed
- ▸ **dxDesc** Format Support: NCDHW DHW-packed
- ▸ Data Type Config Support: All
- ▸ Notes:

  - ▸ **wDesc**'s filter height must equal 16 or less
  - ▸ **wDesc**'s filter width must equal 16 or less
  - ▸ **wDesc**'s filter depth must equal 16 or less
  - ▸ **convDesc**'s must have all filter strides equal to 1
  - ▸ **wDesc**'s filter height must be greater than **convDesc**'s zero-padding height
  - ▸ **wDesc**'s filter width must be greater than **convDesc**'s zero-padding width
  - ▸ **wDesc**'s filter depth must be greater than **convDesc**'s zero-padding width

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The operation was launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ At least one of the following is NULL: `handle`, `dyDesc`, `wDesc`, `convDesc`, `dxDesc`, `dy`, `w`, `dx`, `alpha`, `beta` ▸ `wDesc` and `dyDesc` have a non-matching number of dimensions ▸ `wDesc` and `dxDesc` have a non-matching number of dimensions ▸ `wDesc` has fewer than three number of dimensions ▸ `wDesc`, `dxDesc` and `dyDesc` have a non-matching data type. ▸ `wDesc` and `dxDesc` have a non-matching number of input feature maps per image. ▸ `dyDescs's` spatial sizes do not match with the expected size as determined by `cudnnGetConvolutionNdForwardOutputDim`
`CUDNN_STATUS_NOT_SUPPORTED`	At least one of the following conditions are met: ▸ `dyDesc` or `dxDesc` have negative tensor striding ▸ `dyDesc`, `wDesc` or `dxDesc` has a number of dimensions that is not 4 or 5 ▸ The chosen algo does not support the parameters provided; see above for exhaustive list of parameter support for each algo
`CUDNN_STATUS_MAPPING_ERROR`	An error occurs during the texture binding of the filter data or the input differential tensor data
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.57. cudnnSoftmaxForward

```
cudnnStatus_t
cudnnSoftmaxForward( cudnnHandle_t                    handle,
                     cudnnSoftmaxAlgorithm_t          algorithm,
                     cudnnSoftmaxMode_t               mode,
                     const void                      *alpha,
                     const cudnnTensorDescriptor_t    xDesc,
                     const void                      *x,
                     const void                      *beta,
                     const cudnnTensorDescriptor_t    yDesc,
                     void                            *y );
```

This routine computes the softmax function.

> All tensor formats are supported for all modes and algorithms with 4 and 5D tensors. Performance is expected to be highest with **NCHW fully-packed** tensors. For more than 5 dimensions tensors must be packed in their spatial dimensions

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
algorithm	input	Enumerant to specify the softmax algorithm.
mode	input	Enumerant to specify the softmax mode.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor **xDesc**.
yDesc	input	Handle to the previously initialized output tensor descriptor.
y	output	Data pointer to GPU memory associated with the output tensor descriptor **yDesc**.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The function launched successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The dimensions **n,c,h,w** of the input tensor and output tensors differ. ▸ The **datatype** of the input tensor and output tensors differ. ▸ The parameters **algorithm** or **mode** have an invalid enumerant value.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.58. cudnnSoftmaxBackward

```
cudnnStatus_t
cudnnSoftmaxBackward( cudnnHandle_t                    handle,
                      cudnnSoftmaxAlgorithm_t          algorithm,
                      cudnnSoftmaxMode_t               mode,
                      const void                      *alpha,
                      const cudnnTensorDescriptor_t    yDesc,
                      const void                      *yData,
                      const cudnnTensorDescriptor_t    dyDesc,
                      const void                      *dy,
                      const void                      *beta,
                      const cudnnTensorDescriptor_t    dxDesc,
                      void                            *dx );
```

This routine computes the gradient of the softmax function.

> 💬 In-place operation is allowed for this routine; i.e., `dy` and `dx` pointers may be equal. However, this requires `dyDesc` and `dxDesc` descriptors to be identical (particularly, the strides of the input and output must match for in-place operation to be allowed).

> 💬 All tensor formats are supported for all modes and algorithms with 4 and 5D tensors. Performance is expected to be highest with `NCHW fully-packed` tensors. For more than 5 dimensions tensors must be packed in their spatial dimensions

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
algorithm	input	Enumerant to specify the softmax algorithm.
mode	input	Enumerant to specify the softmax mode.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
yDesc	input	Handle to the previously initialized input tensor descriptor.
y	input	Data pointer to GPU memory associated with the tensor descriptor `yDesc`.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the tensor descriptor `dyData`.
dxDesc	input	Handle to the previously initialized output differential tensor descriptor.
dx	output	Data pointer to GPU memory associated with the output tensor descriptor `dxDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met:

Return Value	Meaning
	- The dimensions `n`,`c`,`h`,`w` of the `yDesc`, `dyDesc` and `dxDesc` tensors differ. - The strides `nStride`, `cStride`, `hStride`, `wStride` of the `yDesc` and `dyDesc` tensors differ. - The `datatype` of the three tensors differs.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.59. cudnnCreatePoolingDescriptor

```
cudnnStatus_t cudnnCreatePoolingDescriptor( cudnnPoolingDescriptor_t*
 poolingDesc )
```

This function creates a pooling descriptor object by allocating the memory needed to hold its opaque structure,

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was created successfully.
`CUDNN_STATUS_ALLOC_FAILED`	The resources could not be allocated.

4.60. cudnnSetPooling2dDescriptor

```
cudnnStatus_t
cudnnSetPooling2dDescriptor( cudnnPoolingDescriptor_t poolingDesc,
                             cudnnPoolingMode_t mode,
                             cudnnNanPropagation_t maxpoolingNanOpt,
                             int windowHeight,
                             int windowWidth,
                             int verticalPadding,
                             int horizontalPadding,
                             int verticalStride,
                             int horizontalStride )
```

This function initializes a previously created generic pooling descriptor object into a 2D description.

Param	In/out	Meaning
poolingDesc	input/output	Handle to a previously created pooling descriptor.
mode	input	Enumerant to specify the pooling mode.
maxpoolingNanOpt	input	Enumerant to specify the Nan propagation mode.
windowHeight	input	Height of the pooling window.
windowWidth	input	Width of the pooling window.
verticalPadding	input	Size of vertical padding.
horizontalPadding	input	Size of horizontal padding

Param	In/out	Meaning
verticalStride	input	Pooling vertical stride.
horizontalStride	input	Pooling horizontal stride.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was set successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the parameters **windowHeight**, **windowWidth**, **verticalStride**, **horizontalStride** is negative or **mode** or **maxpoolingNanOpt** has an invalid enumerant value.

4.61. cudnnGetPooling2dDescriptor

```
cudnnStatus_t
cudnnGetPooling2dDescriptor( const cudnnPoolingDescriptor_t poolingDesc,
                             cudnnPoolingMode_t *mode,
                             cudnnNanPropagation_t *maxpoolingNanOpt,
                             int *windowHeight,
                             int *windowWidth,
                             int *verticalPadding,
                             int *horizontalPadding,
                             int *verticalStride,
                             int *horizontalStride )
```

This function queries a previously created 2D pooling descriptor object.

Param	In/out	Meaning
poolingDesc	input	Handle to a previously created pooling descriptor.
mode	output	Enumerant to specify the pooling mode.
maxpoolingNanOpt	output	Enumerant to specify the Nan propagation mode.
windowHeight	output	Height of the pooling window.
windowWidth	output	Width of the pooling window.
verticalPadding	output	Size of vertical padding.
horizontalPadding	output	Size of horizontal padding.
verticalStride	output	Pooling vertical stride.
horizontalStride	output	Pooling horizontal stride.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was set successfully.

4.62. cudnnSetPoolingNdDescriptor

```
cudnnStatus_t
cudnnSetPoolingNdDescriptor( cudnnPoolingDescriptor_t poolingDesc,
                             cudnnPoolingMode_t mode,
                             cudnnNanPropagation_t maxpoolingNanOpt,
                             int nbDims,
                             int windowDimA[],
                             int paddingA[],
                             int strideA[] )
```

This function initializes a previously created generic pooling descriptor object.

Param	In/out	Meaning
poolingDesc	input/ output	Handle to a previously created pooling descriptor.
mode	input	Enumerant to specify the pooling mode.
maxpoolingNanOpt	input	Enumerant to specify the Nan propagation mode.
nbDims	input	Dimension of the pooling operation.
windowDimA	output	Array of dimension `nbDims` containing the window size for each dimension.
paddingA	output	Array of dimension `nbDims` containing the padding size for each dimension.
strideA	output	Array of dimension `nbDims` containing the striding size for each dimension.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the elements of the arrays `windowDimA, paddingA or strideA` is negative or `mode` or `maxpoolingNanOpt` has an invalid enumerant value.

4.63. cudnnGetPoolingNdDescriptor

```
cudnnStatus_t
cudnnGetPoolingNdDescriptor( const cudnnPoolingDescriptor_t poolingDesc,
                             int nbDimsRequested,
                             cudnnPoolingMode_t *mode,
                             cudnnNanPropagation_t *maxpoolingNanOpt,
                             int *nbDims,
                             int windowDimA[],
                             int paddingA[],
                             int strideA[] )
```

This function queries a previously initialized generic pooling descriptor object.

Param	In/out	Meaning
poolingDesc	input	Handle to a previously created pooling descriptor.
nbDimsRequested	input	Dimension of the expected pooling descriptor. It is also the minimum size of the arrays `windowDimA`, `paddingA` and `strideA` in order to be able to hold the results
mode	output	Enumerant to specify the pooling mode.
maxpoolingNanOpt	input	Enumerant to specify the Nan propagation mode.
nbDims	output	Actual dimension of the pooling descriptor.
windowDimA	output	Array of dimension of at least `nbDimsRequested` that will be filled with the window parameters from the provided pooling descriptor.
paddingA	output	Array of dimension of at least `nbDimsRequested` that will be filled with the padding parameters from the provided pooling descriptor.
strideA	output	Array of dimension at least `nbDimsRequested` that will be filled with the stride parameters from the provided pooling descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was queried successfully.
`CUDNN_STATUS_NOT_SUPPORTED`	The parameter `nbDimsRequested` is greater than CUDNN_DIM_MAX.

4.64. cudnnSetPooling2dDescriptor_v3

```
cudnnStatus_t
cudnnSetPooling2dDescriptor_v3( cudnnPoolingDescriptor_t poolingDesc,
                                cudnnPoolingMode_t mode,
                                int windowHeight,
                                int windowWidth,
                                int verticalPadding,
                                int horizontalPadding,
                                int verticalStride,
                                int horizontalStride )
```

This function initializes a previously created generic pooling descriptor object into a 2D description.

This routine is deprecated, `cudnnSetPooling2dDescriptor` should be used instead.

Param	In/out	Meaning
poolingDesc	input/output	Handle to a previously created pooling descriptor.
mode	input	Enumerant to specify the pooling mode.

Param	In/out	Meaning
windowHeight	input	Height of the pooling window.
windowWidth	input	Width of the pooling window.
verticalPadding	input	Size of vertical padding.
horizontalPadding	input	Size of horizontal padding
verticalStride	input	Pooling vertical stride.
horizontalStride	input	Pooling horizontal stride.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the parameters `windowHeight`, `windowWidth`, `verticalStride`, `horizontalStride` is negative or `mode` has an invalid enumerant value.

4.65. cudnnGetPooling2dDescriptor_v3

```
cudnnStatus_t
cudnnGetPooling2dDescriptor_v3( const cudnnPoolingDescriptor_t poolingDesc,
                                cudnnPoolingMode_t *mode,
                                int *windowHeight,
                                int *windowWidth,
                                int *verticalPadding,
                                int *horizontalPadding,
                                int *verticalStride,
                                int *horizontalStride )
```

This function queries a previously created 2D pooling descriptor object.

> 💬 This routine is deprecated, `cudnnGetPooling2dDescriptor` should be used instead.

Param	In/out	Meaning
poolingDesc	input	Handle to a previously created pooling descriptor.
mode	output	Enumerant to specify the pooling mode.
windowHeight	output	Height of the pooling window.
windowWidth	output	Width of the pooling window.
verticalPadding	output	Size of vertical padding.
horizontalPadding	output	Size of horizontal padding.
verticalStride	output	Pooling vertical stride.
horizontalStride	output	Pooling horizontal stride.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.

4.66. cudnnSetPoolingNdDescriptor_v3

```
cudnnStatus_t
cudnnSetPoolingNdDescriptor_v3( cudnnPoolingDescriptor_t poolingDesc,
                                cudnnPoolingMode_t mode,
                                int nbDims,
                                int windowDimA[],
                                int paddingA[],
                                int strideA[] )
```

This function initializes a previously created generic pooling descriptor object.

> 💬 This routine is deprecated, `cudnnSetPoolingNdDescriptor` should be used instead.

Param	In/out	Meaning
poolingDesc	input/output	Handle to a previously created pooling descriptor.
mode	input	Enumerant to specify the pooling mode.
nbDims	input	Dimension of the pooling operation.
windowDimA	output	Array of dimension `nbDims` containing the window size for each dimension.
paddingA	output	Array of dimension `nbDims` containing the padding size for each dimension.
strideA	output	Array of dimension `nbDims` containing the striding size for each dimension.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the elements of the arrays `windowDimA, paddingA or strideA` is negative or `mode` has an invalid enumerant value.

4.67. cudnnGetPoolingNdDescriptor_v3

```
cudnnStatus_t
cudnnGetPoolingNdDescriptor_v3( const cudnnPoolingDescriptor_t poolingDesc,
                                int nbDimsRequested,
                                cudnnPoolingMode_t *mode,
                                int *nbDims,
                                int windowDimA[],
                                int paddingA[],
                                int strideA[] )
```

This function queries a previously initialized generic pooling descriptor object.

> This routine is deprecated, **cudnnGetPoolingNdDescriptor** should be used instead.

Param	In/out	Meaning
poolingDesc	input	Handle to a previously created pooling descriptor.
nbDimsRequested	input	Dimension of the expected pooling descriptor. It is also the minimum size of the arrays **windowDimA**, **paddingA** and **strideA** in order to be able to hold the results
mode	output	Enumerant to specify the pooling mode.
nbDims	output	Actual dimension of the pooling descriptor.
windowDimA	output	Array of dimension of at least **nbDimsRequested** that will be filled with the window parameters from the provided pooling descriptor.
paddingA	output	Array of dimension of at least **nbDimsRequested** that will be filled with the padding parameters from the provided pooling descriptor.
strideA	output	Array of dimension at least **nbDimsRequested** that will be filled with the stride parameters from the provided pooling descriptor.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was queried successfully.
CUDNN_STATUS_NOT_SUPPORTED	The parameter **nbDimsRequested** is greater than CUDNN_DIM_MAX.

4.68. cudnnSetPooling2dDescriptor_v4

```
cudnnStatus_t
cudnnSetPooling2dDescriptor_v4( cudnnPoolingDescriptor_t poolingDesc,
                                cudnnPoolingMode_t mode,
                                cudnnNanPropagation_t maxpoolingNanOpt,
                                int windowHeight,
                                int windowWidth,
                                int verticalPadding,
                                int horizontalPadding,
                                int verticalStride,
                                int horizontalStride )
```

This function is equivalent to **cudnnSetPooling2dDescriptor**.

4.69. cudnnGetPooling2dDescriptor_v4

```
cudnnStatus_t
cudnnGetPooling2dDescriptor_v4( const cudnnPoolingDescriptor_t poolingDesc,
                                cudnnPoolingMode_t *mode,
                                cudnnNanPropagation_t *maxpoolingNanOpt,
                                int *windowHeight,
                                int *windowWidth,
                                int *verticalPadding,
                                int *horizontalPadding,
                                int *verticalStride,
                                int *horizontalStride )
```

This function is equivalent to **cudnnGetPooling2dDescriptor**.

4.70. cudnnSetPoolingNdDescriptor_v4

```
cudnnStatus_t
cudnnSetPoolingNdDescriptor_v4( cudnnPoolingDescriptor_t poolingDesc,
                                cudnnPoolingMode_t mode,
                                cudnnNanPropagation_t maxpoolingNanOpt,
                                int nbDims,
                                int windowDimA[],
                                int paddingA[],
                                int strideA[] )
```

This function is equivalent to **cudnnSetPoolingNdDescriptor**.

4.71. cudnnGetPoolingNdDescriptor_v4

```
cudnnStatus_t
cudnnGetPoolingNdDescriptor_v4( const cudnnPoolingDescriptor_t poolingDesc,
                                int nbDimsRequested,
                                cudnnPoolingMode_t *mode,
                                cudnnNanPropagation_t *maxpoolingNanOpt,
                                int *nbDims,
                                int windowDimA[],
                                int paddingA[],
                                int strideA[] )
```

This function is equivalent to **cudnnGetPoolingNdDescriptor**.

4.72. cudnnDestroyPoolingDescriptor

```
cudnnStatus_t cudnnDestroyPoolingDescriptor( cudnnPoolingDescriptor_t
 poolingDesc )
```

This function destroys a previously created pooling descriptor object.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was destroyed successfully.

4.73. cudnnGetPooling2dForwardOutputDim

```
cudnnStatus_t
cudnnGetPooling2dForwardOutputDim( const cudnnPoolingDescriptor_t    poolingDesc,
                                   const cudnnTensorDescriptor_t     inputDesc,
                                   int *outN,
                                   int *outC,
                                   int *outH,
                                   int *outW )
```

This function provides the output dimensions of a tensor after 2d pooling has been applied

Each dimension `h and w` of the output images is computed as followed:

```
outputDim = 1 + (inputDim + 2*padding - windowDim)/poolingStride;
```

Param	In/out	Meaning
poolingDesc	input	Handle to a previously inititalized pooling descriptor.
inputDesc	input	Handle to the previously initialized input tensor descriptor.
N	output	Number of images in the output
C	output	Number of channels in the output
H	output	Height of images in the output
W	output	Width of images in the output

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `poolingDesc` has not been initialized. ▸ `poolingDesc` or `inputDesc` has an invalid number of dimensions (2 and 4 respectively are required).

4.74. cudnnGetPoolingNdForwardOutputDim

```
cudnnStatus_t
cudnnGetPoolingNdForwardOutputDim( const cudnnPoolingDescriptor_t    poolingDesc,
                                   const cudnnTensorDescriptor_t     inputDesc,
                                   int nbDims,
                                   int outDimA[] )
```

This function provides the output dimensions of a tensor after Nd pooling has been applied

Each dimension of the `(nbDims-2)-D` images of the output tensor is computed as followed:

```
outputDim = 1 + (inputDim + 2*padding - windowDim)/poolingStride;
```

Param	In/out	Meaning
poolingDesc	input	Handle to a previously inititalized pooling descriptor.
inputDesc	input	Handle to the previously initialized input tensor descriptor.
nbDims	input	Number of dimensions in which pooling is to be applied.
outDimA	output	Array of nbDims output dimensions

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The function launched successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ **poolingDesc** has not been initialized. ▸ The value of **nbDims** is inconsistent with the dimensionality of **poolingDesc** and **inputDesc**.

4.75. cudnnPoolingForward

```
cudnnStatus_t
cudnnPoolingForward( cudnnHandle_t                    handle,
                     const cudnnPoolingDescriptor_t   poolingDesc,
                     const void                       *alpha,
                     const cudnnTensorDescriptor_t    xDesc,
                     const void                       *x,
                     const void                       *beta,
                     const cudnnTensorDescriptor_t    yDesc,
                     void                             *y );
```

This function computes pooling of input values (i.e., the maximum or average of several adjacent values) to produce an output with smaller height and/or width.

> All tensor formats are supported, best performance is expected when using `HW-packed` tensors. Only 2 and 3 spatial dimensions are allowed.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
poolingDesc	input	Handle to a previously initialized pooling descriptor.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor `xDesc`.
yDesc	input	Handle to the previously initialized output tensor descriptor.
y	output	Data pointer to GPU memory associated with the output tensor descriptor `yDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The dimensions `n,c` of the input tensor and output tensors differ. ▸ The `datatype` of the input tensor and output tensors differs.
`CUDNN_STATUS_NOT_SUPPORTED`	The `wStride` of input tensor or output tensor is not 1.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.76. cudnnPoolingBackward

```
cudnnStatus_t
cudnnPoolingBackward( cudnnHandle_t handle,
                      const cudnnPoolingDescriptor_t    poolingDesc,
                      const void                        *alpha,
                      const cudnnTensorDescriptor_t     yDesc,
                      const void                        *y,
                      const cudnnTensorDescriptor_t     dyDesc,
                      const void                        *dy,
                      const cudnnTensorDescriptor_t     xDesc,
                      const void                        *xData,
                      const void                        *beta,
                      const cudnnTensorDescriptor_t     dxDesc,
                      void                              *dx )
```

This function computes the gradient of a pooling operation.

> All tensor formats are supported, best performance is expected when using **HW-packed** tensors. Only 2 and 3 spatial dimensions are allowed

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
poolingDesc	input	Handle to the previously initialized pooling descriptor.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
yDesc	input	Handle to the previously initialized input tensor descriptor.
y	input	Data pointer to GPU memory associated with the tensor descriptor **yDesc**.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the tensor descriptor **dyData**.
xDesc	input	Handle to the previously initialized output tensor descriptor.
x	input	Data pointer to GPU memory associated with the output tensor descriptor **xDesc**.
dxDesc	input	Handle to the previously initialized output differential tensor descriptor.
dx	output	Data pointer to GPU memory associated with the output tensor descriptor **dxDesc**.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The function launched successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The dimensions **n,c,h,w** of the **yDesc** and **dyDesc** tensors differ.

Return Value	Meaning
	▸ The strides `nStride`, `cStride`, `hStride`, `wStride` of the `yDesc` and `dyDesc` tensors differ. ▸ The dimensions `n,c,h,w` of the `dxDesc` and `dxDesc` tensors differ. ▸ The strides `nStride`, `cStride`, `hStride`, `wStride` of the `xDesc` and `dxDesc` tensors differ. ▸ The `datatype` of the four tensors differ.
CUDNN_STATUS_NOT_SUPPORTED	The `wstride` of input tensor or output tensor is not 1.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.77. cudnnActivationForward

```
cudnnStatus_t
    cudnnActivationForward( cudnnHandle_t handle,
                            cudnnActivationDescriptor_t    activationDesc,
                            const void                     *alpha,
                            const cudnnTensorDescriptor_t   srcDesc,
                            const void                     *srcData,
                            const void                     *beta,
                            const cudnnTensorDescriptor_t   destDesc,
                            void                           *destData )
```

This routine applies a specified neuron activation function element-wise over each input value.

💬 In-place operation is allowed for this routine; i.e., `xData` and `yData` pointers may be equal. However, this requires `xDesc` and `yDesc` descriptors to be identical (particularly, the strides of the input and output must match for in-place operation to be allowed).

💬 All tensor formats are supported for 4 and 5 dimensions, however best performance is obtained when the strides of `xDesc` and `yDesc` are equal and `HW-packed`. For more than 5 dimensions the tensors must have their spatial dimensions packed.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
activationDesc,	input	Activation descriptor.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor `xDesc`.

Param	In/out	Meaning
yDesc	input	Handle to the previously initialized output tensor descriptor.
y	output	Data pointer to GPU memory associated with the output tensor descriptor `yDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The parameter `mode` has an invalid enumerant value. ▸ The dimensions `n,c,h,w` of the input tensor and output tensors differ. ▸ The `datatype` of the input tensor and output tensors differs. ▸ The strides `nStride,cStride,hStride,wStride` of the input tensor and output tensors differ and in-place operation is used (i.e., `x` and `y` pointers are equal).
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.78. cudnnActivationBackward

```
cudnnStatus_t
    cudnnActivationBackward( cudnnHandle_t                    handle,
                             cudnnActivationDescriptor_t      activationDesc,
                             const void                       *alpha,
                             const cudnnTensorDescriptor_t    srcDesc,
                             const void                       *srcData,
                             const cudnnTensorDescriptor_t    srcDiffDesc,
                             const void                       *srcDiffData,
                             const cudnnTensorDescriptor_t    destDesc,
                             const void                       *destData,
                             const void                       *beta,
                             const cudnnTensorDescriptor_t    destDiffDesc,
                             void                             *destDiffData)
```

This routine computes the gradient of a neuron activation function.

> 💬 In-place operation is allowed for this routine; i.e. `dy` and `dx` pointers may be equal. However, this requires the corresponding tensor descriptors to be identical (particularly, the strides of the input and output must match for in-place operation to be allowed).

> 💬 All tensor formats are supported for 4 and 5 dimensions, however best performance is obtained when the strides of `yDesc` and `xDesc` are equal and `HW-packed`. For more than 5 dimensions the tensors must have their spatial dimensions packed.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
activationDesc	input	Activation descriptor.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
yDesc	input	Handle to the previously initialized input tensor descriptor.
y	input	Data pointer to GPU memory associated with the tensor descriptor `yDesc`.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the tensor descriptor `dyDesc`.
xDesc	input	Handle to the previously initialized output tensor descriptor.
x	input	Data pointer to GPU memory associated with the output tensor descriptor `xDesc`.
dxDesc	input	Handle to the previously initialized output differential tensor descriptor.
dx	output	Data pointer to GPU memory associated with the output tensor descriptor `dxDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ‣ The strides `nStride`, `cStride`, `hStride`, `wStride` of the input differential tensor and output differential tensors differ and in-place operation is used.
`CUDNN_STATUS_NOT_SUPPORTED`	At least one of the following conditions are met: ‣ The dimensions `n,c,h,w` of the input tensor and output tensors differ. ‣ The `datatype` of the input tensor and output tensors differs. ‣ The strides `nStride`, `cStride`, `hStride`, `wStride` of the input tensor and the input differential tensor differ. ‣ The strides `nStride`, `cStride`, `hStride`, `wStride` of the output tensor and the output differential tensor differ.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.79. cudnnCreateActivationDescriptor

```
cudnnStatus_t
    cudnnCreateActivationDescriptor( cudnnActivationDescriptor_t
 *activationDesc )
```

This function creates a activation descriptor object by allocating the memory needed to hold its opaque structure.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was created successfully.
`CUDNN_STATUS_ALLOC_FAILED`	The resources could not be allocated.

4.80. cudnnSetActivationDescriptor

```
cudnnStatus_t
cudnnSetActivationDescriptor( cudnnActivationDescriptor_t
 activationDesc,

                            cudnnActivationMode_t              mode,
                            cudnnNanPropagation_t              reluNanOpt,
                            double                             reluCeiling )
```

This function initializes a previously created generic activation descriptor object.

Param	In/out	Meaning
activationDesc	input/ output	Handle to a previously created pooling descriptor.
mode	input	Enumerant to specify the activation mode.
reluNanOpt,	input	Enumerant to specify the `Nan` propagation mode.
reluCeiling	input	floating point number to specify the clipping threashod when the activation mode is set to `CUDNN_ACTIVATION_CLIPPED_RELU`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	mode or reluNanOpt has an invalid enumerant value.

4.81. cudnnGetActivationDescriptor

```
cudnnStatus_t
    cudnnGetActivationDescriptor( const cudnnActivationDescriptor_t
activationDesc,
                                  cudnnActivationMode_t              *mode,
                                  cudnnNanPropagation_t
*reluNanOpt,
                                  double
*reluCeiling )
```

This function queries a previously initialized generic activation descriptor object.

Param	In/out	Meaning
activationDesc	input	Handle to a previously created activation descriptor.
mode	output	Enumerant to specify the activation mode.
reluNanOpt,	output	Enumerant to specify the **Nan** propagation mode.
reluCeiling	output	floating point number to specify the clipping threshod when the activation mode is set to **CUDNN_ACTIVATION_CLIPPED_RELU**.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was queried successfully.

4.82. cudnnDestroyActivationDescriptor

```
cudnnStatus_t
    cudnnDestroyActivationDescriptor( cudnnActivationDescriptor_t
activationDesc )
```

This function destroys a previously created activation descriptor object.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was destroyed successfully.

4.83. cudnnActivationForward_v3

```
cudnnStatus_t
cudnnActivationForward_v3( cudnnHandle_t                    handle,
                   cudnnActivationMode_t            mode,
                   const void                      *alpha,
                   const cudnnTensorDescriptor_t    xDesc,
                   const void                      *xData,
                   const void                      *beta,
                   const cudnnTensorDescriptor_t    yDesc,
                   void                            *yData )
```

This routine applies a specified neuron activation function element-wise over each input value.

> 💬 This routine is deprecated, `cudnnActivationForward` should be used instead.

> 💬 In-place operation is allowed for this routine; i.e., `xData` and `yData` pointers may be equal. However, this requires `xDesc` and `yDesc` descriptors to be identical (particularly, the strides of the input and output must match for in-place operation to be allowed).

> 💬 All tensor formats are supported for 4 and 5 dimensions, however best performance is obtained when the strides of `xDesc` and `yDesc` are equal and `HW-packed`. For more than 5 dimensions the tensors must have their spatial dimensions packed.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
mode	input	Enumerant to specify the activation mode.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor `xDesc`.
yDesc	input	Handle to the previously initialized output tensor descriptor.
y	output	Data pointer to GPU memory associated with the output tensor descriptor `yDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The parameter `mode` has an invalid enumerant value. ▸ The dimensions `n,c,h,w` of the input tensor and output tensors differ. ▸ The `datatype` of the input tensor and output tensors differs. ▸ The strides `nStride,cStride,hStride,wStride` of the input tensor and output tensors differ and in-place operation is used (i.e., `x` and `y` pointers are equal).

Return Value	Meaning
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.84. cudnnActivationBackward_v3

```
cudnnStatus_t
cudnnActivationBackward_v3( cudnnHandle_t                    handle,
                        cudnnActivationMode_t          mode,
                        const void                    *alpha,
                        const cudnnTensorDescriptor_t   yDesc,
                        const void                    *y,
                        const cudnnTensorDescriptor_t   dyDesc,
                        const void                    *dy,
                        const cudnnTensorDescriptor_t   xDesc,
                        const void                    *x,
                        const void                    *beta,
                        const cudnnTensorDescriptor_t   dxDesc,
                        void                          *dx );
```

This routine computes the gradient of a neuron activation function.

> This routine is deprecated, `cudnnActivationBackward` should be used instead.

> In-place operation is allowed for this routine; i.e. `dy` and `dx` pointers may be equal. However, this requires the corresponding tensor descriptors to be identical (particularly, the strides of the input and output must match for in-place operation to be allowed).

> All tensor formats are supported for 4 and 5 dimensions, however best performance is obtained when the strides of `yDesc` and `xDesc` are equal and `HW-packed`. For more than 5 dimensions the tensors must have their spatial dimensions packed.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
mode	input	Enumerant to specify the activation mode.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the computation result with prior value in the output layer as follows: dstValue = alpha[0]*result + beta[0]*priorDstValue. Please refer to this section for additional details.
yDesc	input	Handle to the previously initialized input tensor descriptor.
y	input	Data pointer to GPU memory associated with the tensor descriptor `yDesc`.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the tensor descriptor `dyDesc`.
xDesc	input	Handle to the previously initialized output tensor descriptor.
x	input	Data pointer to GPU memory associated with the output tensor descriptor `xDesc`.

Param	In/out	Meaning
dxDesc	input	Handle to the previously initialized output differential tensor descriptor.
dx	output	Data pointer to GPU memory associated with the output tensor descriptor `dxDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The parameter `mode` has an invalid enumerant value. ▸ The strides `nStride`, `cStride`, `hStride`, `wStride` of the input differential tensor and output differential tensors differ and in-place operation is used.
`CUDNN_STATUS_NOT_SUPPORTED`	At least one of the following conditions are met: ▸ The dimensions `n,c,h,w` of the input tensor and output tensors differ. ▸ The `datatype` of the input tensor and output tensors differs. ▸ The strides `nStride`, `cStride`, `hStride`, `wStride` of the input tensor and the input differential tensor differ. ▸ The strides `nStride`, `cStride`, `hStride`, `wStride` of the output tensor and the output differential tensor differ.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.85. cudnnActivationForward_v4

```
cudnnStatus_t
    cudnnActivationForward_v4( cudnnHandle_t handle,
                              cudnnActivationDescriptor_t    activationDesc,
                              const void                     *alpha,
                              const cudnnTensorDescriptor_t   srcDesc,
                              const void                     *srcData,
                              const void                     *beta,
                              const cudnnTensorDescriptor_t   destDesc,
                              void                           *destData )
```

This routine is equivalent to **cudnnActivationForward**.

4.86. cudnnActivationBackward_v4

```
cudnnStatus_t
    cudnnActivationBackward_v4( cudnnHandle_t               handle,
                                cudnnActivationDescriptor_t  activationDesc,
                                const void                  *alpha,
                                const cudnnTensorDescriptor_t  srcDesc,
                                const void                  *srcData,
                                const cudnnTensorDescriptor_t  srcDiffDesc,
                                const void                  *srcDiffData,
                                const cudnnTensorDescriptor_t  destDesc,
                                const void                  *destData,
                                const void                  *beta,
                                const cudnnTensorDescriptor_t  destDiffDesc,
                                void                        *destDiffData)
```

This routine is equivalent to **cudnnActivationBackward**.

4.87. cudnnCreateLRNDescriptor

```
cudnnStatus_t cudnnCreateLRNDescriptor( cudnnLRNDescriptor_t* poolingDesc )
```

This function allocates the memory needed to hold the data needed for LRN and DivisiveNormalization layers operation and returns a descriptor used with subsequent layer forward and backward calls.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was created successfully.
CUDNN_STATUS_ALLOC_FAILED	The resources could not be allocated.

4.88. cudnnSetLRNDescriptor

```
cudnnStatus_t
CUDNNWINAPI cudnnSetLRNDescriptor( cudnnLRNDescriptor_t   normDesc,
                                   unsigned               lrnN,
                                   double                 lrnAlpha,
                                   double                 lrnBeta,
                                   double                 lrnK );
```

This function initializes a previously created LRN descriptor object.

> Macros CUDNN_LRN_MIN_N, CUDNN_LRN_MAX_N, CUDNN_LRN_MIN_K, CUDNN_LRN_MIN_BETA defined in cudnn.h specify valid ranges for parameters.

> Values of double parameters will be cast down to the tensor datatype during computation.

Param	In/out	Meaning
normDesc	output	Handle to a previously created LRN descriptor.
lrnN	input	Normalization window width in elements. LRN layer uses a window [center-lookBehind, center+lookAhead], where lookBehind = floor((lrnN-1)/2), lookAhead = lrnN-lookBehind-1. So for n=10, the window is [k-4...k...k+5] with a total of 10 samples. For DivisiveNormalization layer the window has the same extents as above in all 'spatial' dimensions (dimA[2], dimA[3], dimA[4]). By default lrnN is set to 5 in cudnnCreateLRNDescriptor.
lrnAlpha	input	Value of the alpha variance scaling parameter in the normalization formula. Inside the library code this value is divided by the window width for LRN and by (window width)^#spatialDimensions for DivisiveNormalization. By default this value is set to 1e-4 in cudnnCreateLRNDescriptor.
lrnBeta	input	Value of the beta power parameter in the normalization formula. By default this value is set to 0.75 in cudnnCreateLRNDescriptor.
lrnK	input	Value of the k parameter in normalization formula. By default this value is set to 2.0.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	One of the input parameters was out of valid range as described above.

4.89. cudnnGetLRNDescriptor

```
cudnnStatus_t
CUDNNWINAPI cudnnGetLRNDescriptor( cudnnLRNDescriptor_t    normDesc,
                                   unsigned               *lrnN,
                                   double                 *lrnAlpha,
                                   double                 *lrnBeta,
                                   double                 *lrnK );
```

This function retrieves values stored in the previously initialized LRN descriptor object.

Param	In/out	Meaning
normDesc	output	Handle to a previously created LRN descriptor.
lrnN, lrnAlpha, lrnBeta, lrnK	output	Pointers to receive values of parameters stored in the descriptor object. See cudnnSetLRNDescriptor for more details. Any of these pointers can be NULL (no value is returned for the corresponding parameter).

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	Function completed successfully.

4.90. cudnnDestroyLRNDescriptor

```
cudnnStatus_t cudnnDestroyLRNDescriptor(cudnnLRNDescriptor_t lrnDesc)
```

This function destroys a previously created LRN descriptor object.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was destroyed successfully.

4.91. cudnnLRNCrossChannelForward

```
cudnnStatus_t CUDNNWINAPI cudnnLRNCrossChannelForward(
    cudnnHandle_t                    handle,
    cudnnLRNDescriptor_t             normDesc,
    cudnnLRNMode_t                   lrnMode,
    const void*                      alpha,
    const cudnnTensorDescriptor_t    xDesc,
    const void                       *x,
    const void                       *beta,
    const cudnnTensorDescriptor_t    yDesc,
    void                             *y);
```

This function performs the forward LRN layer computation.

> 💬 Supported formats are: positive-strided, NCHW for 4D x and y, and only NCDHW DHW-packed for 5D (for both x and y). Only non-overlapping 4D and 5D tensors are supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
normDesc	input	Handle to a previously intialized LRN parameter descriptor.
lrnMode	input	LRN layer mode of operation. Currently only CUDNN_LRN_CROSS_CHANNEL_DIM1 is implemented. Normalization is performed along the tensor's dimA[1].
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the layer output value with prior value in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc, yDesc	input	Tensor descriptor objects for the input and output tensors.
x	input	Input tensor data pointer in device memory.
y	output	Output tensor data pointer in device memory.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The computation was performed successfully.

Return Value	Meaning
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the tensor pointers **x**, **y** is NULL. ▸ Number of input tensor dimensions is 2 or less. ▸ LRN descriptor parameters are outside of their valid ranges. ▸ One of tensor parameters is 5D but is not in NCDHW DHW-packed format.
`CUDNN_STATUS_NOT_SUPPORTED`	At least one of the following conditions are met: ▸ Any of the input tensor datatypes is not the same as any of the output tensor datatype. ▸ x and y tensor dimensions mismatch. ▸ Any tensor parameters strides are negative.

4.92. cudnnLRNCrossChannelBackward

```
cudnnStatus_t CUDNNWINAPI cudnnLRNCrossChannelBackward(
    cudnnHandle_t                 handle,
    cudnnLRNDescriptor_t          normDesc,
    cudnnLRNMode_t                lrnMode,
    const void*                   alpha,
    const cudnnTensorDescriptor_t  yDesc,
    const void                    *y,
    const cudnnTensorDescriptor_t  dyDesc,
    const void                    *dy,
    const cudnnTensorDescriptor_t  xDesc,
    const void                    *x,
    const void                    *beta,
    const cudnnTensorDescriptor_t  dxDesc,
    void                          *dx);
```

This function performs the backward LRN layer computation.

💬 Supported formats are: positive-strided, NCHW for 4D x and y, and only NCDHW DHW-packed for 5D (for both x and y). Only non-overlapping 4D and 5D tensors are supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
normDesc	input	Handle to a previously intialized LRN parameter descriptor.
lrnMode	input	LRN layer mode of operation. Currently only CUDNN_LRN_CROSS_CHANNEL_DIM1 is implemented. Normalization is performed along the tensor's dimA[1].
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the layer output value with prior value in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.

Param	In/out	Meaning
yDesc, y	input	Tensor descriptor and pointer in device memory for the layer's y data.
dyDesc, dy	input	Tensor descriptor and pointer in device memory for the layer's input cumulative loss differential data dy (including error backpropagation).
xDesc, x	input	Tensor descriptor and pointer in device memory for the layer's x data. Note that these values are not modified during backpropagation.
dxDesc, dx	output	Tensor descriptor and pointer in device memory for the layer's resulting cumulative loss differential data dx (including error backpropagation).

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The computation was performed successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ One of the tensor pointers **x**, **y** is NULL. ▸ Number of input tensor dimensions is 2 or less. ▸ LRN descriptor parameters are outside of their valid ranges. ▸ One of tensor parameters is 5D but is not in NCDHW DHW-packed format.
CUDNN_STATUS_NOT_SUPPORTED	At least one of the following conditions are met: ▸ Any of the input tensor datatypes is not the same as any of the output tensor datatype. ▸ Any pairwise tensor dimensions mismatch for x,y,dx,dy. ▸ Any tensor parameters strides are negative.

4.93. cudnnDivisiveNormalizationForward

```
cudnnStatus_t CUDNNWINAPI cudnnDivisiveNormalizationForward(
    cudnnHandle_t                    handle,
    cudnnLRNDescriptor_t             normDesc,
    cudnnDivNormMode_t               mode,
    const void                      *alpha,
    const cudnnTensorDescriptor_t    xDesc,
    const void                      *x,
    const void                      *means,
    void                            *temp,
    void                            *temp2,
    const void                      *beta,
    const cudnnTensorDescriptor_t    yDesc,
    void                            *y );
```

This function performs the forward spatial DivisiveNormalization layer computation. It divides every value in a layer by the standard deviation of it's spatial neighbors as

described in *"What is the Best Multi-Stage Architecture for Object Recognition"*, Jarrett 2009, Local Contrast Normalization Layer section. Note that Divisive Normalization only implements the x/max(c, sigma_x) portion of the computation, where sigma_x is the variance over the spatial neighborhood of x. The full LCN (Local Contrastive Normalization) computation can be implemented as a two-step process:

x_m = x-mean(x);

y = x_m/max(c, sigma(x_m));

The "x-mean(x)" which is often referred to as "subtractive normalization" portion of the computation can be implemented using cuDNN average pooling layer followed by a call to addTensor.

> Supported tensor formats are NCHW for 4D and NCDHW for 5D with any non-overlapping non-negative strides. Only 4D and 5D tensors are supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
normDesc	input	Handle to a previously intialized LRN parameter descriptor. This descriptor is used for both LRN and DivisiveNormalization layers.
divNormMode	input	DivisiveNormalization layer mode of operation. Currently only CUDNN_DIVNORM_PRECOMPUTED_MEANS is implemented. Normalization is performed using the means input tensor that is expected to be precomputed by the user.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the layer output value with prior value in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc, yDesc	input	Tensor descriptor objects for the input and output tensors. Note that xDesc is shared between x, means, temp and temp2 tensors.
x	input	Input tensor data pointer in device memory.
means	input	Input means tensor data pointer in device memory. Note that this tensor can be NULL (in that case it's values are assumed to be zero during the computation). This tensor also doesn't have to contain means, these can be any values, a frequently used variation is a result of convolution with a normalized positive kernel (such as Gaussian).
temp, temp2	workspace	Temporary tensors in device memory. These are used for computing intermediate values during the forward pass. These tensors do not have to be preserved as inputs from forward to the backward pass. Both use xDesc as their descriptor.
y	output	Pointer in device memory to a tensor for the result of the forward DivisiveNormalization computation.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The computation was performed successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▶ One of the tensor pointers `x`, `y`, `temp`, `temp2` is NULL. ▶ Number of input tensor or output tensor dimensions is outside of [4,5] range. ▶ A mismatch in dimensions between any two of the input or output tensors. ▶ For in-place computation when pointers x == y, a mismatch in strides between the input data and output data tensors. ▶ Alpha or beta pointer is NULL. ▶ LRN descriptor parameters are outside of their valid ranges. ▶ Any of the tensor strides are negative.
`CUDNN_STATUS_UNSUPPORTED`	At least one of the following conditions are met: ▶ Any of the input and output tensor strides mismatch (for the same dimension).

4.94. cudnnDivisiveNormalizationBackward

```
 cudnnStatus_t
CUDNNWINAPI cudnnDivisiveNormalizationBackward(
    cudnnHandle_t                   handle,
    cudnnLRNDescriptor_t            normDesc,
    cudnnDivNormMode_t              mode,
    const void                     *alpha,
    const cudnnTensorDescriptor_t   xDesc,
    const void                     *x,
    const void                     *means,
    const void                     *dy,
    void                           *temp,
    void                           *temp2,
    const void                     *beta,
    const cudnnTensorDescriptor_t   dxDesc,
    void                           *dx,
    void                           *dMeans );
```

This function performs the backward DivisiveNormalization layer computation.

> Supported tensor formats are NCHW for 4D and NCDHW for 5D with any non-overlapping non-negative strides. Only 4D and 5D tensors are supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
normDesc	input	Handle to a previously intialized LRN parameter descriptor (this descriptor is used for both LRN and DivisiveNormalization layers).

Param	In/out	Meaning
mode	input	DivisiveNormalization layer mode of operation. Currently only CUDNN_DIVNORM_PRECOMPUTED_MEANS is implemented. Normalization is performed using the means input tensor that is expected to be precomputed by the user.
alpha, beta	input	Pointers to scaling factors (in host memory) used to blend the layer output value with prior value in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc, x, means	input	Tensor descriptor and pointers in device memory for the layer's x and means data. Note: the means tensor is expected to be precomputed by the user. It can also contain any valid values (not required to be actual means, and can be for instance a result of a convolution with a Gaussian kernel).
dy	input	Tensor pointer in device memory for the layer's dy cumulative loss differential data (error backpropagation).
temp, temp2	workspace	Temporary tensors in device memory. These are used for computing intermediate values during the backward pass. These tensors do not have to be preserved from forward to backward pass. Both use xDesc as a descriptor.
dxDesc	input	Tensor descriptor for dx and dMeans.
dx, dMeans	output	Tensor pointers (in device memory) for the layer's resulting cumulative gradients dx and dMeans (dLoss/dx and dLoss/dMeans). Both share the same descriptor.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The computation was performed successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the tensor pointers `x, dx, temp, tmep2, dy` is NULL. ▸ Number of any of the input or output tensor dimensions is not within the [4,5] range. ▸ Either alpha or beta pointer is NULL. ▸ A mismatch in dimensions between xDesc and dxDesc. ▸ LRN descriptor parameters are outside of their valid ranges. ▸ Any of the tensor strides is negative.
`CUDNN_STATUS_UNSUPPORTED`	At least one of the following conditions are met: ▸ Any of the input and output tensor strides mismatch (for the same dimension).

4.95. cudnnBatchNormalizationForwardInference

```
cudnnStatus_t CUDNNWINAPI cudnnBatchNormalizationForwardInference(
    cudnnHandle_t                    handle,
    cudnnBatchNormMode_t             mode,
    const void                      *alpha,
    const void                      *beta,
    const cudnnTensorDescriptor_t    xDesc,
    const void                      *x,
    const cudnnTensorDescriptor_t    yDesc,
    void                            *y,
    const cudnnTensorDescriptor_t    bnScaleBiasMeanVarDesc,
    const void                      *bnScale,
    const void                      *bnBias,
    const void                      *estimatedMean,
    const void                      *estimatedVariance,
    double                           epsilon );
```

This function performs the forward BatchNormalization layer computation for inference phase. This layer is based on the paper "*Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift*", S. Ioffe, C. Szegedy, 2015.

Only 4D and 5D tensors are supported.

The input transformation performed by this function is defined as: y := alpha*y + beta *(bnScale * (x-estimatedMean)/sqrt(epsilon + estimatedVariance)+bnBias)

The epsilon value has to be the same during training, backpropagation and inference.

For training phase use cudnnBatchNormalizationForwardTraining.

Much higher performance when HW-packed tensors are used for all of x, dy, dx.

Param	Meaning
handle	Input. Handle to a previously created cuDNN library descriptor.
mode	Input. Mode of operation (spatial or per-activation). cudnnBatchNormMode_t
alpha, beta	Inputs. Pointers to scaling factors (in host memory) used to blend the layer output value with prior value in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc, yDesc, x, y	Tensor descriptors and pointers in device memory for the layer's x and y data.

Param	Meaning
bnScaleBiasMeanVarDesc, bnScaleData, bnBiasData	Inputs. Tensor descriptor and pointers in device memory for the batch normalization scale and bias parameters (in the original paper bias is referred to as beta and scale as gamma).
estimatedMean, estimatedVariance	Inputs. Mean and variance tensors (these have the same descriptor as the bias and scale). It is suggested that resultRunningMean, resultRunningVariance from the cudnnBatchNormalizationForwardTraining call accumulated during the training phase are passed as inputs here.
epsilon	Input. Epsilon value used in the batch normalization formula. Minimum allowed value is CUDNN_BN_MIN_EPSILON defined in cudnn.h.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The computation was performed successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the pointers `alpha`, `beta`, `x`, `y`, `bnScaleData`, `bnBiasData`, `estimatedMean`, `estimatedInvVariance` is NULL. ▸ Number of xDesc or yDesc tensor descriptor dimensions is not within the [4,5] range. ▸ bnScaleBiasMeanVarDesc dimensions are not 1xC(x1)x1x1 for spatial or 1xC(xD)xHxW for per-activation mode (parenthesis for 5D). ▸ epsilon value is less than CUDNN_BN_MIN_EPSILON ▸ Dimensions or data types mismatch for xDesc, yDesc

4.96. cudnnBatchNormalizationForwardTraining

```
cudnnStatus_t CUDNNWINAPI cudnnBatchNormalizationForwardTraining(
    cudnnHandle_t                    handle,
    cudnnBatchNormMode_t             mode,
    const void                      *alpha,
    const void                      *beta,
    const cudnnTensorDescriptor_t    xDesc,
    const void                      *x,
    const cudnnTensorDescriptor_t    yDesc,
    void                            *y,
    const cudnnTensorDescriptor_t    bnScaleBiasMeanVarDesc,
    const void                      *bnScale,
    const void                      *bnBias,
    double                           exponentialAverageFactor,
    void                            *resultRunningMean,
    void                            *resultRunningInvVariance,
    double                           epsilon,
    void                            *resultSaveMean,
    void                            *resultSaveVariance );
```

This function performs the forward BatchNormalization layer computation for training phase.

> Only 4D and 5D tensors are supported.

> The epsilon value has to be the same during training, backpropagation and inference.

> For inference phase use cudnnBatchNormalizationForwardInference.

> Much higher performance for HW-packed tensors for both x and y.

Param	Meaning
handle	Handle to a previously created cuDNN library descriptor.
mode	Mode of operation (spatial or per-activation). cudnnBatchNormMode_t
alpha, beta	Inputs. Pointers to scaling factors (in host memory) used to blend the layer output value with prior value in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc, yDesc, x, y	Tensor descriptors and pointers in device memory for the layer's x and y data.
bnScaleBiasMeanVarDesc	Shared tensor descriptor desc for all the 6 tensors below in the argument list. The dimensions for this tensor descriptor are dependent on the normalization mode.
bnScale, bnBias	Inputs. Pointers in device memory for the batch normalization scale and bias parameters (in original paper bias is referred to as beta and scale as gamma). Note that bnBias parameter can replace the previous layer's bias parameter for improved efficiency.
exponentialAverageFactor	Input. Factor used in the moving average computation runningMean = newMean*factor + runningMean*(1-factor). Use a factor=1/(1+n) at N-th call to the function to get Cumulative Moving Average (CMA) behavior CMA[n] = (x[1]+...+x[n])/n. Since CMA[n+1] = (n*CMA[n]+x[n+1])/(n+1)= ((n+1)*CMA[n]-CMA[n])/(n+1) + x[n+1]/(n+1) = CMA[n]*(1-1/(n+1))+x[n+1]*1/(n+1)
resultRunningMean, resultRunningVariance	Inputs/outputs. Running mean and variance tensors (these have the same descriptor as the bias and scale). Both of these pointers can be NULL but only at the same time. The value stored in resultRunningVariance (or passed as an input in inference mode) is the moving average of variance[x] where variance is computed either over batch or spatial +batch dimensions depending on the mode. If these pointers are not NULL, the tensors should be initialized to some reasonable values or to 0.
epsilon	Epsilon value used in the batch normalization formula. Minimum allowed value is CUDNN_BN_MIN_EPSILON defined in cudnn.h. Same epsilon value should be used in forward and backward functions.

Param	Meaning
resultSaveMean, resultSaveInvVariance	Outputs. Optional cache to save intermediate results computed during the forward pass - these can then be reused to speed up the backward pass. For this to work correctly, the bottom layer data has to remain unchanged until the backward function is called. Note that both of these parameters can be NULL but only at the same time. It is recommended to use this cache since memory overhead is relatively small because these tensors have a much lower product of dimensions than the data tensors.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The computation was performed successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ One of the pointers `alpha`, `beta`, `x`, `y`, `bnScaleData`, `bnBiasData` is NULL. ▸ Number of xDesc or yDesc tensor descriptor dimensions is not within the [4,5] range. ▸ bnScaleBiasMeanVarDesc dimensions are not 1xC(x1)x1x1 for spatial or 1xC(xD)xHxW for per-activation mode (parens for 5D). ▸ Exactly one of resultSaveMean, resultSaveInvVariance pointers is NULL. ▸ Exactly one of resultRunningMean, resultRunningInvVariance pointers is NULL. ▸ epsilon value is less than CUDNN_BN_MIN_EPSILON ▸ Dimensions or data types mismatch for xDesc, yDesc

4.97. cudnnBatchNormalizationBackward

```
cudnnStatus_t CUDNNWINAPI cudnnBatchNormalizationBackward(
        cudnnHandle_t                    handle,
        cudnnBatchNormMode_t             mode,
        const void                      *alphaDataDiff,
        const void                      *betaDataDiff,
        const void                      *alphaParamDiff,
        const void                      *betaParamDiff,
        const cudnnTensorDescriptor_t    xDesc,
        const void                      *x,
        const cudnnTensorDescriptor_t    dyDesc,
        const void                      *dy,
        const cudnnTensorDescriptor_t    dxDesc,
        void                            *dx,
        const cudnnTensorDescriptor_t    bnScaleBiasDiffDesc,
        const void                      *bnScale,
        void                            *resultBnScaleDiff,
        void                            *resultBnBiasDiff,
        double                           epsilon,
        const void                      *savedMean,
        const void                      *savedInvVariance
        );
```

This function performs the backward BatchNormalization layer computation.

> Only 4D and 5D tensors are supported.

> The epsilon value has to be the same during training, backpropagation and inference.

> Much higher performance when HW-packed tensors are used for all of x, dy, dx.

Param	Meaning
handle	Handle to a previously created cuDNN library descriptor.
mode	Mode of operation (spatial or per-activation). cudnnBatchNormMode_t
alphaDataDiff, betaDataDiff	Inputs. Pointers to scaling factors (in host memory) used to blend the gradient output dx with a prior value in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.
alphaParamDiff, betaParamDiff	Inputs. Pointers to scaling factors (in host memory) used to blend the gradient outputs dBnScaleResult and dBnBiasResult with prior values in the destination tensor as follows: dstValue = alpha[0]*resultValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc, x, dyDesc, dy, dxDesc, dx	Tensor descriptors and pointers in device memory for the layer's x data, backpropagated differential dy (inputs) and resulting differential with respect to x, dx (output).

Param	Meaning
bnScaleBiasDiffDesc	Shared tensor descriptor for all the 5 tensors below in the argument list (bnScale, resultBnScaleDiff, resultBnBiasDiff, savedMean, savedInvVariance). The dimensions for this tensor descriptor are dependent on normalization mode. Note: The data type of this tensor descriptor must be 'float' for FP16 and FP32 input tensors, and 'double' for FP64 input tensors.
bnScale	Input. Pointers in device memory for the batch normalization scale parameter (in original paper bias is referred to as gamma). Note that bnBias parameter is not needed for this layer's computation.
resultBnScaleDiff, resultBnBiasDiff	Outputs. Pointers in device memory for the resulting scale and bias differentials computed by this routine. Note that scale and bias gradients are not backpropagated below this layer (since they are dead-end computation DAG nodes).
epsilon	Epsilon value used in batch normalization formula. Minimum allowed value is CUDNN_BN_MIN_EPSILON defined in cudnn.h. Same epsilon value should be used in forward and backward functions.
savedMean, savedInvVariance	Inputs. Optional cache parameters containing saved intermediate results computed during the forward pass. For this to work correctly, the layer's x and bnScale, bnBias data has to remain unchanged until the backward function is called. Note that both of these parameters can be NULL but only at the same time. It is recommended to use this cache since the memory overhead is relatively small.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The computation was performed successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ Any of the pointers `alpha`, `beta`, `x`, `dy`, `dx`, `bnScale`, `resultBnScaleDiff`, `resultBnBiasDiff` is NULL. ▸ Number of xDesc or yDesc or dxDesc tensor descriptor dimensions is not within the [4,5] range. ▸ bnScaleBiasMeanVarDesc dimensions are not 1xC(x1)x1x1 for spatial or 1xC(xD)xHxW for per-activation mode (parentheses for 5D). ▸ Exactly one of savedMean, savedInvVariance pointers is NULL. ▸ epsilon value is less than CUDNN_BN_MIN_EPSILON ▸ Dimensions or data types mismatch for any pair of xDesc, dyDesc, dxDesc

4.98. cudnnDeriveBNTensorDescriptor

```
cudnnStatus_t CUDNNWINAPI cudnnDeriveBNTensorDescriptor(
      cudnnTensorDescriptor_t derivedBnDesc,
      const cudnnTensorDescriptor_t xDesc,
      cudnnBatchNormMode_t mode);
```

Derives a secondary tensor descriptor for BatchNormalization scale, invVariance, bnBias, bnScale subtensors from the layer's x data descriptor. Use the tensor descriptor produced by this function as the bnScaleBiasMeanVarDesc and bnScaleBiasDiffDesc parameters in Spatial and Per-Activation Batch Normalization forward and backward functions. Resulting dimensions will be **1xC(x1)x1x1** for BATCHNORM_MODE_SPATIAL and **1xC(xD)xHxW** for BATCHNORM_MODE_PER_ACTIVATION (parentheses for 5D). For HALF input data type the resulting tensor descriptor will have a FLOAT type. For other data types it will have the same type as the input data.

> Only 4D and 5D tensors are supported.

> derivedBnDesc has to be first created using cudnnCreateTensorDescriptor

> xDesc is the descriptor for the layer's x data and has to be setup with proper dimensions prior to calling this function.

Param	In/out	Meaning
derivedBnDesc	output	Handle to a previously created tensor descriptor.
xDesc	input	Handle to a previously created and initialized layer's x data descriptor.
mode	input	Batch normalization layer mode of operation.

Possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The computation was performed successfully.
`CUDNN_STATUS_BAD_PARAM`	Invalid Batch Normalization mode.

4.99. cudnnCreateRNNDescriptor

```
cudnnStatus_t cudnnCreateRNNDescriptor(cudnnRNNDescriptor_t * rnnDesc)
```

This function creates a generic RNN descriptor object by allocating the memory needed to hold its opaque structure.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was created successfully.
`CUDNN_STATUS_ALLOC_FAILED`	The resources could not be allocated.

4.100. cudnnDestroyRNNDescriptor

```
cudnnStatus_t cudnnDestroyRNNDescriptor(cudnnRNNDescriptor_t rnnDesc)
```

This function destroys a previously created RNN descriptor object.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was destroyed successfully.

4.101. cudnnSetRNNDescriptor

```
cudnnStatus_t
cudnnSetRNNDescriptor( cudnnRNNDescriptor_t rnnDesc,
                       int hiddenSize,
                       int numLayers,
                       cudnnDropoutDescriptor_t dropoutDesc,
                       cudnnRNNInputMode_t inputMode,
                       cudnnDirectionMode_t direction,
                       cudnnRNNMode_t mode,
                       cudnnDataType_t dataType )
```

This function initializes a previously created RNN descriptor object.

> Larger networks (eg. longer sequences, more layers) are expected to be more efficient than smaller networks.

Param	In/out	Meaning
rnnDesc	input/output	A previously created RNN descriptor.
hiddenSize	input	Size of the internal hidden state for each layer.
numLayers	input	Number of stacked layers.
dropoutDesc	input	Handle to a previously created and initialized dropout descriptor.
inputMode	input	Specifies the behavior at the input to the first layer
direction	input	Specifies the recurrence pattern. (eg. bidirectional)
mode	input	The type of RNN to compute.
dataType	input	Math precision.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was set successfully.
`CUDNN_STATUS_BAD_PARAM`	Either at least one of the parameters `hiddenSize, numLayers` was zero or negative, one of `inputMode, direction, mode, dataType` has an invalid enumerant value, `dropoutDesc` is an invalid dropout descriptor or `rnnDesc` has not been created correctly.

4.102. cudnnGetRNNWorkspaceSize

```
cudnnStatus_t
cudnnGetRNNWorkspaceSize( cudnnHandle_t              handle,
                          const cudnnRNNDescriptor_t    rnnDesc,
                          const int seqLength,
                          const cudnnTensorDescriptor_t *xDesc,
                          size_t                        *sizeInBytes)
```

This function is used to query the amount of work space required to execute the RNN described by **rnnDesc** with inputs dimensions defined by **xDesc**.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
rnnDesc	input	A previously initialized RNN descriptor.
seqLength	input	Number of iterations to unroll over.
xDesc	input	An array of tensor descriptors describing the input to each recurrent iteration. Each tensor descriptor must have the same second dimension. The first dimension of the tensors may decrease from element **n** to element **n+1** but may not increase.
sizeInBytes	output	Minimum amount of GPU memory needed as workspace to be able to execute an RNN with the specified descriptor and input tensors.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The descriptor **rnnDesc** is invalid. ▸ At least one of the descriptors in **xDesc** is invalid. ▸ The descriptors in **xDesc** have inconsistent second dimensions, strides or data types. ▸ The descriptors in **xDesc** have increasing first dimensions. ▸ The descriptors in **xDesc** is not fully packed.
CUDNN_STATUS_NOT_SUPPORTED	The the data types in tensors described by xDesc is not supported.

4.103. cudnnGetRNNTrainingReserveSize

```
cudnnStatus_t
cudnnGetRNNTrainingReserveSize( cudnnHandle_t               handle,
                                const cudnnRNNDescriptor_t   rnnDesc,
                                const int seqLength,
                                const cudnnTensorDescriptor_t *xDesc,
                                size_t                       *sizeInBytes)
```

This function is used to query the amount of reserved space required for training the RNN described by **rnnDesc** with inputs dimensions defined by **xDesc**. The same reserve space must be passed to **cudnnRNNForwardTraining**, **cudnnRNNBackwardData** and **cudnnRNNBackwardWeights**. Each of these calls overwrites the contents of the reserve space, however it can safely be copied if reuse is required.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
rnnDesc	input	A previously initialized RNN descriptor.
seqLength	input	Number of iterations to unroll over.
xDesc	input	An array of tensor descriptors describing the input to each recurrent iteration. Each tensor descriptor must have the same second dimension. The first dimension of the tensors may decrease from element **n** to element **n+1** but may not increase.
sizeInBytes	output	Minimum amount of GPU memory needed as reserve space to be able to train an RNN with the specified descriptor and input tensors.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The descriptor **rnnDesc** is invalid. ▸ At least one of the descriptors in **xDesc** is invalid. ▸ The descriptors in **xDesc** have inconsistent second dimensions, strides or data types. ▸ The descriptors in **xDesc** have increasing first dimensions. ▸ The descriptors in **xDesc** is not fully packed.
CUDNN_STATUS_NOT_SUPPORTED	The the data types in tensors described by xDesc is not supported.

4.104. cudnnGetRNNParamsSize

```
cudnnStatus_t
cudnnGetRNNParamsSize( cudnnHandle_t                 handle,
                       const cudnnRNNDescriptor_t    rnnDesc,
                       const cudnnTensorDescriptor_t xDesc,
                       size_t                        *sizeInBytes,
                       cudnnDataType_t dataType)
```

This function is used to query the amount of parameter space required to execute the RNN described by **rnnDesc** with inputs dimensions defined by **xDesc**.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
rnnDesc	input	A previously initialized RNN descriptor.
xDesc	input	A fully packed tensor descriptor describing the input to one recurrent iteration.
sizeInBytes	output	Minimum amount of GPU memory needed as parameter space to be able to execute an RNN with the specified descriptor and input tensors.
dataType	input	The data type of the parameters.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The descriptor **rnnDesc** is invalid. ▸ The descriptor **xDesc** is invalid. ▸ The descriptor **xDesc** is not fully packed. ▸ The combination of **dataType** and tensor descriptor data type is invalid.
CUDNN_STATUS_NOT_SUPPORTED	The combination of the RNN descriptor and tensor descriptors is not supported.

4.105. cudnnGetRNNLinLayerMatrixParams

```
cudnnStatus_t
cudnnGetRNNLinLayerMatrixParams( cudnnHandle_t                 handle,
                                 const cudnnRNNDescriptor_t rnnDesc,
                                 const int layer,
                                 const cudnnTensorDescriptor_t xDesc,
                                 const cudnnFilterDescriptor_t wDesc,
                                 const void * w,
                                 const int linLayerID,
                                 cudnnFilterDescriptor_t linLayerMatDesc,
                                 void ** linLayerMat)
```

This function is used to obtain a pointer and descriptor for the matrix parameters in **layer** within the RNN described by **rnnDesc** with inputs dimensions defined by **xDesc**.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
rnnDesc	input	A previously initialized RNN descriptor.
layer	input	The layer to query.
xDesc	input	A fully packed tensor descriptor describing the input to one recurrent iteration.
wDesc	input	Handle to a previously initialized filter descriptor describing the weights for the RNN.
w	input	Data pointer to GPU memory associated with the filter descriptor **wDesc**.

linLayerID	input	The linear layer to obtain information about:
		
		 ▸ If **mode** in **rnnDesc** was set to **CUDNN_RNN_RELU** or **CUDNN_RNN_TANH** a value of 0 references the matrix multiplication applied to the input from the previous layer, a value of 1 references the matrix multiplication applied to the recurrent input.
		 ▸ If **mode** in **rnnDesc** was set to **CUDNN_LSTM** values of 0-3 reference matrix multiplications applied to the input from the previous layer, value of 4-7 reference matrix multiplications applied to the recurrent input.
		
		     ▸ Values 0 and 4 reference the input gate.
		     ▸ Values 1 and 5 reference the forget gate.
		     ▸ Values 2 and 6 reference the new memory gate.
		     ▸ Values 3 and 7 reference the output gate.
		 ▸ If **mode** in **rnnDesc** was set to **CUDNN_GRU** values of 0-2 reference matrix multiplications applied to the input from the previous layer, value of 3-5 reference matrix multiplications applied to the recurrent input.
		
		     ▸ Values 0 and 3 reference the reset gate.
		     ▸ Values 1 and 4 reference the update gate.
		     ▸ Values 2 and 5 reference the new memory gate.
		
		 Please refer to this section for additional details on modes.

linLayerMatDesc	output	Handle to a previously created filter descriptor.
linLayerMat	output	Data pointer to GPU memory associated with the filter descriptor **linLayerMatDesc**.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The query was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met:

Return Value	Meaning
	▸ The descriptor `rnnDesc` is invalid.
	▸ One of the descriptors `xDesc`, `wDesc`, `linLayerMatDesc` is invalid.
	▸ One of `layer`, `linLayerID` is invalid.

4.106. cudnnGetRNNLinLayerBiasParams

```
cudnnStatus_t
cudnnGetRNNLinLayerBiasParams( cudnnHandle_t            handle,
                               const cudnnRNNDescriptor_t rnnDesc,
                               const int layer,
                               const cudnnTensorDescriptor_t xDesc,
                               const cudnnFilterDescriptor_t wDesc,
                               const void * w,
                               const int linLayerID,
                               cudnnFilterDescriptor_t linLayerBiasDesc,
                               void ** linLayerBias
```

This function is used to obtain a pointer and descriptor for the bias parameters in **layer** within the RNN described by **rnnDesc** with inputs dimensions defined by **xDesc**.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN library descriptor.
rnnDesc	input	A previously initialized RNN descriptor.
layer	input	The layer to query.
xDesc	input	A fully packed tensor descriptor describing the input to one recurrent iteration.
wDesc	input	Handle to a previously initialized filter descriptor describing the weights for the RNN.
w	input	Data pointer to GPU memory associated with the filter descriptor `wDesc`.
linLayerID	input	The linear layer to obtain information about: ▸ If `mode` in `rnnDesc` was set to `CUDNN_RNN_RELU` or `CUDNN_RNN_TANH` a value of 0 references the bias applied to the input from the previous layer, a value of 1 references the bias applied to the recurrent input. ▸ If `mode` in `rnnDesc` was set to `CUDNN_LSTM` values of 0, 1, 2 and 3 reference bias applied to the input from the previous layer, value of 4, 5, 6 and 7 reference bias applied to the recurrent input. 　▸ Values 0 and 4 reference the input gate. 　▸ Values 1 and 5 reference the forget gate. 　▸ Values 2 and 6 reference the new memory gate. 　▸ Values 3 and 7 reference the output gate. ▸ If `mode` in `rnnDesc` was set to `CUDNN_GRU` values of 0, 1 and 2 reference bias applied to the input from the previous layer, value of 3, 4 and 5 reference bias applied to the recurrent input.

Param	In/out	Meaning
		▸ Values 0 and 3 reference the reset gate. ▸ Values 1 and 4 reference the update gate. ▸ Values 2 and 5 reference the new memory gate. Please refer to this section for additional details on modes.
linLayerBiasDesc	output	Handle to a previously created filter descriptor.
linLayerBias	output	Data pointer to GPU memory associated with the filter descriptor `linLayerMatDesc`.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The descriptor `rnnDesc` is invalid. ▸ One of the descriptors `xDesc`, `wDesc`, `linLayerBiasDesc` is invalid. ▸ One of `layer`, `linLayerID` is invalid.

4.107. cudnnRNNForwardInference

```
cudnnStatus_t
cudnnRNNForwardInference( cudnnHandle_t handle,
                          const cudnnRNNDescriptor_t rnnDesc,
                          const int seqLength,
                          const cudnnTensorDescriptor_t * xDesc,
                          const void * x,
                          const cudnnTensorDescriptor_t hxDesc,
                          const void * hx,
                          const cudnnTensorDescriptor_t cxDesc,
                          const void * cx,
                          const cudnnFilterDescriptor_t wDesc,
                          const void * w,
                          const cudnnTensorDescriptor_t *yDesc,
                          void * y,
                          const cudnnTensorDescriptor_t hyDesc,
                          void * hy,
                          const cudnnTensorDescriptor_t cyDesc,
                          void * cy,
                          void * workspace,
                          size_t workSpaceSizeInBytes)
```

This routine executes the recurrent neural network described by `rnnDesc` with inputs `x, hx, cx`, weights `w` and outputs `y, hy, cy`. `workspace` is required for intermediate storage. This function does not store data required for training; `cudnnRNNForwardTraining` should be used for that purpose.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.

Param	In/out	Meaning
rnnDesc	input	A previously initialized RNN descriptor.
seqLength	input	Number of iterations to unroll over.
xDesc	input	An array of fully packed tensor descriptors describing the input to each recurrent iteration. Each tensor descriptor must have the same second dimension. The first dimension of the tensors may decrease from element `n` to element `n+1` but may not increase.
x	input	Data pointer to GPU memory associated with the tensor descriptors in the array `xDesc`. The data are expected to be packed contiguously with the first element of iteration n+1 following directly from the last element of iteration n.
hxDesc	input	A fully packed tensor descriptor describing the initial hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
hx	input	Data pointer to GPU memory associated with the tensor descriptor `hxDesc`. If a NULL pointer is passed, the initial hidden state of the network will be initialized to zero.
cxDesc	input	A fully packed tensor descriptor describing the initial cell state for LSTM networks. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
cx	input	Data pointer to GPU memory associated with the tensor descriptor `cxDesc`. If a NULL pointer is passed, the initial cell state of the network will be initialized to zero.
wDesc	input	Handle to a previously initialized filter descriptor describing the weights for the RNN.

Param	In/out	Meaning
w	input	Data pointer to GPU memory associated with the filter descriptor `wDesc`.
yDesc	input	An array of fully packed tensor descriptors describing the output from each recurrent iteration. The second dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the second dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the second dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The first dimension of the tensor `n` must match the first dimension of the tensor `n` in `xDesc`.
y	output	Data pointer to GPU memory associated with the output tensor descriptor `yDesc`. The data are expected to be packed contiguously with the first element of iteration n+1 following directly from the last element of iteration n.
hyDesc	input	A fully packed tensor descriptor describing the final hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
hy	output	Data pointer to GPU memory associated with the tensor descriptor `hyDesc`. If a NULL pointer is passed, the final hidden state of the network will not be saved.
cyDesc	input	A fully packed tensor descriptor describing the final cell state for LSTM networks. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the

Param	In/out	Meaning
		numLayers argument passed to the **cudnnSetRNNDescriptor** call used to initialize **rnnDesc**. The tensor must be fully packed.
cy	output	Data pointer to GPU memory associated with the tensor descriptor **cyDesc**. If a NULL pointer is passed, the final cell state of the network will be not be saved.
workspace	input	Data pointer to GPU memory to be used as a workspace for this call.
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided **workspace**

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The function launched successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The descriptor **rnnDesc** is invalid. ▸ At least one of the descriptors **hxDesc**, **cxDesc**, **wDesc**, **hyDesc**, **cyDesc** or one of the descriptors in **xDesc**, **yDesc** is invalid. ▸ The descriptors in one of **xDesc**, **hxDesc**, **cxDesc**, **wDesc**, **yDesc**, **hyDesc**, **cyDesc** have incorrect strides or dimensions. ▸ **workSpaceSizeInBytes** is too small.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.
CUDNN_STATUS_ALLOC_FAILED	The function was unable to allocate memory.

4.108. cudnnRNNForwardTraining

```
cudnnStatus_t
cudnnRNNForwardTraining( cudnnHandle_t handle,
                         const cudnnRNNDescriptor_t rnnDesc,
                         const int seqLength,
                         const cudnnTensorDescriptor_t *xDesc,
                         const void * x,
                         const cudnnTensorDescriptor_t hxDesc,
                         const void * hx,
                         const cudnnTensorDescriptor_t cxDesc,
                         const void * cx,
                         const cudnnFilterDescriptor_t wDesc,
                         const void * w,
                         const cudnnTensorDescriptor_t *yDesc,
                         void * y,
                         const cudnnTensorDescriptor_t hyDesc,
                         void * hy,
                         const cudnnTensorDescriptor_t cyDesc,
                         void * cy,
                         void * workspace,
                         size_t workSpaceSizeInBytes,
                         void * reserveSpace,
                         size_t reserveSpaceSizeInBytes)
```

This routine executes the recurrent neural network described by `rnnDesc` with inputs `x, hx, cx`, weights `w` and outputs `y, hy, cy. workspace` is required for intermediate storage. `reserveSpace` stores data required for training. The same reserveSpace data must be used for future calls to `cudnnRNNBackwardData` and `cudnnRNNBackwardWeights` if these execute on the same input data.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
rnnDesc	input	A previously initialized RNN descriptor.
xDesc	input	An array of fully packed tensor descriptors describing the input to each recurrent iteration. Each tensor descriptor must have the same second dimension. The first dimension of the tensors may decrease from element `n` to element `n+1` but may not increase.
seqLength	input	Number of iterations to unroll over.
x	input	Data pointer to GPU memory associated with the tensor descriptors in the array `xDesc`.
hxDesc	input	A fully packed tensor descriptor describing the initial hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▪ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▪ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
hx	input	Data pointer to GPU memory associated with the tensor descriptor `hxDesc`. If a NULL pointer is passed, the initial hidden state of the network will be initialized to zero.
cxDesc	input	A fully packed tensor descriptor describing the initial cell state for LSTM networks. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▪ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▪ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.

Param	In/out	Meaning
cx	input	Data pointer to GPU memory associated with the tensor descriptor `cxDesc`. If a NULL pointer is passed, the initial cell state of the network will be initialized to zero.
wDesc	input	Handle to a previously initialized filter descriptor describing the weights for the RNN.
w	input	Data pointer to GPU memory associated with the filter descriptor `wDesc`.
yDesc	input	An array of fully packed tensor descriptors describing the output from each recurrent iteration. The second dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the second dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the second dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The first dimension of the tensor `n` must match the first dimension of the tensor `n` in `xDesc`.
y	output	Data pointer to GPU memory associated with the output tensor descriptor `yDesc`.
hyDesc	input	A fully packed tensor descriptor describing the final hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
hy	output	Data pointer to GPU memory associated with the tensor descriptor `hyDesc`. If a NULL pointer is passed, the final hidden state of the network will not be saved.
cyDesc	input	A fully packed tensor descriptor describing the final cell state for LSTM networks. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`.

Param	In/out	Meaning
		▸ If **direction** is **CUDNN_BIDIRECTIONAL** the third dimension should match double the **hiddenSize** argument passed to **cudnnSetRNNDescriptor**.
		The second dimension must match the second dimension of the tensors described in **xDesc**. The first dimension must match the **numLayers** argument passed to the **cudnnSetRNNDescriptor** call used to initialize **rnnDesc**. The tensor must be fully packed.
cy	output	Data pointer to GPU memory associated with the tensor descriptor **cyDesc**. If a NULL pointer is passed, the final cell state of the network will be not be saved.
workspace	input	Data pointer to GPU memory to be used as a workspace for this call.
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided **workspace**
reserveSpace	input/ output	Data pointer to GPU memory to be used as a reserve space for this call.
reserveSpaceSizeInBytes	input	Specifies the size in bytes of the provided **reserveSpace**

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The function launched successfully.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The descriptor **rnnDesc** is invalid. ▸ At least one of the descriptors **hxDesc**, **cxDesc**, **wDesc**, **hyDesc**, **cyDesc** or one of the descriptors in **xDesc**, **yDesc** is invalid. ▸ The descriptors in one of **xDesc**, **hxDesc**, **cxDesc**, **wDesc**, **yDesc**, **hyDesc**, **cyDesc** have incorrect strides or dimensions. ▸ **workSpaceSizeInBytes** is too small. ▸ **reserveSpaceSizeInBytes** is too small.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.
CUDNN_STATUS_ALLOC_FAILED	The function was unable to allocate memory.

4.109. cudnnRNNBackwardData

```
cudnnStatus_t
cudnnRNNBackwardData( cudnnHandle_t handle,
                      const cudnnRNNDescriptor_t rnnDesc,
                      const int seqLength,
                      const cudnnTensorDescriptor_t * yDesc,
                      const void * y,
                      const cudnnTensorDescriptor_t * dyDesc,
                      const void * dy,
                      const cudnnTensorDescriptor_t dhyDesc,
                      const void * dhy,
                      const cudnnTensorDescriptor_t dcyDesc,
                      const void * dcy,
                      const cudnnFilterDescriptor_t wDesc,
                      const void * w,
                      const cudnnTensorDescriptor_t hxDesc,
                      const void * hx,
                      const cudnnTensorDescriptor_t cxDesc,
                      const void * cx,
                      const cudnnTensorDescriptor_t * dxDesc,
                      void * dx,
                      const cudnnTensorDescriptor_t dhxDesc,
                      void * dhx,
                      const cudnnTensorDescriptor_t dcxDesc,
                      void * dcx,
                      void * workspace,
                      size_t workSpaceSizeInBytes,
                      const void * reserveSpace,
                      size_t reserveSpaceSizeInBytes )
```

This routine executes the recurrent neural network described by **rnnDesc** with output gradients **dy, dhy, dhc**, weights **w** and input gradients **dx, dhx, dcx**. **workspace** is required for intermediate storage. The data in **reserveSpace** must have previously been generated by **cudnnRNNForwardTraining**. The same reserveSpace data must be used for future calls to **cudnnRNNBackwardWeights** if they execute on the same input data.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
rnnDesc	input	A previously initialized RNN descriptor.
seqLength	input	Number of iterations to unroll over.
yDesc	input	An array of fully packed tensor descriptors describing the output from each recurrent iteration. The second dimension of the tensor depends on the **direction** argument passed to the **cudnnSetRNNDescriptor** call used to initialize **rnnDesc**: - If **direction** is **CUDNN_UNIDIRECTIONAL** the second dimension should match the **hiddenSize** argument passed to **cudnnSetRNNDescriptor**. - If **direction** is **CUDNN_BIDIRECTIONAL** the second dimension should match double the **hiddenSize** argument passed to **cudnnSetRNNDescriptor**. The first dimension of the tensor **n** must match the first dimension of the tensor **n** in **dyDesc**.

Param	In/out	Meaning
y	input	Data pointer to GPU memory associated with the output tensor descriptor `yDesc`.
dyDesc	input	An array of fully packed tensor descriptors describing the gradient at the output from each recurrent iteration. The second dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the second dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the second dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The first dimension of the tensor `n` must match the second dimension of the tensor `n` in `dxDesc`.
dy	input	Data pointer to GPU memory associated with the tensor descriptors in the array `dyDesc`.
dhyDesc	input	A fully packed tensor descriptor describing the gradients at the final hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
dhy	input	Data pointer to GPU memory associated with the tensor descriptor `dhyDesc`. If a NULL pointer is passed, the gradients at the final hidden state of the network will be initialized to zero.
dcyDesc	input	A fully packed tensor descriptor describing the gradients at the final cell state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.

Param	In/out	Meaning
dcy	input	Data pointer to GPU memory associated with the tensor descriptor `dcyDesc`. If a NULL pointer is passed, the gradients at the final cell state of the network will be initialized to zero.
wDesc	input	Handle to a previously initialized filter descriptor describing the weights for the RNN.
w	input	Data pointer to GPU memory associated with the filter descriptor `wDesc`.
hxDesc	input	A fully packed tensor descriptor describing the initial hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
hx	input	Data pointer to GPU memory associated with the tensor descriptor `hxDesc`. If a NULL pointer is passed, the initial hidden state of the network will be initialized to zero.
cxDesc	input	A fully packed tensor descriptor describing the initial cell state for LSTM networks. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
cx	input	Data pointer to GPU memory associated with the tensor descriptor `cxDesc`. If a NULL pointer is passed, the initial cell state of the network will be initialized to zero.
dxDesc	input	An array of fully packed tensor descriptors describing the gradient at the input of each recurrent iteration. Each tensor descriptor must have the same second dimension. The first dimension of the tensors may decrease from element `n` to element `n+1` but may not increase.

Param	In/out	Meaning
dx	output	Data pointer to GPU memory associated with the tensor descriptors in the array `dxDesc`.
dhxDesc	input	A fully packed tensor descriptor describing the gradient at the initial hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
dhx	output	Data pointer to GPU memory associated with the tensor descriptor `dhxDesc`. If a NULL pointer is passed, the gradient at the hidden input of the network will not be set.
dcxDesc	input	A fully packed tensor descriptor describing the gradient at the initial cell state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
dcx	output	Data pointer to GPU memory associated with the tensor descriptor `dcxDesc`. If a NULL pointer is passed, the gradient at the cell input of the network will not be set.
workspace	input	Data pointer to GPU memory to be used as a workspace for this call.
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workspace`
reserveSpace	input/ output	Data pointer to GPU memory to be used as a reserve space for this call.
reserveSpaceSizeInBytes	input	Specifies the size in bytes of the provided `reserveSpace`

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.

Return Value	Meaning
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The descriptor `rnnDesc` is invalid. ▸ At least one of the descriptors `dhxDesc`, `wDesc`, `hxDesc`, `cxDesc`, `dcxDesc`, `dhyDesc`, `dcyDesc` or one of the descriptors in `yDesc`, `dxdesc`, `dydesc` is invalid. ▸ The descriptors in one of `yDesc`, `dxDesc`, `dyDesc`, `dhxDesc`, `wDesc`, `hxDesc`, `cxDesc`, `dcxDesc`, `dhyDesc`, `dcyDesc` has incorrect strides or dimensions. ▸ `workSpaceSizeInBytes` is too small. ▸ `reserveSpaceSizeInBytes` is too small.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.
`CUDNN_STATUS_ALLOC_FAILED`	The function was unable to allocate memory.

4.110. cudnnRNNBackwardWeights

```
cudnnStatus_t
cudnnRNNBackwardWeights( cudnnHandle_t handle,
                         const cudnnRNNDescriptor_t rnnDesc,
                         const int seqLength,

                         const cudnnTensorDescriptor_t * xDesc,
                         const void * x,
                         const cudnnTensorDescriptor_t hxDesc,
                         const void * hx,
                         const cudnnTensorDescriptor_t * yDesc,
                         const void * y,
                         const void * workspace,
                         size_t workSpaceSizeInBytes,
                         const cudnnFilterDescriptor_t dwDesc,
                         void * dw,
                         const void * reserveSpace,
                         size_t reserveSpaceSizeInBytes )
```

This routine accumulates weight gradients `dw` from the recurrent neural network described by `rnnDesc` with inputs `x, hx`, and outputs `y`. The mode of operation in this case is additive, the weight gradients calculated will be added to those already existing in `dw`. `workspace` is required for intermediate storage. The data in `reserveSpace` must have previously been generated by `cudnnRNNBackwardData`.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
rnnDesc	input	A previously initialized RNN descriptor.
seqLength	input	Number of iterations to unroll over.
xDesc	input	An array of fully packed tensor descriptors describing the input to each recurrent iteration. Each tensor descriptor must have the same second dimension. The first dimension of the tensors may decrease from element `n` to element `n+1` but may not increase.

Param	In/out	Meaning
x	input	Data pointer to GPU memory associated with the tensor descriptors in the array `xDesc`.
hxDesc	input	A fully packed tensor descriptor describing the initial hidden state of the RNN. The third dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the third dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the third dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The second dimension must match the second dimension of the tensors described in `xDesc`. The first dimension must match the `numLayers` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`. The tensor must be fully packed.
hx	input	Data pointer to GPU memory associated with the tensor descriptor `hxDesc`. If a NULL pointer is passed, the initial hidden state of the network will be initialized to zero.
yDesc	input	An array of fully packed tensor descriptors describing the output from each recurrent iteration. The second dimension of the tensor depends on the `direction` argument passed to the `cudnnSetRNNDescriptor` call used to initialize `rnnDesc`: ▸ If `direction` is `CUDNN_UNIDIRECTIONAL` the second dimension should match the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. ▸ If `direction` is `CUDNN_BIDIRECTIONAL` the second dimension should match double the `hiddenSize` argument passed to `cudnnSetRNNDescriptor`. The first dimension of the tensor `n` must match the first dimension of the tensor `n` in `dyDesc`.
y	input	Data pointer to GPU memory associated with the output tensor descriptor `yDesc`.
workspace	input	Data pointer to GPU memory to be used as a workspace for this call.
workSpaceSizeInBytes	input	Specifies the size in bytes of the provided `workspace`
dwDesc	input	Handle to a previously initialized filter descriptor describing the gradients of the weights for the RNN.
dw	input/output	Data pointer to GPU memory associated with the filter descriptor `dwDesc`.
reserveSpace	input	Data pointer to GPU memory to be used as a reserve space for this call.
reserveSpaceSizeInBytes	input	Specifies the size in bytes of the provided `reserveSpace`

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The function launched successfully.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The descriptor `rnnDesc` is invalid. ▸ At least one of the descriptors `hxDesc`, `dwDesc` or one of the descriptors in `xDesc`, `yDesc` is invalid. ▸ The descriptors in one of `xDesc`, `hxDesc`, `yDesc`, `dwDesc` has incorrect strides or dimensions. ▸ `workSpaceSizeInBytes` is too small. ▸ `reserveSpaceSizeInBytes` is too small.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.
`CUDNN_STATUS_ALLOC_FAILED`	The function was unable to allocate memory.

4.111. cudnnCreateDropoutDescriptor

```
cudnnStatus_t cudnnCreateDropoutDescriptor(cudnnRNNDescriptor_t * rnnDesc)
```

This function creates a generic dropout descriptor object by allocating the memory needed to hold its opaque structure.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was created successfully.
`CUDNN_STATUS_ALLOC_FAILED`	The resources could not be allocated.

4.112. cudnnDestroyDropoutDescriptor

```
cudnnStatus_t cudnnDestroyDropoutDescriptor(cudnnDropoutDescriptor_t rnnDesc)
```

This function destroys a previously created dropout descriptor object.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The object was destroyed successfully.

4.113. cudnnDropoutGetStatesSize

```
cudnnStatus_t
cudnnDropoutGetStatesSize( cudnnHandle_t handle,
                           size_t * sizeInBytes);
```

This function is used to query the amount of space required to store the states of the random number generators used by **cudnnDropoutForward** function.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
sizeInBytes	output	Amount of GPU memory needed to store random generator states.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.

4.114. cudnnDropoutGetReserveSpaceSize

```
cudnnStatus_t
cudnnDropoutGetReserveSpaceSize( cudnnTensorDescriptor_t xDesc,
                                 size_t * sizeInBytes);
```

This function is used to query the amount of reserve needed to run dropout with the input dimensions given by **xDesc**. The same reserve space is expected to be passed to **cudnnDropoutForward** and **cudnnDropoutBackward**, and its contents is expected to remain unchanged between **cudnnDropoutForward** and **cudnnDropoutBackward** calls.

Param	In/out	Meaning
xDesc	input	Handle to a previously initialized tensor descriptor, describing input to a dropout operation.
sizeInBytes	output	Amount of GPU memory needed as reserve space to be able to run dropout with an input tensor descriptor specified by xDesc.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The query was successful.

4.115. cudnnSetDropoutDescriptor

```
cudnnStatus_t
cudnnSetDropoutDescriptor( cudnnDropoutDescriptor_t dropoutDesc,
                           cudnnHandle_t handle,
                           float dropout,
                           void * states,
                           size_t stateSizeInBytes,
                           unsigned long long seed)
```

This function initializes a previously created dropout descriptor object. If **states** argument is equal to NULL, random number generator states won't be initialized, and only **dropout** value will be set. No other function should be writing to the memory

pointed at by **states** argument while this function is running. The user is expected not to change memory pointed at by **states** for the duration of the computation.

Param	In/out	Meaning
dropoutDesc	input/ output	Previously created dropout descriptor object.
handle	input	Handle to a previously created cuDNN context.
dropout	input	The probability with which the value from input would be propagated through the dropout layer.
states	output	Pointer to user-allocated GPU memory that will hold random number generator states.
sizeInBytes	input	Specifies size in bytes of the provided memory for the states
seed	input	Seed used to initialize random number generator states.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The call was successful.
CUDNN_STATUS_INVALID_VALUE	**sizeInBytes** is less than the value returned by **cudnnDropoutGetStatesSize.**
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU

4.116. cudnnDropoutForward

```
cudnnStatus_t
cudnnDropoutForward( cudnnHandle_t handle,
                    const cudnnDropoutDescriptor_t dropoutDesc,
                    const cudnnTensorDescriptor_t xdesc,
                    const void * x,
                    const cudnnTensorDescriptor_t ydesc,
                    void * y,
                    void * reserveSpace,
                    size_t reserveSpaceSizeInBytes)
```

This function performs forward dropout operation over **x** returning results in **y**. If **dropout** was used as a parameter to **cudnnSetDropoutDescriptor**, the approximately **dropout** fraction of **x** values will be replaces by **0**, and the rest will be scaled by **1/(1-dropout)** This function should not be running concurrently with another **cudnnDropoutForward** function using the same **states**.

💬 Better performance is obtained for fully packed tensors

💬 Should not be called during inference

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
dropoutDesc	input	Previously created dropout descriptor object.
xDesc	input	Handle to a previously initialized tensor descriptor.
x	input	Pointer to data of the tensor described by the `xDesc` descriptor.
yDesc	input	Handle to a previously initialized tensor descriptor.
y	output	Pointer to data of the tensor described by the `yDesc` descriptor.
reserveSpace	output	Pointer to user-allocated GPU memory used by this function. It is expected that contents of `reserveSpace` doe not change between `cudnnDropoutForward` and `cudnnDropoutBackward` calls.
reserveSpaceSizeInBytes	input	Specifies size in bytes of the provided memory for the reserve space

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The call was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ The number of elements of input tensor and output tensors differ. ▸ The `datatype` of the input tensor and output tensors differs. ▸ The strides of the input tensor and output tensors differ and in-place operation is used (i.e., `x` and `y` pointers are equal). ▸ The provided `reserveSpaceSizeInBytes` is less then the value returned by `cudnnDropoutGetReserveSpaceSize` ▸ `cudnnSetDropoutDescriptor` has not been called on `dropoutDesc` with the non-NULL `states` argument
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

4.117. cudnnDropoutBackward

```
cudnnStatus_t
cudnnDropoutBackward( cudnnHandle_t handle,
                      const cudnnDropoutDescriptor_t dropoutDesc,
                      const cudnnTensorDescriptor_t dydesc,
                      const void * dy,
                      const cudnnTensorDescriptor_t dxdesc,
                      void * dx,
                      void * reserveSpace,
                      size_t reserveSpaceSizeInBytes)
```

This function performs backward dropout operation over **dy** returning results in **dx**. If during forward dropout operation value from **x** was propagated to **y** then during backward operation value from **dy** will be propagated to **dx**, otherwise, **dx** value will be set to **0**.

💬 Better performance is obtained for fully packed tensors

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
dropoutDesc	input	Previously created dropout descriptor object.
dyDesc	input	Handle to a previously initialized tensor descriptor.
dy	input	Pointer to data of the tensor described by the **dyDesc** descriptor.
dxDesc	input	Handle to a previously initialized tensor descriptor.
dx	output	Pointer to data of the tensor described by the **dxDesc** descriptor.
reserveSpace	input	Pointer to user-allocated GPU memory used by this function. It is expected that **reserveSpace** was populated during a call to **cudnnDropoutForward** and has not been changed.
reserveSpaceSizeInBytes	input	Specifies size in bytes of the provided memory for the reserve space

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The call was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ The number of elements of input tensor and output tensors differ. ▸ The **datatype** of the input tensor and output tensors differs. ▸ The strides of the input tensor and output tensors differ and in-place operation is used (i.e., **x** and **y** pointers are equal). ▸ The provided **reserveSpaceSizeInBytes** is less then the value returned by **cudnnDropoutGetReserveSpaceSize** ▸ **cudnnSetDropoutDescriptor** has not been called on **dropoutDesc** with the non-NULL **states** argument
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.118. cudnnCreateSpatialTransformerDescriptor

```
cudnnStatus_t
cudnnCreateSpatialTransformerDescriptor(
                cudnnSpatialTransformerDescriptor_t *stDesc)
```

This function creates a generic spatial transformer descriptor object by allocating the memory needed to hold its opaque structure.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was created successfully.
CUDNN_STATUS_ALLOC_FAILED	The resources could not be allocated.

4.119. cudnnDestroySpatialTransformerDescriptor

```
cudnnStatus_t
cudnnDestroySpatialTransformerDescriptor(
                cudnnSpatialTransformerDescriptor_t stDesc)
```

This function destroys a previously created spatial transformer descriptor object.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The object was destroyed successfully.

4.120. cudnnSetSpatialTransformerNdDescriptor

```
cudnnStatus_t
cudnnSetSpatialTransformerNdDescriptor(
                cudnnSpatialTransformerDescriptor_t        stDesc,
                cudnnSamplerType_t                         samplerType,
                cudnnDataType_t                            dataType,
                const int                                  nbDims,
                const int                                  dimA[]);
```

This function initializes a previously created generic spatial transformer descriptor object.

Param	In/out	Meaning
stDesc	input/output	Previously created spatial transformer descriptor object.
samplerType	input	Enumerant to specify the sampler type.
dataType	input	Data type.
nbDims	input	Dimension of the transformed tensor.
dimA	input	Array of dimension **nbDims** containing the size of the transformed tensor for every dimension.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The call was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ Either **stDesc** or **dimA** is NULL. ▸ Either **dataType** or **samplerType** has an invalid enumerant value

4.121. cudnnSpatialTfGridGeneratorForward

```
cudnnStatus_t
cudnnSpatialTfGridGeneratorForward(
        cudnnHandle_t                           handle,
        const cudnnSpatialTransformerDescriptor_t stDesc,
        const void*                             theta,
        void*                                   grid)
```

This function generates a grid of coordinates in the input tensor corresponding to each pixel from the output tensor.

Only 2d transformation is supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
stDesc	input	Previously created spatial transformer descriptor object.
theta	input	Affine transformation matrix. It should be of size n*2*3 for a 2d transformation, where n is the number of images specified in **stDesc**.
grid	output	A grid of coordinates. It is of size n*h*w*2 for a 2d transformation, where n, h, w is specified in **stDesc**. In the 4th dimension, the first coordinate is x, and the second coordinate is y.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The call was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ **handle** is NULL. ▸ One of the parameters **grid**, **theta** is NULL.
CUDNN_STATUS_NOT_SUPPORTED	The dimension of transformed tensor specified in **stDesc** > 4.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.122. cudnnSpatialTfGridGeneratorBackward

```
cudnnStatus_t
cudnnSpatialTfGridGeneratorBackward(
        cudnnHandle_t                               handle,
        const cudnnSpatialTransformerDescriptor_t   stDesc,
        const void*                                 dgrid,
        void*                                       dtheta)
```

This function computes the gradient of a grid generation operation.

> Only 2d transformation is supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
stDesc	input	Previously created spatial transformer descriptor object.
dgrid	input	Data pointer to GPU memory contains the input differential data.
dtheta	output	Data pointer to GPU memory contains the output differential data.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The call was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ **handle** is NULL. ▸ One of the parameters **dgrid**, **dtheta** is NULL.
CUDNN_STATUS_NOT_SUPPORTED	The dimension of transformed tensor specified in **stDesc** > 4.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.123. cudnnSpatialTfSamplerForward

```
cudnnStatus_t
cudnnSpatialTfSamplerForward(
        cudnnHandle_t                               handle,
        const cudnnSpatialTransformerDescriptor_t   stDesc,
        const void*                                 alpha,
        const cudnnTensorDescriptor_t               xDesc,
        const void*                                 x,
        const void*                                 grid,
        const void*                                 beta,
        cudnnTensorDescriptor_t                     yDesc,
        void*                                       y)
```

This function performs a sampler operation and generates the output tensor using the grid given by the grid generator.

Only 2d transformation is supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
stDesc	input	Previously created spatial transformer descriptor object.
alpha,beta	input	Pointers to scaling factors (in host memory) used to blend the source value with prior value in the destination tensor as follows: dstValue = alpha[0]*srcValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor **xDesc**.
grid	input	A grid of coordinates generated by **cudnnSpatialTfGridGeneratorForward**.
yDesc	input	Handle to the previously initialized output tensor descriptor.
y	output	Data pointer to GPU memory associated with the output tensor descriptor **yDesc**.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
CUDNN_STATUS_SUCCESS	The call was successful.
CUDNN_STATUS_BAD_PARAM	At least one of the following conditions are met: ▸ **handle** is NULL. ▸ One of the parameters **x**, **y**, **grid** is NULL.
CUDNN_STATUS_NOT_SUPPORTED	The dimension of transformed tensor > 4.
CUDNN_STATUS_EXECUTION_FAILED	The function failed to launch on the GPU.

4.124. cudnnSpatialTfSamplerBackward

```
cudnnStatus_t
cudnnSpatialTfSamplerBackward(
        cudnnHandle_t                             handle,
        const cudnnSpatialTransformerDescriptor_t stDesc,
        const void*                               alpha,
        const cudnnTensorDescriptor_t             xDesc,
        const void*                               x,
        const void*                               beta,
        const cudnnTensorDescriptor_t             dxDesc,
        void*                                     dx,
        const void*                               alphaDgrid,
        const cudnnTensorDescriptor_t             dyDesc,
        const void*                               dy,
        const void*                               grid,
        const void*                               betaDgrid,
        void*                                     dgrid)
```

This function computes the gradient of a sampling operation.

> Only 2d transformation is supported.

Param	In/out	Meaning
handle	input	Handle to a previously created cuDNN context.
stDesc	input	Previously created spatial transformer descriptor object.
alpha,beta	input	Pointers to scaling factors (in host memory) used to blend the source value with prior value in the destination tensor as follows: dstValue = alpha[0]*srcValue + beta[0]*priorDstValue. Please refer to this section for additional details.
xDesc	input	Handle to the previously initialized input tensor descriptor.
x	input	Data pointer to GPU memory associated with the tensor descriptor `xDesc`.
dxDesc	input	Handle to the previously initialized output differential tensor descriptor.
dx	output	Data pointer to GPU memory associated with the output tensor descriptor `dxDesc`.
alphaDgrid,betaDgrid	input	Pointers to scaling factors (in host memory) used to blend the gradient outputs dgrid with prior value in the destination pointer as follows: dstValue = alpha[0]*srcValue + beta[0]*priorDstValue. Please refer to this section for additional details.
dyDesc	input	Handle to the previously initialized input differential tensor descriptor.
dy	input	Data pointer to GPU memory associated with the tensor descriptor `dyDesc`.

Param	In/out	Meaning
grid	input	A grid of coordinates generated by `cudnnSpatialTfGridGeneratorForward`.
dgrid	output	Data pointer to GPU memory contains the output differential data.

The possible error values returned by this function and their meanings are listed below.

Return Value	Meaning
`CUDNN_STATUS_SUCCESS`	The call was successful.
`CUDNN_STATUS_BAD_PARAM`	At least one of the following conditions are met: ▸ `handle` is NULL. ▸ One of the parameters `x,dx,y,dy,grid,dgrid` is NULL. ▸ The dimension of `dy` differs from those specified in `stDesc`
`CUDNN_STATUS_NOT_SUPPORTED`	The dimension of transformed tensor > 4.
`CUDNN_STATUS_EXECUTION_FAILED`	The function failed to launch on the GPU.

Chapter 5.
ACKNOWLEDGMENTS

Some of the cuDNN library routines were derived from code developed by others and are subject to the following:

5.1. University of Tennessee

```
Copyright (c) 2010 The University of Tennessee.

All rights reserved.

Redistribution and use in source and binary forms, with or without
modification, are permitted provided that the following conditions are
met:
    * Redistributions of source code must retain the above copyright
      notice, this list of conditions and the following disclaimer.
    * Redistributions in binary form must reproduce the above
      copyright notice, this list of conditions and the following
      disclaimer listed in this license in the documentation and/or
      other materials provided with the distribution.
    * Neither the name of the copyright holders nor the names of its
      contributors may be used to endorse or promote products derived
      from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS
"AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT
LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR
A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT
OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL,
SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT
LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE,
DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY
THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT
(INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE
OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.
```

5.2. University of California, Berkeley

5.3. Facebook AI Research, New York